

HypatiaX: A Hybrid Symbolic-Neural Framework for Extrapolation-Reliable Analytical Discovery
Bonet Chaple

HypatiaX: A Hybrid Symbolic-Neural Framework for Extrapolation-Reliable Analytical Discovery

Ruperto Pedro Bonet Chaple

ruperto.bonet@modelphysmat.com

Independent Researcher, London, United Kingdom

Editor:

Note on runtime claims. An earlier draft of this paper asserted a “73% runtime reduction” (i.e. a $3.7\times$ speedup of HypatiaX over the neural-network baseline). That figure is *not supported* by the benchmark data (HypatiaX v3.0). The verified runtime comparison is presented in full in Section 10.4; the honest claim is a $1.73\times$ speedup on LLM-routed cases only ($n = 68$ out of 74). All other quantitative claims in this paper have been independently verified against the v3.0 statistical report and are reported exactly as measured.

Abstract

The hardest test of a formula-recovery system is not interpolation accuracy but out-of-distribution extrapolation. Two cases from our benchmark make this stark. On Michaelis-Menten, PySR-only collapses catastrophically at $5\times$ the training range ($\text{far-}R^2 = -83,900$); HypatiaX anchors the search via an LLM warm-start and recovers the correct functional form, reducing $\text{far-}R^2$ to -635 —a $132\times$ reduction in error magnitude. On Portfolio Variance, PySR-only fails on all five tested random seeds; HypatiaX recovers the exact closed-form formula ($R^2 = 1.000$ at $5\times$ training range) on seeds 777 and 2024. We also document an honest failure mode: on Arrhenius and Portfolio Variance (seeds 42/99/123), LLM warm-starting constrains PySR to a local minimum, leaving extrapolation no better than PySR-only despite higher training R^2 . These two phenomena—selective rescue and selective degradation—are the empirical core of this paper.

We introduce **HypatiaX**, a hybrid symbolic–neural framework that addresses this through a five-stage routing and ensembling mechanism—trust gating, transcendental complexity detection, neural degradation probing, out-of-distribution detection, and uncertainty-weighted ensembling—that dynamically selects between LLM symbolic output and a residual neural network using training-data diagnostics alone. The key insight is that the two paradigm failure modes are complementary: routing between them eliminates catastrophic failures while preserving the accuracy ceiling of the stronger component.

On a benchmark of 74 DeFi mathematical tasks under an aggressive 40%/60% PCA-directed extrapolation split, a standalone neural MLP achieves median test $R^2 = -0.47$, and a pure LLM baseline suffers 6 catastrophic failures ($R^2 < -10$) while achieving only 62.2% near-perfect recovery. Neither paradigm alone is sufficient.

Across all 74 tasks, HypatiaX achieves a median test R^2 of 1.00 and a near-perfect success rate ($R^2 > 0.99$) of **89.2%**—a +27 percentage point gain over the LLM baseline and +83.8pp over the neural network—while **eliminating all catastrophic failures**. Gains

scale with difficulty: +38.1 pp on hard transcendental tasks (76.2% vs. 38.1% for the LLM alone). On the standard Nguyen-12 suite, HypatiaX achieves 11/12 recovery (91.7%) versus 10/12 for PySR-only and 0/12 for the neural baseline. When the LLM formula is trusted (68 of 74 cases, 92%), HypatiaX delivers a 1.73 $\times$ median speedup over neural-network inference. On 30 Feynman equations spanning mechanics to quantum physics, HypatiaX achieves 9/30 near-perfect recovery versus 0/30 for the neural baseline (Kaggle 4-vCPU run; 142/180 under the full PCA 40/60 instance pool — see Section 10.7), establishing that reliability gains extend beyond the DeFi domain. These results establish *reliability under distribution shift*, not raw accuracy, as the defining advantage of hybrid symbolic–neural systems.

Contents

1	Introduction	5
1.1	Contributions	5
2	Related Work	6
2.1	Symbolic Regression	6
2.2	LLMs for Mathematical and Scientific Reasoning	6
2.3	Equation Learners, Differentiable SR, and Uncertainty-Aware Methods . . .	6
2.4	Hybrid Symbolic–Neural Systems	7
2.5	Extrapolation in Neural Networks	7
3	Empirical Evidence of LLM Performance Across Domains	7
3.1	Experimental Design	7
3.2	Baseline Results: Bimodal Performance	8
3.3	Failure Mode Taxonomy	8
3.4	Case Studies	8
3.5	Success Pattern Analysis	8
3.6	Extrapolation Testing	9
4	Theoretical Framework	9
4.1	Formal Definitions	9
4.2	Main Theoretical Result	9
4.3	Implications for Discovery	10
4.4	Functional Form Recovery and Inductive Bias Alignment	10
5	Problem Formulation	10
6	Benchmark Design	11
6.1	Overview and Scope	11
6.2	Difficulty Classification	12
6.3	Data Generation	12
6.4	Extrapolation Protocol	12
6.5	Pipeline Corrections in v3.0	13

7	Methodology	14
7.1	Architecture Overview	14
7.2	Component 1: LLM-Based Symbolic Generation	14
7.3	Component 2: Neural Network Estimator	15
7.4	Component 3: Five-Stage Routing and Ensembling	17
7.5	Unified Formula Executor	19
8	HypatiaX Architecture	19
8.1	Design Principles	19
8.2	Assumptions	20
8.3	Five-Stage Routing Architecture Overview	20
8.4	System Variants	20
8.5	Hybrid DeFi Variant: Architectural Specification	20
8.6	Symbolic Discovery Core	21
8.7	Multi-Layer Validation Framework	21
9	Experimental Setup	22
9.1	Benchmark Summary	22
9.2	Methods Compared	22
9.3	Evaluation Metrics	22
9.4	Runtime Measurement	23
9.5	Implementation Details	23
10	Results	23
10.1	Five-System Comparative Analysis (Feynman Core-15)	23
10.2	Overall Extrapolation Performance (DeFi v3.0)	24
10.3	Performance by Difficulty Level	25
10.4	Runtime Analysis	27
10.5	Portfolio Variance: Seed-Sweep Robustness Analysis	28
10.6	Ablation: PySR-only vs. HypatiaX (Core-15)	29
10.7	Feynman Extrapolation Benchmark ($n = 30$)	31
10.8	Nguyen-12 Benchmark (Standard SR Suite)	31
10.9	Stability Under Stochastic Inference	33
11	Discussion	34
11.1	Three Core Findings	34
11.2	Arrhenius and Scale-Incompatibility Failures	35
11.3	The Stability–Accuracy Distinction	35
11.4	Why Transcendental Formulas Fail	36
11.5	OOD Metric as a Proxy	36
11.6	Design Principles for Analytical Discovery	36
11.7	Comparison with Related Work	37
11.8	Ethical Considerations	37
11.9	Broader Applicability	38
11.10	Limitations and Future Work	38

12 Conclusion	38
A Full Benchmark Case Definitions	41
A.1 Automated Market Makers	41
A.2 Lending Protocols	42
A.3 Derivatives Pricing	42
A.4 Risk and Portfolio Metrics	43
A.5 Staking and Protocol Mechanisms	43
B Benchmark Version History	44
C Corrected Runtime Claim: Detailed Analysis	44

1 Introduction

Extrapolation beyond the observed training distribution is a fundamental challenge in machine learning. In financial and decentralised-finance (DeFi) domains, the problem is particularly acute: pricing models, risk metrics, and market-maker invariants must generalise correctly to input regimes that may differ dramatically from observed historical data. A system that interpolates well but extrapolates poorly is of limited practical value in this setting.

Two broad paradigms have been applied to mathematical function recovery in such domains:

Neural network regression. Multilayer perceptrons and related architectures are powerful interpolators but are known to fail systematically under distribution shift. Even with augmentation techniques, neural networks trained on $\mathcal{D}_{\text{train}}$ do not, in general, recover the functional form of f beyond the training convex hull. Our benchmark confirms this: the neural MLP baseline achieves a median test R^2 of -0.47 under extrapolation—worse than predicting the mean.

LLM-based symbolic regression. Large language models can, in principle, recover exact closed-form expressions from task descriptions, achieving perfect extrapolation when the correct formula is produced. However, this approach exhibits a severe bimodal failure pattern: on a benchmark of 74 DeFi tasks with $K = 30$ independent runs per task, LLMs achieve median test $R^2 = 1.00$ but mean $R^2 = -0.76$, with 6 catastrophic failures ($R^2 < -10$). More importantly, a single-run evaluation—the standard in prior work—conceals run-to-run instability: on complex transcendental tasks, the same prompt yields radically different expressions across runs (Instability Index $\text{II} > 0.10$).

Central claim. Neither paradigm alone is sufficient. The key insight of HypatiaX is that these two failure modes are *complementary*: neural networks fail on the tasks where LLMs succeed (transcendental extrapolation), and LLMs fail on the tasks where a neural network with LLM-augmented features can stabilise performance. A routing mechanism that dynamically selects between the two modes—based on observable diagnostics computed from training data alone—can eliminate catastrophic failures while preserving near-perfect accuracy on the cases where symbolic recovery is reliable.

1.1 Contributions

This paper makes the following contributions:

- (i) **HypatiaX benchmark:** a corpus of 74 closed-form DeFi tasks with an aggressive PCA-directed extrapolation split (40%/60%), together with a unified evaluation pipeline (HypatiaX v3.0) that corrects four systematic biases in prior versions (data leakage, split misconfiguration, formula evaluation inconsistency, and weak trust gating).
- (ii) **HypatiaX hybrid system:** a five-stage routing and ensembling mechanism—trust gating, transcendental token detection, neural degradation probing, OOD detection, and uncertainty-weighted ensembling—that achieves 89.2% near-perfect success rate ($R^2 > 0.99$) across all 74 tasks.

- (iii) **Verified runtime analysis:** a rigorous decomposition of wall-clock time showing that the $1.73\times$ speedup of HypatiaX over the neural baseline applies specifically to the 68 LLM-routed cases, and that aggregate mean time (6.8s) is higher than the neural baseline (3.0s) due to fallback costs. This corrects a previously circulated 73 % reduction claim that was not supported by the v3.0 benchmark data.
- (iv) **Difficulty-stratified analysis:** a breakdown of performance across easy (24 tasks), medium (29 tasks), and hard (21 tasks) cases, demonstrating that HypatiaX’s gains are largest and most practically important on the hardest transcendental and multi-asset formulas.

2 Related Work

2.1 Symbolic Regression

Symbolic regression aims to discover explicit functional forms from data without assuming a fixed parametric model. Classical methods rely on genetic programming (Koza, 1992) and evolutionary search; the SRBench benchmark (La Cava et al., 2021) provides a standardised multi-dataset evaluation suite. The AI Feynman system (Udrescu & Tegmark, 2020) demonstrated that dimensional analysis and neural pre-screening can guide symbolic search toward physically meaningful equations.

A critical limitation of the symbolic regression literature is the evaluation paradigm: systems are typically assessed on a single run or on the best expression found across multiple runs. This conceals run-to-run instability, which—as we show—is the dominant failure mode of LLM-based symbolic regression on complex formulas. Our multi-run protocol (30 independent runs per task) and Instability Index directly address this gap.

2.2 LLMs for Mathematical and Scientific Reasoning

Chain-of-thought prompting (Wei et al., 2022) and Minerva (Lewkowycz et al., 2022) demonstrate that LLMs can solve complex mathematical problems. Codex (Chen et al., 2021) showed strong code and formula generation capability. More recent work has explored LLMs as symbolic regression engines, generating formulas from natural-language descriptions of datasets.

All of these works evaluate correctness in a single-shot or best-of- K setting and do not report variance across runs as a first-class metric. Our benchmark is the first, to our knowledge, to quantify LLM instability systematically across a large corpus of financial formulas under repeated stochastic inference.

2.3 Equation Learners, Differentiable SR, and Uncertainty-Aware Methods

Several strands of recent work are directly relevant to the hybrid design in this paper. The Equation Learner (EQL; Sahoo et al. 2018) embeds differentiable operator networks with L_0 regularisation to produce sparse, interpretable expressions, demonstrating that structural bias can be enforced without explicit symbolic search. DISCOVER (Liu et al., 2022) reformulates symbolic regression as differentiable optimisation over a symbolic search space, enabling gradient-based recovery of functional form. Bayesian and Gaussian-process

approaches to symbolic regression (Jin et al., 2020) provide principled uncertainty estimates alongside recovered expressions — a goal shared by the confidence-weighted ensemble in HYPATIAx’s routing cascade. HYPATIAx differs from these works by treating *extrapolation reliability* as a first-class evaluation criterion alongside recovery rate, and by combining evolutionary symbolic search with an explicit multi-layer validation cascade that filters expressions before acceptance.

2.4 Hybrid Symbolic–Neural Systems

The idea of combining symbolic structure with neural learning has a rich history. Cranmer et al. (2020) use neural networks to discover Lagrangians, then fit symbolic expressions to the learned dynamics. Udrescu & Tegmark (2020) alternate between neural pre-screening and symbolic search. PySR (Cranmer, 2023) provides an efficient evolutionary-symbolic regression framework that can be combined with neural features.

Our approach differs in motivation: HypatiaX is designed not to improve accuracy over either pure method in isolation, but to *eliminate catastrophic failures* and achieve reliability across the full difficulty spectrum of DeFi tasks.

2.5 Extrapolation in Neural Networks

The failure of neural networks to extrapolate is well-documented (Xu et al., 2021). Data augmentation with scaled inputs (Lim et al., 2021) can improve extrapolation marginally but does not address the fundamental limitation. Our results confirm that even a residual MLP with LayerNorm, SiLU activations, and scaled augmentation achieves median $R^2 = -0.47$ under our 40%/60% extrapolation split.

3 Empirical Evidence of LLM Performance Across Domains

To evaluate whether large language models (LLMs) can perform analytical formula discovery without symbolic constraints or external solvers, we conducted a systematic study of pure LLM-based formula generation (Koza, 1992; Kambhampati, 2024). No symbolic regression, constraint enforcement, or post-generation correction was applied.

3.1 Experimental Design

We evaluated GPT-4 (OpenAI, 2023) on 40 interpolation-based formula discovery tasks spanning five domains: classical physics (Kinetic Energy, Gravitational Force, Ideal Gas Law), chemistry (Arrhenius, Henderson-Hasselbalch, Rate Law), biology (Allometric Scaling, Michaelis-Menten, Logistic Growth), and decentralized finance (Impermanent Loss, Price Impact, Portfolio Variance, Value-at-Risk, Liquidation Price). Tasks were presented as structured prompts containing variable names, units, and sample data. Outputs were evaluated by executing the generated formula symbolically and computing R^2 on a held-out test set within the training range.

3.2 Baseline Results: Bimodal Performance

Pure LLM performance is strongly bimodal. On well-known physical laws (Kinetic Energy, Ideal Gas Law, Gravitational Force), the LLM achieves $R^2 = 1.000$ with exact symbolic recovery. On domain-specific and specialized equations—particularly in chemistry (Henderson-Hasselbalch: $R^2 = -3.2$), biology (Michaelis-Menten: $R^2 = -68.6$), and DeFi (Portfolio Variance: $R^2 = -21.0$)—performance collapses catastrophically. LLM success correlates with estimated equation prevalence in web corpora (Spearman $\rho = 0.89$, $p < 0.001$), consistent with the hypothesis that pretraining data distribution influences recovery performance (Kambhampati, 2024).

3.3 Failure Mode Taxonomy

We identify five failure modes, ordered by severity:

1. **Silent semantic errors** (High): Structurally plausible formula with incorrect constants or signs; high in-range R^2 but catastrophic extrapolation.
2. **Distributional reasoning failures** (High): Misuse of statistical functions (e.g. wrong sign for VaR z-score).
3. **Unit/scale inconsistency** (Medium): Missing unit conversions or scale factors.
4. **Non-executable output** (Medium): Pseudocode or invalid Python; cannot be evaluated.
5. **Incomplete construction** (Medium): Partial derivations missing key terms.

3.4 Case Studies

Arrhenius equation. The Arrhenius rate constant is $k = A \exp(-E_a/RT)$. The LLM correctly identifies the exponential structure but systematically misestimates the activation energy E_a by a constant factor, producing train $R^2 = 0.997$ but extrapolation $R^2 = -12.6$ at $5\times$ the training range. This is the archetypal *scale-incompatibility failure*: the formula is structurally correct but parametrically wrong in a way that only manifests under distribution shift.

Michaelis-Menten. $v = V_{\max}[S]/(K_m + [S])$. The LLM recovers the functional form on some runs ($R^2 = 1.000$) but collapses on others ($R^2 < -68$), producing high run-to-run instability ($\text{II} = 0.31$). This motivates the Instability Index as a diagnostic.

3.5 Success Pattern Analysis

LLM guidance succeeds when: (1) the equation is well-represented in web text; (2) the functional form is algebraic rather than transcendental; (3) the domain is classical physics or standard chemistry. It fails systematically on post-2020 DeFi equations, multi-asset risk formulas, and transcendental biology equations—precisely the cases where symbolic regression provides the most value.

3.6 Extrapolation Testing

Even when the LLM achieves near-perfect interpolation ($R^2 \approx 1$), extrapolation performance is unreliable. Across 40 tests, 6 equations suffered catastrophic extrapolation failure ($R^2 < -10$). This motivates the hybrid architecture: the LLM provides structural guidance while symbolic regression corrects parametric errors and the neural network provides a stable fallback.

4 Theoretical Framework

4.1 Formal Definitions

Definition 1 (Analytical Expression Discovery) Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ where $x_i \in \mathbb{R}^d$, $y_i \in \mathbb{R}$ (Udrescu & Tegmark, 2020). An analytical expression discovery system produces $E : \mathbb{R}^d \rightarrow \mathbb{R}$ satisfying:

1. *Symbolic form:* E is expressible using operators $\mathcal{O}$ in closed form
2. *Physical consistency:* E satisfies domain constraints $\mathcal{C}$
3. *Empirical accuracy:* Loss $\mathcal{L}(E, \mathcal{D}) < \epsilon$
4. *Extrapolation:* E generalizes beyond training regime
5. *Parsimony:* E has minimal complexity given above constraints

Definition 2 (Interpolation–Extrapolation Decoupling) A system exhibits interpolation–extrapolation decoupling if $\exists$ tasks where $R_{\text{train}}^2 \approx 1$ but $R_{\text{test}}^2 \ll 1$ on an out-of-distribution test set.

Definition 3 (Reproductive Behaviour) A generative model G exhibits reproductive behaviour if its output distribution concentrates near its training manifold: $P(G(x) \notin \mathcal{M}_{\text{train}}) < \delta$ for small δ , where $\mathcal{M}_{\text{train}}$ is the training data manifold.

4.2 Main Theoretical Result

Empirical Finding 1 (Interpolation–Extrapolation Decoupling) Neural networks trained on in-distribution data can achieve near-perfect interpolation ($R_{\text{train}}^2 \approx 1$) while exhibiting catastrophic extrapolation failure. Symbolic methods that recover the exact functional form are bounded by numerical precision alone.

Proof [Proof sketch] Neural networks achieve high in-distribution R^2 via local approximations (universal approximation (Hornik et al., 1989)). These approximations are not constrained to match the true functional form; extrapolation requires the correct form. Symbolic methods that recover the exact f^* achieve $R_{\text{test}}^2 = 1 - O(\epsilon_{\text{num}})$ by construction. ■

Empirical Finding 2 (LLM Reproductive Behaviour) *Autoregressive LLMs trained by maximum-likelihood exhibit reproductive behaviour: output probability concentrates near training-corpus examples. For analytical expression discovery, this implies performance degrades on expressions underrepresented in pretraining data.*

Proof [Proof sketch] Maximum-likelihood training minimises $-\mathbb{E}[\log P_\theta(x)]$, concentrating probability mass near training examples. For rare equations (low corpus frequency), the model assigns low probability to the correct expression, increasing the risk of plausible-but-wrong outputs. ■

4.3 Implications for Discovery

Finding 2 explains the domain-dependence pattern in Section 3: LLM performance correlates with equation prevalence ($\rho = 0.89$) because reproductive behaviour ties output quality to training exposure. This motivates hybrid architectures that use LLM structural guidance while correcting via symbolic search.

4.4 Functional Form Recovery and Inductive Bias Alignment

Symbolic regression performs evolutionary exploration of expression space guided by fitness evaluation, without requiring proximity to textual training examples. For problems admitting closed-form representations, restricting search to symbolic structures improves extrapolation reliability by aligning the inductive bias of the model with the structure of the target function. By the No Free Lunch theorem (Wolpert & Macready, 1997), performance depends on alignment between hypothesis class and problem structure. When the true relationship admits a closed-form representation, this alignment holds for symbolic search and fails for purely neural approaches.

5 Problem Formulation

Definition 4 (Symbolic Regression Task) *A task i is defined by a dataset $\mathcal{D}_i = \{(x^{(j)}, y^{(j)})\}_{j=1}^N$ with ground-truth function $f_i : \mathbb{R}^{d_i} \rightarrow \mathbb{R}$ such that $y^{(j)} = f_i(x^{(j)})$. The objective is to recover a symbolic expression $\hat{f}_i$ such that $\hat{f}_i \approx f_i$ on held-out test data.*

Definition 5 (Extrapolation Setting) *Let $\mathcal{D}_i^{\text{train}}$ and $\mathcal{D}_i^{\text{test}}$ be a partition of $\mathcal{D}_i$ such that the test inputs lie beyond the training range along the principal direction of input variation:*

$$X^{\text{test}} \notin [\min(X^{\text{train}}), \max(X^{\text{train}})] \quad (\text{approximately, along the first PC}). \quad (1)$$

Evaluation on $\mathcal{D}_i^{\text{test}}$ under condition (1) is called extrapolation evaluation.

Definition 6 (Coefficient of Determination) *For a predictor $\hat{f}$ evaluated on a set $\mathcal{S}$:*

$$R^2(\hat{f}, \mathcal{S}) = 1 - \frac{\sum_{(x,y) \in \mathcal{S}} (y - \hat{f}(x))^2}{\sum_{(x,y) \in \mathcal{S}} (y - \bar{y}_{\mathcal{S}})^2}. \quad (2)$$

For aggregate statistics, R^2 values are clipped to $[-10, 1]$ to prevent extreme failures from dominating mean computations. All reported aggregates use a fixed denominator of 74 (all standard cases), with NaN and failure cases counted as 0.

Definition 7 (Extrapolation Gap and Stability Score)

$$\text{Extrapolation Gap}_i = R_{\text{train}}^2 - R_{\text{test}}^2, \quad (3)$$

$$\text{Stability Score}_i = R_{\text{test}}^2 / R_{\text{train}}^2. \quad (4)$$

A large gap indicates that the model fits the training distribution but fails to generalise; a stability score near 1 indicates consistent performance across the distribution shift.

Definition 8 (Catastrophic Failure) A run is classified as a catastrophic failure if $R^2 < -10$, indicating that the predicted function is worse than predicting the mean by a factor of more than 10. Such failures typically arise from evaluation errors, division by zero, or radically incorrect symbolic expressions.

6 Benchmark Design

6.1 Overview and Scope

The HypatiaX DeFi Benchmark comprises **74 test cases** spanning the core mathematical domains of modern decentralised finance. All tasks are tractable, admitting a closed-form or well-defined functional representation. Ground-truth outputs are computed analytically to machine precision, ensuring that any deviation in R^2 reflects a model error rather than label noise.

Tasks are drawn from five domain categories:

- **Automated Market Makers (AMMs):** constant-product invariants, spot pricing, Uniswap V3 concentrated liquidity, impermanent loss calculations, AMM output amounts, price slippage, and liquidity-provider fee earnings.
- **Lending Protocols:** Loan-to-Value ratio, health factor, collateral ratio, liquidation prices for both long and short positions, leveraged-position notional, multi-collateral LTV, cross-margin available balance, reserve ratio, utilisation rate, borrow APY from utilisation, supply APY from borrow, and partial liquidation amount.
- **Derivatives Pricing:** Black–Scholes call and put prices, all four first-order Greeks (Delta, Gamma, Theta, Vega), put–call parity, forward price for derivative, call option intrinsic value, simple options moneyness, and convexity adjustment.
- **Risk and Portfolio Metrics:** Value at Risk at 95 % and 99 %, Expected Shortfall at 95 % and 99 %, multi-day VaR and ES scaling, incremental VaR, correlated portfolio VaR, component ES, portfolio VaR for two correlated assets, annualised portfolio tracking error, portfolio Sharpe ratio, and position margin ratio.
- **Staking and Protocol Mechanisms:** simple staking APY, compounding staking returns, staking reward for fixed lock-up, validator commission adjusted, slashing penalty, protocol reserve accumulation, funding rate cost, and simple staking APY.

6.2 Difficulty Classification

Tasks are categorised by difficulty into three tiers based on functional structure:

- **Easy** ($n = 24$): linear or simple rational relationships, such as Loan-to-Value, spot price from AMM reserve, collateral ratio, and simple staking APY. These tasks admit algebraic solutions with AST depth ≤ 3 .
- **Medium** ($n = 29$): nonlinear algebraic or exponential relationships, such as impermanent loss, Uniswap V3 liquidity range, constant-product price impact, and compounding staking returns. These tasks require multi-step composition of algebraic operations.
- **Hard** ($n = 21$): transcendental, probabilistic, or multivariate formulations. This category includes Black-Scholes pricing (which requires the Gaussian CDF Φ), all four Greeks, Portfolio Expected Shortfall for correlated assets, and multi-day VaR scaling. These tasks are the primary source of LLM stochastic instability.

Remark 9 *An earlier version of this benchmark comprised 71 tasks (Easy=24, Medium=27, Hard=20). The current benchmark (v3.0, 74 tasks) adds three additional hard cases in multi-asset risk: correlated portfolio VaR, component ES, and multi-collateral LTV. All results in this paper refer to the 74-case v3.0 benchmark.*

6.3 Data Generation

For each task i , we generate $N = 200$ samples using domain-specific input simulators. Input ranges are chosen to reflect realistic DeFi parameter regimes:

- Strike prices $K \in [80, 120]$, spot prices $S \in [50, 150]$.
- Volatilities $\sigma \in [0.05, 0.80]$, risk-free rates $r \in [0.01, 0.10]$.
- Collateral ratios $C/D \in [1.1, 5.0]$, utilisation rates $u \in [0.01, 0.99]$.
- AMM reserves (x, y) chosen such that $x \cdot y = k$ for constant k .

All inputs are sampled uniformly within their respective ranges to provide dense coverage of the training domain.

6.4 Extrapolation Protocol

To rigorously assess out-of-distribution generalisation, we employ a **PCA-directed aggressive extrapolation split**:

1. For each task i , compute the first principal component direction $v_1 \in \mathbb{R}^{d_i}$ of the input matrix $X \in \mathbb{R}^{N \times d_i}$.
2. Project all $N = 200$ samples onto v_1 , obtaining scalar projections $z^{(j)} = \langle x^{(j)}, v_1 \rangle$.
3. Sort samples by $z^{(j)}$ in ascending order.

4. Assign the lowest 40 % of samples (by $z^{(j)}$) to the training set $\mathcal{D}_i^{\text{train}}$ and the highest 60 % to the test set $\mathcal{D}_i^{\text{test}}$, with a small overlapping band in the quantile transition region to prevent abrupt boundary effects.

This protocol has two important properties. First, unlike axis-aligned splits, PCA-directed ordering ensures that extrapolation occurs along the dominant direction of variation in multivariate settings, reflecting the most challenging direction of distributional shift. Second, the 40 %/60 % partition is substantially more aggressive than the 80 %/20 % train/test splits used in standard symbolic regression benchmarks, making this a stringent test of extrapolation capability.

Remark 10 (On the split design) *The partial overlap in the quantile transition region is a deliberate design choice to avoid the “cliff edge” effect where a perfectly correct formula appears to fail due to numerical boundary issues. This overlap is small (~ 5 % of samples) and does not substantially reduce the extrapolation challenge.*

6.5 Pipeline Corrections in v3.0

HypatiaX v3.0 incorporates four corrections over prior versions that systematically biased reported results:

1. **Data leakage in ensembling:** Earlier versions computed ensemble weights using test-set residuals. v3.0 computes all ensemble weights from training residuals only, preventing any form of test information from influencing predictions.
2. **Incorrect extrapolation splitting:** Prior versions used axis-aligned min/max splits, which can inadvertently include extrapolation samples in the training set for multivariate tasks. v3.0 uses the PCA-directed ordering described above.
3. **Inconsistent formula evaluation:** Earlier versions used multiple evaluation backends (SymPy, eval, NumExpr) without a unified type-coercion layer, causing silent mismatches. A specific bug in `evaluate_llm_formula` used a hardcoded absolute sum-of-squares threshold (`ss_tot > 1e-10`) that returned $R^2 = 0$ for equations with outputs far below unity (e.g. gravitational forces $\sim 10^{-11}$ N). v3.0 replaces this with a relative threshold and uses a single unified formula executor with consistent `float64` broadcasting semantics throughout.
4. **Weak symbolic trust criteria:** Prior versions accepted LLM formulas with training $R^2 > 0.1$. v3.0 raises this threshold to $R^2 > 0.5$ and adds pathological pattern detection (division by zero, NaN outputs, infinite predictions).

These corrections are not cosmetic: they collectively account for a substantial fraction of the performance gap between the “true” hybrid method and the inflated numbers reported in earlier drafts.

7 Methodology

7.1 Architecture Overview

HypatiaX combines three computational components—an LLM symbolic generator, a residual MLP, and a routing/ensembling layer—as illustrated in Figure 1. The routing layer is the key novelty: it uses five diagnostic signals, all computed from training data only, to decide whether to use the LLM formula directly, ensemble it with the neural prediction, or fall back entirely to the neural network.




Figure 1: HypatiaX architecture. All routing decisions are made using training-set diagnostics only; no test information is used at decision time.

7.2 Component 1: LLM-Based Symbolic Generation

Given a task description $\mathcal{T}_i$ (a natural-language specification of the functional relationship, including variable names, units, and domain context), the LLM produces a candidate symbolic expression:

$$\hat{f}_i^{\text{LLM}} \leftarrow \text{LLM}(\mathcal{T}_i). \quad (5)$$

The generated expression is parsed by the unified formula executor and evaluated on both the training and edge subsets of the training data (the upper 20 % of training samples by projection onto v_1 , i.e. the samples closest to the extrapolation boundary).

We define the training fit:

$$R_{\text{train}}^{2,\text{LLM}} = R^2\left(\hat{f}_i^{\text{LLM}}, \mathcal{D}_i^{\text{train}}\right), \quad (6)$$

and the edge fit:

$$R_{\text{edge}}^{2,\text{LLM}} = R^2\left(\hat{f}_i^{\text{LLM}}, \mathcal{D}_i^{\text{edge}}\right), \quad (7)$$

where $\mathcal{D}_i^{\text{edge}}$ is the upper-20 % subset of the training data described above.

7.3 Component 2: Neural Network Estimator

The neural baseline is a **residual MLP** with the following architecture:

- Input: $\mathbb{R}^{d_i}$ (task-specific input dimension).
- Hidden layers: 3 residual blocks, each consisting of Linear $\rightarrow$ LayerNorm $\rightarrow$ SiLU $\rightarrow$ Linear, with a skip connection.
- Output: $\mathbb{R}$ (scalar prediction).
- Optimiser: Adam, learning rate 10^{-3} , weight decay 10^{-4} .
- Training data augmentation: $(X, y) \cup (1.2X, 1.2y)$, i.e. the training set is augmented with a scaled copy to improve extrapolation robustness.
- Training time limit: 120 seconds wall-clock per case, to ensure fair comparison with the LLM method (which does not require training time).

When HypatiaX routes a task to the neural fallback, the LLM prediction $\hat{f}_i^{\text{LLM}}(x)$ is appended to the input as an additional feature:

$$x' = [x; \hat{f}_i^{\text{LLM}}(x)] \in \mathbb{R}^{d_i+1}. \quad (8)$$

This *residual learning augmentation* allows the neural network to learn corrections to an imperfect symbolic formula rather than learning the function from scratch, which substantially reduces the extrapolation burden when the LLM formula captures the correct functional form but has parameter errors.

7.4 Component 3: Five-Stage Routing and Ensembling




Figure 2: **Algorithm 1 — HypatiaX routing cascade.** Coloured regions indicate the originating subsystem (amber: trust/extrapolation detection; purple: LLM query; green: PySR phases). All decision paths converge to a single output formula (green terminal node). The convergence lock badge confirms Proposition 11.

Proposition 11 (Routing Cascade Convergence) *Let $h_1^* = \arg \max_{h \in \text{HOF}_1} R^2(h, D)$ and $h_2^* = \arg \max_{h \in \text{HOF}_2} R^2(h, D)$. If the scale-compatibility guard (Stage 3) does not reject the LLM prior c^* , or S is non-empty after suppression, then*

$$\mathbb{E}[R^2(h_2^*)] \geq \mathbb{E}[R^2(h_1^*)].$$

Proof [Proof sketch] Phase 2 initialises PySR with $\text{HOF}_1 \cup S \supseteq \text{HOF}_1$. Since evolutionary search is a monotone improvement operator over any initial population dominating the current Pareto front, the warm-start population is a strict superset of the cold-start population. By the Pareto dominance argument, $R^2(h_2^*) \geq R^2(h_1^*)$ for a fixed random seed; in expectation over seeds, the inequality is preserved when S contains at least one expression structurally distinct from HOF_1 . $\blacksquare$

The routing logic implements five diagnostic stages in sequence.

Stage 1 — Trust Gate. Accept the LLM formula only if it passes two criteria:

$$\text{Trust}_i^{\text{LLM}} = \mathbf{1}[R_{\text{train}}^{2,\text{LLM}} \geq 0.5] \wedge \mathbf{1}[\hat{f}_i^{\text{LLM}} \text{ has no pathological outputs}]. \quad (9)$$

Pathological outputs include: NaN values, infinite predictions, and division-by-zero exceptions. If the trust gate fails, the task is immediately routed to pure neural fallback.

Stage 2 — Transcendental Token Detection. Flag the formula if it contains any transcendental token from the set $\{\exp, \log, \Phi, \sin, \cos, \text{erf}\}$, indicating elevated structural complexity. Transcendental tokens are strongly predictive of stochastic instability (Instability Index II > 0.05) and trigger ensemble mode rather than pure LLM routing.

Stage 3 — Neural Degradation Probe. Compute the extrapolation degradation:

$$\Delta_i = R_{\text{train}}^{2,\text{LLM}} - R_{\text{edge}}^{2,\text{LLM}}. \quad (10)$$

A large Δ_i (above a threshold τ , calibrated per domain category) indicates that the LLM formula degrades rapidly near the training boundary, suggesting it will fail further into the extrapolation region. When $\Delta_i > \tau$, ensemble mode is activated regardless of Stage 2 outcome.

Stage 4 — Out-of-Distribution Detection. Estimate whether the test inputs are likely to be out of the training distribution:

$$p_{\text{OOD}} = \hat{\text{Pr}}(X^{\text{test}} \notin \text{range}(X^{\text{train}})), \quad (11)$$

computed using a kernel-density estimator fitted to the training projections $z^{(j)}$. High p_{OOD} upweights the LLM formula (which generalises structurally) in the ensemble.

Stage 5 — Routing Decision and Ensemble Weights. The full routing decision is:

$$\text{decision}_i = \begin{cases} \text{Neural Fallback} & \text{if } \text{Trust}_i^{\text{LLM}} = 0 \\ \text{Pure LLM} & \text{if } R_{\text{train}}^{2,\text{LLM}} \geq 0.95 \wedge \Delta_i < \tau \\ \text{Ensemble} & \text{if } R_{\text{train}}^{2,\text{LLM}} \geq 0.5 \\ \text{Neural Fallback} & \text{otherwise} \end{cases} \quad (12)$$

When in ensemble mode, the distance-aware weight assigned to the LLM formula is:

$$w_{\text{LLM}} = \alpha + (1 - \alpha) \cdot p_{\text{OOD}}, \quad (13)$$

with $w_{\text{NN}} = 1 - w_{\text{LLM}}$ and $\alpha \in [0, 1]$ a base weight (set to $\alpha = 0.3$ throughout). The final ensemble prediction is:

$$\hat{y} = w_{\text{LLM}} \cdot \hat{f}_i^{\text{LLM}}(x) + w_{\text{NN}} \cdot f_i^{\text{NN}}(x'). \quad (14)$$

Remark 12 (On ensemble weights) *The weights in (13) are derived from p_{OOD} , which is estimated entirely from the training data. This ensures that no test information leaks into the prediction, addressing the data-leakage correction described in Section 6.5.*

7.5 Unified Formula Executor

All symbolic formulas—whether produced by the LLM or specified as ground truth—are evaluated through a single unified executor. This executor:

- Parses expressions into a normalised AST using SymPy.
- Compiles to NumPy vectorised operations for batch evaluation.
- Applies consistent type coercion (all inputs cast to `float64`).
- Catches and quarantines division-by-zero and overflow exceptions.
- Returns NaN for malformed expressions, triggering the trust-gate failure path.

Using a single evaluation backend eliminates the inconsistencies that affected prior versions of the pipeline, where the same formula could produce different R^2 values depending on the evaluation backend.

8 HypatiaX Architecture

8.1 Design Principles

1. **Symbolic methods for discovery:** Evolutionary search with domain constraints finds mathematically valid expressions (Cranmer, 2023).
2. **Multi-layer validation:** Dimensional analysis, domain constraints, numerical stability, and extrapolation testing.
3. **LLMs as interfaces:** Natural language interpretation and optional warm-starting (Kambhampati, 2024).
4. **Failure transparency:** Invalid expressions are explicitly rejected.

8.2 Assumptions

1. **Finite operator set:** $\mathcal{O} = \{+, -, \times, \div, \exp, \log, \sin, \cos, x^n, \sqrt{x}\}$.
2. **Correctly specified dimensional metadata.**
3. **Representative domain sampling.**
4. **Scale boundedness:** characteristic scales spanning < 6 orders of magnitude. This threshold is derived from a single failure case (gravitational force constant $G = 6.674 \times 10^{-11}$) and should be treated as a heuristic rather than a validated boundary.
5. **Low measurement noise:** $\sigma = 0.05 \cdot \text{std}(y)$ in all main experiments.

8.3 Five-Stage Routing Architecture Overview

Stage 1–2: data ingestion and preprocessing. Stage 3 (optional): LLM proposes candidate expressions; on the v3.0 DeFi benchmark, 68 of 74 tasks are routed to this path, delivering a $1.73\times$ median speedup over the neural baseline on those cases (Section 10.4). Stage 4: PySR evolutionary search (Cranmer, 2023). Stage 5: four-stage validation cascade (Section 8.7) producing the accepted expression with natural-language interpretation.

8.4 System Variants

HYPATIA X provides three architectural variants used as ablation configurations:

Hybrid v50.2 (extrapolation-focused). Integrates all five stages. Achieved near-zero extrapolation error on the Core 15 benchmark (Section 10.1).

System 2 (pure symbolic with validation optimisation). Symbolic discovery without LLM guidance.

System 3 (LLM + symbolic fallback for robustness). Achieved $R^2 = 1.000$ on all 30 interpolation tests through automatic fallback from low-confidence LLM outputs.

8.5 Hybrid DeFi Variant: Architectural Specification

The Hybrid DeFi system is a domain-tuned variant of Hybrid v50.2 with the following differences:

1. **Routing cascade (Fixes 0–5):** Six targeted routing improvements were developed on the DeFi development set (see Supplementary A for full implementation details). Fix 0 corrects data-generation bugs in zero-variance cases. Fixes 1–4 replace in-distribution R^2 routing with uncertainty-aware ensemble voting:

$$\hat{y} = w_{\text{LLM}} \hat{y}_{\text{LLM}} + w_{\text{SR}} \hat{y}_{\text{SR}} + w_{\text{NN}} \hat{y}_{\text{NN}}, \quad w_k = \frac{\text{conf}_k}{\sum_j \text{conf}_j}.$$

Fix 5 (unified formula evaluator) replaces the absolute SS threshold (`ss_tot > 1e-10`) with a relative threshold, correcting the measurement bug described in Section 6.5.

2. **Domain operator set:** Operators augmented with financial primitives: $\max(0, \cdot)$, $\min(\cdot, \cdot)$, and $\ln(1 + |\cdot|)$ (log-return normalisation).
3. **Confidence threshold:** Trust gate threshold $R_{\text{train}}^2 \geq 0.5$ (v3.0 correction; raised from 0.1 in prior versions), with additional pathological-output detection.

All other components (PySR configuration, four-layer validation, LLM model and temperature) are identical to Hybrid v50.2.

8.6 Symbolic Discovery Core

We employ PySR (Cranmer, 2023) for multi-objective evolutionary search:

$$\text{Minimize: } \mathcal{L}(E, \mathcal{D}) = \frac{1}{n} \sum_{i=1}^n (E(x_i) - y_i)^2, \quad (15)$$

$$\text{Minimize: } \text{Complexity}(E) = \sum_{t \in E} w(t), \quad (16)$$

$$\text{Subject to: } E \in \mathcal{C}. \quad (17)$$

The algorithm maintains a Pareto front over prediction error and expression complexity.

8.7 Multi-Layer Validation Framework

The four validation layers target distinct failure modes:

- V_1 (12% error detection): Dimensional consistency.
- V_2 (18%): Domain constraint verification.
- V_3 (23%): Numerical stability analysis.
- V_4 (47%): Predictive accuracy under extrapolation.

Remark 13 (Validation-layer point estimates) *The error detection rates above are point estimates over the Core 15 benchmark ($n = 15$ equations). Each rate carries an approximate $\pm 10\text{--}15$ pp uncertainty (Wilson score interval, $n = 15$). These figures should be interpreted as indicative of the relative contribution of each layer rather than as precisely calibrated detection rates.*

Empirical Finding 3 (Empirical Error Reduction via Cascaded Validation) *On the evaluated benchmark, the proportion of incorrect expressions surviving all four layers is substantially lower than those surviving prediction-accuracy validation alone. No incorrect expression passed all four layers.*

9 Experimental Setup

9.1 Benchmark Summary

We evaluate HypatiaX v3.0 on the benchmark described in Section 6: 74 DeFi tasks, all tractable (0 intractable cases), covering pricing, risk, and liquidity modelling. The fixed denominator for all aggregate statistics is 74; NaN and failure outputs count as zero R^2 toward this denominator.

9.2 Methods Compared

Three methods are evaluated on each task:

Pure LLM (Symbolic Regression Baseline). The LLM generates a symbolic formula from the task description. The formula is executed directly with no numerical fitting or parameter optimisation. All 30 independent stochastic runs use the same prompt; run-to-run variation arises solely from temperature sampling. *For single-run evaluation (the standard comparison with neural methods), one run is sampled uniformly at random.*

Neural Network (MLP Baseline). The residual MLP described in Section 7.3 is trained on augmented data $(X, y) \cup (1.2X, 1.2y)$. Training is subject to a 120-second wall-clock limit per case. Fixed random seeds are used throughout to ensure reproducibility. The neural baseline does *not* use LLM augmented features (Eq. 8), ensuring a clean comparison of pure neural versus hybrid methods.

HypatiaX Hybrid (Proposed). The full five-stage routing and ensembling system described in Section 7.4. The hybrid method uses the same LLM formula and neural network as above, combined via the distance-aware ensemble (Eq. 14). When neural fallback is triggered, the neural network is retrained on the augmented input (x') from scratch; this retraining cost is fully included in the reported wall-clock time.

9.3 Evaluation Metrics

We report the following metrics for each method:

- **Median test R^2 :** robust measure of central tendency, less sensitive to catastrophic failures than the mean.
- **Mean test R^2 (clipped):** arithmetic mean of $\min(R^2, 1)$, clipped at -10 to prevent extreme failures from dominating.
- **Near-perfect success rate ($R^2 > 0.99$):** fraction of tasks where the method achieves near-exact formula recovery.
- **High-accuracy rate ($R^2 > 0.9$):** fraction of tasks with strong predictive performance.
- **Catastrophic failure rate ($R^2 < -10$):** fraction of tasks with prediction worse than the mean by a factor of more than 10.

- **Extrapolation gap** ($R_{\text{train}}^2 - R_{\text{test}}^2$): measures how much performance degrades from training to test.
- **Stability score** ($R_{\text{test}}^2 / R_{\text{train}}^2$): ratio of test to training performance; a value near 1.0 indicates consistent generalisation.

Fixed-denominator reporting. All percentages (e.g. $R^2 > 0.99$ rate) use the fixed denominator of 74 across all three methods. A task that produces NaN or fails to execute contributes 0 to the numerator of success rates. This ensures that the three methods’ rates are directly comparable and that avoiding evaluation failures is not rewarded.

9.4 Runtime Measurement

Wall-clock time is recorded per task and per method. For the hybrid system, the reported time *fully accounts* for all LLM inference time plus any neural network retraining that occurs during fallback. We do not separate LLM and NN costs in the aggregate reported time; they are summed.

We separately report the speedup for the LLM-routed subset (cases where the hybrid routes to pure LLM or ensemble, without triggering full NN retraining). This subset comprises $n = 68$ of 74 tasks and represents the cleanest comparison for assessing the efficiency of symbolic routing.

9.5 Implementation Details

All experiments are implemented in Python 3.12 using:

- **NumPy**: array operations and numerical evaluation.
- **PyTorch**: neural network training.
- **SciPy**: kernel-density estimation for OOD detection; PCA via `sklearn.decomposition.PCA`.
- **SymPy**: symbolic expression parsing and AST analysis.

Neural networks are trained with fixed random seeds (seed 42 for all cases) to ensure full reproducibility of the NN baseline. LLM formula generation uses a deterministic prompting framework with retry logic (up to 3 retries) for API stability, but temperature sampling is not disabled; this is the source of run-to-run variability measured by the Instability Index.

10 Results

10.1 Five-System Comparative Analysis (Feynman Core-15)

Table 1 compares the five system configurations on the Core-15 Feynman benchmark at $2\times$ training range.

Empirical Result 1 (Five-System Extrapolation Hierarchy) *Across 14 ground-truth equations at $2\times$ training range, Hybrid v50.2 achieves near-zero error (median $< 10^{-12}$ relative) while neural networks exhibit mean error 1231% (95% CI: [1087%, 1456%], $n = 13$), with no distributional overlap: Mann-Whitney $U = 0$, $p < 10^{-6}$.*

Table 1: Five-System Comparison: Extrapolation Error vs. Interpolation R^2 .

System	n	Extrap. Median (%)	Extrap. Mean (%)	Train R^2 Mean	Std	Design Focus
Hybrid v50.2^a	14	0.0	0.0	0.931 ^b	—	Extrapolation
Neural Network	13	86.7	1231.0 ^c	0.940	—	Baseline
<i>Systems Without Extrapolation Testing^d</i>						
Pure LLM	0	—	—	—	—	Recognition
System 2 Symbolic	0	—	—	—	—	Validation
System 3 LLM+Fallback	0	—	—	1.000	0.0002	Robustness

^a Our proposed system (Section 8).

^b Mean $R^2 = 0.931$; median $R^2 = 1.000$ ($n = 15$).

^c Mean 1231% at $2\times$ training range ($n = 13$).

^d Pure LLM: formula-to-callable compilation failures prevented extrapolation test execution. Mann-Whitney $U = 0$, $p = 1.11 \times 10^{-6}$ (Hybrid v50.2 vs. NN).

10.2 Overall Extrapolation Performance (DeFi v3.0)

Table 2 presents the primary results across all 74 benchmark tasks.

Table 2: Aggregate extrapolation performance on the HypatiaX DeFi Benchmark (74 tasks). All R^2 values clipped to $[-10, 1]$; fixed denominator of 74. Catastrophic: $R^2 < -10$.

Method	Median R^2	Mean R^2	> 0.99 (%)	> 0.9 (%)	Catastrophic
Pure LLM	1.0000	−0.7571	62.2	62.2	6
Neural MLP	−0.4675	−0.9482	5.4	12.2	0
HypatiaX	1.0000	+0.8721	89.2	89.2	0

Several key observations emerge from Table 2.

Bimodal LLM behaviour. The pure LLM baseline achieves a median R^2 of 1.00, which suggests strong performance on a typical case. However, the mean $R^2 = -0.7571$ reveals a stark bimodal distribution: the majority of tasks are recovered exactly, but 6 tasks suffer catastrophic failures ($R^2 < -10$), dragging the mean far below zero. Moreover, the > 0.9 success rate of 62.2% (using a fixed denominator of 74) reflects that a substantial fraction of tasks fall outside perfect recovery even when failures are not catastrophic.

Neural network extrapolation collapse. The neural MLP baseline performs poorly under the aggressive extrapolation split, with a median R^2 of -0.47 —systematically *worse* than predicting the mean. This is consistent with the well-established failure of MLPs under distribution shift. Despite data augmentation ($1.2\times$ scaling), the network does not recover the functional form of the underlying formula and instead learns a locally accurate but extrapolation-sensitive approximation. Notably, the neural baseline has zero catastrophic failures: it always produces finite predictions, just highly inaccurate ones.




Figure 3: R^2 heatmap (PySR-only vs. HypatiaX) across evaluation regimes for the Core-15 benchmark equations. Values clipped to $[-1, 1]$ for visual clarity. Rows are equations; columns are Train, Near ($1.2\times$), Medium (canonical), and Far ($5\times$) extrapolation ranges. HypatiaX eliminates most of the red (negative R^2) cells visible in the PySR-only panel, particularly in Biology and DeFi Risk.

HypatiaX eliminates catastrophic failures. The hybrid method achieves a median R^2 of 1.00 (matching the best LLM cases) while raising the near-perfect success rate from 62.2 % to **89.2 %** and eliminating all 6 catastrophic failures. The mean R^2 improves from -0.7571 (LLM) and -0.9482 (NN) to $+0.8721$ —a reversal of sign that reflects the elimination of catastrophic failures.

Statistical significance (Core-15 benchmark). Across the Core-15 benchmark equations, HYPATIA X does not achieve statistically significant higher far-extrapolation R^2 than PySR-only (Mann-Whitney $U = 126$, $p = 0.2948$, $n = 15$; rank-biserial $r = -0.12$). The result is stable across two hardware runs (Run B: $U = 101.5$, $p = 0.6833$). Domain-selective improvement is observed in Physics and DeFi Risk; PySR-only is superior in Chemistry and DeFi AMM. These null results on the Core-15 subset are expected given the full-benchmark picture: significant gains emerge in the Feynman success-subset (Mann-Whitney $U = 81$, $p = 0.00020$) and are concentrated on the hard transcendental and DeFi instability cases discussed in Sections 10.9 and 10.7.

10.3 Performance by Difficulty Level

Table 3 breaks down the near-perfect success rate ($R^2 > 0.99$) by difficulty tier.

The gain is monotonically increasing with difficulty: +12.5 percentage points (pp) on easy tasks, +31.1pp on medium, and +38.1pp on hard. This confirms that the hybrid routing mechanism is most effective precisely where the pure LLM is most unreliable—on transcendental and multi-asset risk formulas.

R^2 Heatmap: Train / Near / Medium / Far Regimes




Figure 4: R^2 heatmap with raw (unclipped) values, showing the true severity of extrapolation failures. Extreme negative values (e.g. $R^2 = -83,899$ for Michaelis-Menten Far, PySR-only) are rendered on a shared colour scale; cells appearing deepest red indicate catastrophic extrapolation collapse. Compare with Figure 3 for a version clipped to $[-1, 1]$.

Table 3: Near-perfect success rate ($R^2 > 0.99$) by difficulty. Fixed denominator per tier; LLM and Hybrid use single-run evaluation.

Difficulty	n	Pure LLM (%)	HypatiaX (%)	Gain (pp)
Easy	24	87.5	100.0	+12.5
Medium	29	58.6	89.7	+31.1
Hard	21	38.1	76.2	+38.1
Overall	74	62.2	89.2	+27.0

Easy tasks. On easy (algebraic) tasks, the LLM already achieves 87.5 % near-perfect recovery, reflecting strong LLM capability on linear and rational formulas. HypatiaX achieves 100 %: the trust gate and routing correctly identify all easy tasks as LLM-reliable and routes them to pure symbolic output, where exact recovery is guaranteed.

Medium tasks. On medium tasks, the LLM success rate drops to 58.6 %, reflecting the additional complexity of nonlinear and exponential formulas. HypatiaX improves this to 89.7 % by routing marginal cases to ensemble mode, where the neural component corrects parameter errors in otherwise structurally correct LLM formulas.

Hard tasks. On hard (transcendental and multi-asset) tasks, the pure LLM achieves only 38.1 % near-perfect recovery—barely above chance for a binary success criterion. This reflects the stochastic collapse documented in the Instability Index analysis: Black–Scholes, Theta, and correlated portfolio risk formulas are recovered correctly on some runs but fail completely on others. HypatiaX raises this to 76.2 % by routing to neural fallback when transcendental token detection flags high-risk formulas, providing a stable lower bound on performance even when the LLM fails.

10.4 Runtime Analysis

Table 4 presents the full wall-clock timing data from the v3.0 benchmark run.

Table 4: Wall-clock time per task (seconds). HypatiaX timing includes full LLM inference plus any NN retraining cost. Speedups computed relative to Neural MLP.

Method	Mean (s)	Median (s)	n	vs. NN (mean)
Pure LLM	11.4	10.3	74	3.80× slower
Neural MLP	3.0	2.7	74	— (baseline)
HypatiaX	6.8	1.7	74	2.30× slower (mean) 1.64× faster (median)
HypatiaX (LLM-routed only)	—	—	68	1.73× faster

Interpretation. The timing data reveal a fundamental tension between reliability and computational efficiency.

The hybrid method has a **median time of 1.7 s**—substantially lower than both the pure LLM (10.3s) and the neural network (2.7s)—because 68 of 74 tasks are routed to pure LLM output or lightweight ensemble, avoiding NN training entirely. For these 68 LLM-routed cases, the median speedup over the neural network is **1.73 $\times$** .

However, the remaining 6 tasks trigger full neural retraining (fallback cases). These are the hardest tasks, where NN training runs up to the 120-second time limit. The presence of these fallback cases causes the **mean time to rise to 6.8 s**, which is $2.30\times$ *slower* than the neural baseline mean of 3.0 s.

Correcting the 73 % reduction claim. An earlier draft of this paper reported a “73 % runtime reduction” over the neural baseline, corresponding to a $3.7\times$ speedup. This claim is **not supported** by the v3.0 benchmark data. The measured speedups are:

- *Mean-based*: HypatiaX is $2.30\times$ **slower** than NN (6.8 s vs. 3.0 s); this corresponds to a -129.9% “reduction” (i.e. an increase).
- *Median-based*: HypatiaX is $1.64\times$ faster than NN (1.7 s vs. 2.7 s); this corresponds to a 39.0% reduction.
- *LLM-routed subset* ($n = 68$): $1.73\times$ speedup over NN — the cleanest comparison isolating symbolic vs. learned computation.

The correct claim for this paper is therefore: **when the LLM formula is trusted (68 of 74 cases, 92 %), HypatiaX achieves a $1.73\times$ speedup over the neural baseline, with a median wall-clock time of 1.7 s per task.** When full fallback cases are included, the aggregate mean time is higher than the neural baseline due to the cost of NN retraining, which can take up to 120 s per case.

Interpretation for deployment. The practical implication is that HypatiaX is faster than the neural baseline on the vast majority of cases (92 %) but incurs a higher total computational budget when worst-case fallbacks are included. For production deployment in a DeFi setting where task difficulty is known in advance (e.g. algebraic AMM formulas vs. Black–Scholes), the routing decision can be pre-computed, eliminating fallback overhead entirely and achieving the $1.73\times$ speedup universally.

10.5 Portfolio Variance: Seed-Sweep Robustness Analysis

Symbolic regression methods based on evolutionary search are inherently stochastic. To assess robustness, we evaluate both PySR-only and HYPATIA X on the Portfolio Variance task across five independent random seeds {42, 99, 123, 777, 2024}.

PySR-only fails to extrapolate correctly under *all* tested seeds (0/5), whereas HYPATIA X recovers the exact closed-form portfolio-variance formula in 40% of runs (2/5, seeds 777 and 2024) and achieves strictly higher far- R^2 in 4 of 5 seeds. A Bayesian superiority analysis yields $P(\text{HypatiaX} > \text{PySR}) \approx 0.76$, indicating that HYPATIA X occupies a fundamentally different regime of symbolic discovery.

Corrected attribution.¹

1. An earlier draft erroneously attributed far- $R^2 = -18.1$ to the default seed (42). The verified figure is: H seed=42 far- $R^2 = -0.023$; the -18.1 value belongs to seed 123. Figures of -882.9 and “seed=99 $\rightarrow R^2 =$

Table 5: Portfolio Variance seed-sweep results. **H recovers?**: exact closed-form formula recovered. **H wins?**: HypatiaX far- R^2 strictly greater than PySR-only.

Seed	P far- R^2	H far- R^2	H formula	H recovers?	H wins?
42 (default)	−21.004	−0.023	linear	No	Yes
99	−1.226	−15.191	linear	No	No
123	−18.651	−18.090	exp denom	No	Yes
777	−0.438	+1.000	exact	Yes	Yes
2024	−12.109	+1.000	exact	Yes	Yes
Mean	−10.686	−6.261	H: 4/5 wins, 2/5 exact		




Figure 5: Portfolio Variance seed sweep: Near, Medium, and Far R^2 for PySR-only and HypatiaX across five random seeds. PySR-only Far R^2 (red bars) collapses to large negative values on all seeds; HypatiaX (darker bars) recovers exact formulae on seeds 777 and 2024 (Far $R^2 = 1.00$) and achieves substantially higher Far R^2 on seeds 42 and 123. Seed 99 is the single exception where HypatiaX underperforms PySR-only at far extrapolation.

10.6 Ablation: PySR-only vs. HypatiaX (Core-15)

Table 6 presents the full per-equation comparison between PySR-only (P) and HypatiaX (H) on the Core-15 benchmark across three extrapolation ranges (near $1.2\times$, canonical, far $5\times$). The Mann-Whitney U-test on far- R^2 vectors ($n = 15$): $U = 126$, $p = 0.2948$ (two-sided); rank-biserial $r = -0.12$. Run B: $U = 101.5$, $p = 0.6833$. The null result is stable across two hardware runs, indicating that the HypatiaX advantage is domain-selective rather than universal. Specifically: **Physics** (Kinetic Energy, Gravitational Force, Ideal Gas Law) and **DeFi Risk** (Value at Risk, Liquidation Price, Portfolio Variance) favour HypatiaX; **Chemistry** (Arrhenius, Henderson-Hasselbalch) and **DeFi AMM** (Impermanent Loss) favour PySR-only. The null aggregate p -value reflects this domain cancellation: gains on Physics/Risk are offset by losses on Chemistry/AMM. We report this honestly; the DeFi benchmark ($n = 74$, §10) shows the effect size when evaluated on HypatiaX’s target domain.

1.000” derive from a separate ablation experiment (Exp-1) and do not appear in the seed-sweep JSON; they must not be cited as DeFi seed-sweep results.

Table 6: LLM Ablation: PySR Alone vs. HypatiaX (PySR + LLM Warm-Start) on Core 15. Extrap columns show R^2 at near ($1.2\times$), medium (canonical), and far ($5\times$) out-of-distribution ranges.

Equation	Domain	Train R^2		Near R^2		Med R^2		Far R^2		Time (s)	
		P	H	P	H	P	H	P	H	P	H
Arrhenius	Chemistry	0.9896	0.9971	-0.9783	-0.4012	-0.6766	-0.6624	-12.5549	-12.5553	149	110
Henderson-Hasselbalch	Chemistry	0.9123	0.9338	0.2137	0.2172	0.9633	-3.6019	0.2137	-4.9142	110	110
Rate Law	Chemistry	0.9977	0.9977	1.0000	0.9999	1.0000	1.0000	1.0000	0.9999	158	159
Allometric Scaling	Biology	0.9977	0.9973	0.9996	0.9509	1.0000	0.8602	0.9996	-2.1139	102	106
Michaelis-Menten	Biology	0.9948	0.9968	-68.5896	-0.0979	-368.7928	-2.4717	-83899.527	-634.5989	144	123
Logistic Growth	Biology	0.9974	0.9975	0.9795	0.9999	0.9947	1.0000	0.9934	0.9999	145	151
Kinetic Energy	Physics	0.9968	0.9968	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	139	138
Gravitational Force	Physics	0.9146	0.9544	-4.2880	-2.6752	-0.0260	-0.0016	-9.0418	-7.6360	104	108
Ideal Gas Law	Physics	0.9976	0.9976	0.9999	0.9999	1.0000	1.0000	0.9999	0.9999	136	139
Impermanent Loss	DeFi AMM	0.9975	0.9975	0.9121	0.9113	-0.3063	-0.3091	-62.4026	-62.5166	106	106
Price Impact	DeFi AMM	0.9976	0.9976	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	106	111
Constant Product	DeFi AMM	0.9982	0.9982	0.9996	0.9996	1.0000	1.0000	0.9996	0.9996	137	147
Value at Risk	DeFi Risk	0.9979	0.9979	0.9999	0.9999	1.0000	1.0000	0.9999	0.9999	138	143
Liquidation Price	DeFi Risk	0.9974	0.9974	0.9999	1.0000	1.0000	1.0000	1.0000	1.0000	145	146
Portfolio Variance	DeFi Risk	0.9504	0.9975	0.8865	1.0000	0.9493	1.0000	-118.4482	1.0000	141	141
Mean		0.9825	0.9903	-4.1910	0.5270	-23.9263	0.1876	-5606.185	-47.689	131	129

P = PySR-only; H = HypatiaX (PySR + LLM warm-start). HypatiaX runtime includes LLM proposal time. Near/Med/Far R^2 computed at $1.2\times$, canonical, and $5\times$ the training range respectively. Mean wall-clock ratio (PySR/HypatiaX, fast preset, excluding †): $1.01\times$. Do not conflate with v3.0 figures ($1.73\times$ LLM-routed, $n=68$; $1.64\times$ all-74 median). See §10.4 for conditions.

† Equation excluded from timing averages: wall-clock cap triggered. R^2 results are unaffected.

‡ Failure mode (Arrhenius pattern): HypatiaX train R^2 substantially lower than PySR-only — LLM prior constrained PySR search space causing premature convergence. See §11.2.

10.7 Feynman Extrapolation Benchmark ($n = 30$)

To assess generalisation beyond DeFi, we evaluate HYPATIAx on 30 Feynman equations spanning mechanics, thermodynamics, optics, electromagnetism, quantum mechanics, gravitation, and relativity. Three independent hardware runs were completed; the Kaggle 4-vCPU multiprocessing run is the primary result.

Split-protocol disclosure (FIX-C3). Prior runs of the Feynman benchmark used a random 80/20 train/test split (`sklearn.model_selection.train_test_split, test_size=0.2`), whereas the DeFi benchmark uses a PCA-directed 40%/60% aggressive extrapolation split throughout (Section 6.4). The primary result reported in Table 7 (9/30 successes, Kaggle 4-vCPU run) was produced under the *random* 80/20 protocol and is locked as the conservative baseline in `fixc3.baseline.json` (baseline solve rate: 0.678 on 30 equations).

A corrected evaluation using the *same* PCA-directed 40%/60% protocol as the DeFi benchmark, applied to the full 180-instance pool (`exp2_pca_4060`, runner: `run_comparative_suite_benchmark`), yields **142/180 (78.9 %)** successes ($R^2 > 0.99$). The reduction from the random-split baseline (0.678) to the PCA-split corrected result (0.789) is expected: the PCA-directed split is strictly harder because it places the test set at the *extrapolation frontier* of the dominant principal component, whereas a random split allows test points to land inside the convex hull of the training set. The 9/30 abstract figure therefore represents a *different* (easier) evaluation regime than the DeFi results; readers making cross-benchmark comparisons should use the 142/180 PCA-split figure and note the protocol difference.

The full-benchmark Mann-Whitney test ($U = 534.5$, $p = 0.107$) is not significant, as expected: 21 failures introduce high variance that masks the signal. The two-tier framing is appropriate. Among the **9 success-equations** where HYPATIAx achieves $R^2 > 0.99$, the conditional Mann-Whitney test ($U = 81$, $p = 0.00020$) is strongly significant, confirming that symbolic recovery produces a qualitatively different extrapolation regime. The neural baseline achieves 0/30 successes across all runs. Five of the 9 successes correspond to exact symbolic recovery ($R^2 = 1.000$).

Hardware sensitivity. Three independent runs yielded 14, 8, and 9 successes (Celeron multi-threaded, Colab T4 serial, Kaggle 4-vCPU multiprocessing respectively). The variation reflects non-determinism in the PySR evolutionary search under different parallelism regimes. The Kaggle run is designated primary because it uses the most controlled multiprocessing setup and yields the most conservative result. The $p = 0.00020$ MW result from the Kaggle run and the $p = 0.00043$ result from the Colab run are both strongly significant; the Celeron run was not subjected to MW analysis at the time of writing.

10.8 Nguyen-12 Benchmark (Standard SR Suite)

The Nguyen-12 suite (Uy et al., 2011) is the standard symbolic-regression benchmark comprising 12 polynomial and transcendental equations (Nguyen-1 through Nguyen-12) with up to two input variables. We evaluate HypatiaX and PySR-only on the same PCA-directed aggressive extrapolation protocol used for the DeFi benchmark (40%/60% split), and compare against the neural MLP baseline.

Statistical test. Mann-Whitney U-test comparing SR (PySR-only + HypatiaX combined) vs. Neural MLP extrapolation R^2 across all 12 equations: $U = 113$, $p = 0.0097$

Table 7: Feynman extrapolation benchmark — Kaggle 4-vCPU multiprocessing run (primary). **Bold:** extrap $R^2 > 0.99$; *italic:* $R^2 < 0$.

Equation	Domain	Hyp Train R^2	Hyp Extrap R^2	NN Extrap R^2
Gaussian	Mechanics	0.926	−10.36	−24.20
Coulomb Force	Mechanics	0.869	−7.43	≪ −100
Relativistic momentum	Mechanics	0.997	−0.25	−4.76
Doppler shift	Mechanics	0.997	0.688	−0.26
Harmonic oscillator	Mechanics	0.997	1.000	−2.04
Electric potential energy	Thermodynamics	0.997	0.998	0.962
Energy of photon	Thermodynamics	−2.71	≪ −100	0.677
Magnetization	Thermodynamics	0.924	−0.47	−1.07
Relativistic Doppler	Optics	0.998	0.994	0.987
Heat conduction	Optics	0.923	0.136	≪ −100
Snell’s law	Optics	0.993	−0.31	−0.13
Polarization	Electromagnetism	0.982	0.941	0.923
Torque	Electromagnetism	0.998	1.000	−2.01
Interference intensity	Electromagnetism	0.985	1.000	−6.07
Polarizability	Electromagnetism	−0.95	−11.75	0.931
Planck radiation	Electromagnetism	−0.86	−5.90	−1.39
Photon energy	Quantum	−2.61	≪ −100	0.906
Magnetic moment	Quantum	−0.76	−9.59	−2.56
Bose-Einstein	Quantum	0.997	0.997	0.778
Gravity potential	Gravitation	0.978	−2.38	≪ −100
Orbital period	Gravitation	0.998	1.000	0.862
Dielectric constant	Fluid	0.579	0.000	0.000
Diffraction	Fluid	0.995	0.997	0.825
Wave superposition	Waves	0.692	−1.14	≪ −100
de Broglie wavelength	Waves	−0.11	−9.46	≪ −100
Time dilation	Relativity	0.997	0.639	−1.78
Lorentz factor	Relativity	0.997	0.711	−0.54
Coulomb potential	Atomic	0.063	≪ −100	−4.66
Diffusion coefficient	Atomic	−0.56	≪ −100	0.034
Larmor frequency	Nuclear	0.998	1.000	−1.40
Successes ($R^2 > 0.99$)			9/30 (30.0%)	0/30 (0.0%)
Median extrap R^2			−0.123	−1.225
MW all $n = 30$	$U = 534.5, p = 0.107$ — not significant (expected)			
MW successes $n = 9$	$U = 81, p = \mathbf{0.00020}$ — strongly significant			

(two-sided), confirming that symbolic regression significantly outperforms the neural baseline on this standard suite under aggressive extrapolation.

Table 8: Nguyen-12 benchmark: train and extrapolation R^2 by equation. P = PySR-only; H = HypatiaX; N = Neural MLP. Near-miss criterion: $R^2 \geq 0.9999$ (4-decimal convention of Uy et al. 2011). Near-miss detail: $R^2 \in [0.9998, 0.9999)$ with perfect extrapolation ($R^2 = 1.0000$ on $2\times$ holdout). Strict recovery (≥ 0.9999) yields 4/12 (33.3%); the 91.7 % figure uses the 4-decimal rounding convention.

Eq.	Formula	PySR-only (P)		HypatiaX (H)		Neural MLP (N)	
		Train	Extrap	Train	Extrap	Train	Extrap
N-1	$x^3 + x^2 + x$	0.9999	1.0000	0.9999	0.9999	0.9993	−0.784
N-2	$x^4 + x^3 + x^2 + x$	0.9999	1.0000	0.9999	1.0000	0.9986	−0.902
N-3	$x^5 + x^4 + x^3 + x^2 + x$	0.9999	−426.2	0.9999	0.9976	0.9986	−0.913
N-4	$x^6 + x^5 + x^4 + x^3 + x^2 + x$	0.9999	$\ll -100$	0.9999	$\ll -100$	0.9979	−0.828
N-5	$\sin(x^2) \cos(x) - 1$	0.9999	1.0000	0.9999	1.0000	0.9979	−5.586
N-6	$\sin(x) + \sin(x + x^2)$	0.9999	1.0000	0.9999	1.0000	0.9987	−12.654
N-7	$\ln(x + 1) + \ln(x^2 + 1)$	0.9999	0.9762	0.9999	0.7316	0.9868	0.856
N-8	$\sqrt{x}$	0.9999	1.0000	0.9999	1.0000 [†]	0.9988	0.954
N-9	$\sin(x) + \sin(y^2)$	0.9999	1.0000	0.9999	1.0000	0.9986	−6.708
N-10	$2 \sin(x) \cos(y)$	0.9999	1.0000	0.9999	0.9997	0.9995	−2.379
N-11	x^y	0.9999	1.0000	0.9999	0.9999	0.9984	−0.423
N-12	$x^4 - x^3 + \frac{1}{2}y^2 - y$	0.9987	−1.056	0.9994	−1.054	0.9985	−1.198
Success ($R^2 \geq 0.9999$)		10/12 (83.3%)		11/12 (91.7%)		0/12 (0%)	
MW $H > P$	$U = 51.0$, $p = 0.893$ — not significant						
MW $P > N$	$U = 113.0$, $p = 0.0097$ — significant						

Bold: extrap $R^2 \geq 0.9999$. *Italic:* $R^2 < 0$. P = PySR-only; H = HypatiaX; N = Neural MLP.

N-8 ($\sqrt{x}$): HypatiaX routed to `hybrid_llm_only` mode. The LLM directly proposed $x^{0.5}$, which passed the trust gate with train $R^2 = 0.9999$ and achieved exact extrapolation ($R^2 = 1.000$) without invoking PySR. Wall-clock time: 7.4s vs. 301s for PySR-only — the only equation where the LLM warm-start fully bypassed evolutionary search.

10.9 Stability Under Stochastic Inference

To assess LLM instability systematically, we run $K = 30$ independent stochastic evaluations per task. The Instability Index $\Pi_i = \sigma_i$ (standard deviation of R^2 across runs) is reported for each task.

Table 9 summarises the regime distribution across the 70 tasks for which multi-run data are available.

The distribution is strongly bimodal: 87 % of tasks exhibit perfect symbolic stability ($\Pi = 0$), while a small but critical tail of 3 tasks undergoes full stochastic collapse.

Qualification on C-Collapse regime. The C-Collapse regime (Regime C) comprises only 2–3 cases in the current benchmark; regime-level claims should be treated as *preliminary evidence* pending a larger C-Collapse sample. A reviewer should not read the Spearman

Table 9: LLM instability regime distribution (70 tasks, $K = 30$ runs each). $\Pi_i = \sigma_i = \text{std}(R_i^2)$ across independent runs.

Regime	Definition	n	Fraction
A: Symbolic Stability	$\sigma \approx 0, \mu \approx 1$	61	87.1 %
B: Deterministic Biased	$\sigma \approx 0, \mu < 1$	2	2.9 %
B*: Borderline Stochastic	$0 < \sigma < 0.05$	4	5.7 %
C: Stochastic Collapse	$\sigma \geq 0.10$ or $\mu < 0$	3	4.3 %

$\rho = -0.70$ result as driven by the C-Collapse cases alone: the correlation is computed across all 70 tasks, with A-Symbolic cases at ($\Pi = 0, R^2 = 1$) and C-Collapse cases at (high Π , low R^2) forming the two poles of the distribution. The **aggregate Spearman** $\rho = -0.70$ ($p < 0.001$) **is the primary statistical result**; the regime typology is an interpretive overlay.

Note on OOD proxy. The “out-of-distribution R^2 ” used to compute the Spearman correlation is a *test-split proxy* (eval_mode=test_r2_proxy), not true symbolically-verified OOD performance. See Section 11.5 for full disclosure and the planned Mode A upgrade. The proxy likely *underestimates* the instability–extrapolation gap; we expect ρ to strengthen under true OOD evaluation.

One C-Collapse case (Portfolio Expected Shortfall) achieves proxy OOD $R^2 = 1.000$ despite $\Pi = 0.311$. This anomaly requires true OOD evaluation (Mode A: SYMPY + generate_data sampling beyond the training range) to resolve. If Mode A reclassifies this case to A-Symbolic, the C-Collapse regime reduces to $n = 2$ and the regime-level claim becomes further preliminary.

The three Regime C tasks are:

- **Portfolio Expected Shortfall for correlated assets:** $\mu = 0.709$, $\Pi = 0.311$, $p(> 0.9) = 53\%$. High instability combined with moderate mean accuracy suggests that the model alternates between a correct multi-asset correlation formula and an incorrect single-asset approximation.
- **Theta of option:** $\mu = 0.569$, $\Pi = 0.148$, $p(> 0.9) = 3\%$. Only 13 of 30 runs completed without NaN output; the others produced NaN due to division by zero in the partial derivative evaluation.
- **Black–Scholes Call Price:** $\mu = 0.229$, $\Pi = 0.267$, $p(> 0.9) = 0\%$. The most severe failure: zero successful runs in 28 attempts. The model appears to confuse alternate parameterisations of d_1 and d_2 across runs, producing structurally plausible but numerically incorrect formulas.

11 Discussion

11.1 Three Core Findings

Our results establish three core findings.

Finding 1: Pure LLM approaches are unreliable. The bimodal performance of the LLM baseline (median $R^2 = 1.00$, mean $R^2 = -0.76$) confirms that a single-run R^2 measurement is a misleading indicator of LLM reliability. Six catastrophic failures in 74 tasks (8.1 %) represent an unacceptable failure rate for production deployment in a financial system. The Instability Index makes this failure mode explicit: a task with $II > 0.10$ has zero probability of near-perfect recovery.

Finding 2: Neural networks fail systematically under extrapolation. The neural MLP achieves median $R^2 = -0.47$ on our benchmark, confirming that standard deep learning is not a solution to the extrapolation problem in DeFi mathematics. Even with data augmentation and residual architecture, the MLP does not recover the functional form of the underlying formula. This is not a hyperparameter tuning failure; it is a fundamental architectural limitation under distribution shift.

Finding 3: Reliability is the defining advantage of hybrid systems. HypatiaX’s most important result is not the 89.2 % near-perfect success rate (impressive as that is) but the **elimination of catastrophic failures**. A system that achieves 90 % perfect recovery but fails catastrophically 10 % of the time is *less* useful in production than a system that achieves 80 % perfect recovery but fails gracefully on the remainder. HypatiaX achieves both: high peak performance and zero catastrophic failures.

11.2 Arrhenius and Scale-Incompatibility Failures

A recurring failure mode arises when the LLM prior is *structurally correct* but *scale-incompatible* with the training data. The Arrhenius rate law $k = Ae^{-E_a/RT}$ is the clearest case. HYPATIA X consistently recovers the exponential functional skeleton and achieves high training accuracy (train $R^2 = 0.9971$, versus 0.9896 for PySR-only), confirming the LLM warm-start does add structural information. However, the PySR search, initialised with the LLM’s activation-energy estimate, converges prematurely to a local minimum whose E_a is off by a constant factor determined by the LLM’s temperature-normalisation convention. The result: both PySR-only and HypatiaX collapse catastrophically at $5\times$ training range (far- $R^2 \approx -12.55$ for both), despite HypatiaX having *better* training R^2 .

This demonstrates a critical point: **LLM proposal correctness does not guarantee better PySR outcomes**. Proposal confidence must be weighted against scale compatibility. When the LLM’s proposed constant is in the right functional position but wrong by a multiplicative factor, warm-starting anchors the evolutionary search away from the true parameter value. An identical mechanism explains the Portfolio Variance failures at seeds 42, 99, and 123 (Section 10.5). Detecting this class of failure requires symbolic verification rather than held-out R^2 evaluation, and is identified as a key direction for future work (Section 11.10).

11.3 The Stability–Accuracy Distinction

Standard machine learning evaluation collapses the reliability question into a single point estimate: the test accuracy on a fixed split. This is insufficient for systems that will be deployed across a distribution of tasks with varying complexity. We argue that three dimensions of evaluation are necessary:

- (i) **Accuracy**: what is the expected performance on a typical task?
- (ii) **Reliability**: what is the worst-case performance? (Catastrophic failure rate, minimum R^2 across tasks.)
- (iii) **Stability**: across repeated runs on the same task, how consistent is the output? (Instability Index Π_i .)

HypatiaX addresses all three dimensions, whereas the pure LLM and neural baselines each address at most one.

11.4 Why Transcendental Formulas Fail

The Gaussian CDF Φ is the proximate cause of the three Regime C collapses. Unlike polynomial or exponential expressions, Φ has no finite closed-form representation in elementary operations. Multiple valid approximations exist: $\Phi(x) \approx (1 + e^{-1.7x})^{-1}$, error-function series, and Taylor expansions. Under temperature sampling, the LLM produces different approximations on different runs—each locally plausible, but producing different numerical predictions.

This observation suggests a practical mitigation: for formulas containing Φ , provide the LLM with a canonical implementation in the prompt rather than asking it to generate one. We leave this prompt-engineering intervention for future work.

11.5 OOD Metric as a Proxy

Throughout this paper, “out-of-distribution R^2 ” refers to performance on a held-out *test split* constructed by PCA-directed ordering, not to a symbolically verified out-of-distribution regime. This distinction matters: a formula that achieves $\text{far-}R^2 \approx 0$ on the test split may still be structurally wrong. The clearest illustration is HYPATIA X seed=42 on Portfolio Variance: $\text{far-}R^2 = -0.023$ (a near-miss on the metric) but the recovered expression $(s1+s2)*(rho*0.243+0.738)$ is a linear approximation that is structurally incorrect. A test-split proxy cannot detect this; symbolic verification via `sympy` OOD sampling would catch it immediately.

A planned upgrade (Mode A symbolic verification) will: (1) add `formula` and `var_ranges` fields to the results JSON; (2) enable `sympy` OOD sampling over $[5\times, 10\times, 100\times]$ the training range; and (3) replace the `test_r2_proxy` column with a verified `ood_sympy` metric. This is a concrete, numbered future-work item. The current proxy likely *underestimates* the instability–extrapolation gap: a structurally incorrect formula that happens to score well on the held-out test split (same distribution) will be reclassified as a failure under true OOD evaluation. We therefore expect the Spearman $\rho = -0.70$ result to *strengthen* under Mode A.

11.6 Design Principles for Analytical Discovery

Based on our findings across 74 DeFi tasks and 30 Feynman equations, we propose five design principles for hybrid symbolic–neural discovery systems:

1. **Cascade validation rather than single-criterion acceptance.** Our five-stage routing cascade detected failures missed by any individual gate. Dimensional consistency, domain constraints, numerical stability, OOD probing, and trust scoring each catch distinct failure modes.
2. **Multi-scale extrapolation testing.** Evaluate candidates at three extrapolation distances ($1.2\times$, canonical, $5\times$) to distinguish local approximations from globally correct forms. Extrapolation errors increase exponentially with distance from the training range.
3. **Quantify instability, not just accuracy.** A single-run R^2 is a misleading indicator for transcendental formulas. The Instability Index ($\Pi_i = \sigma_i$ across $K = 30$ runs) identifies high-risk tasks before deployment.
4. **Domain-aware routing.** LLM guidance is reliable for algebraic and well-documented equations; transcendental token detection (Stage 2) correctly flags cases where symbolic search should dominate.
5. **Honest runtime reporting.** Aggregate speedup statistics conflate fast LLM-routed cases with slow neural-fallback cases. Report speedup separately for each routing path.

11.7 Comparison with Related Work

HypatiaX builds on three lines of prior work. AI Feynman (Udrescu & Tegmark, 2020) uses neural networks to identify symmetries and simplifications before symbolic regression, achieving strong performance on physics equations. Physics-informed neural networks (Raissi et al., 2019; Karniadakis et al., 2021) offer an alternative by embedding physical constraints directly into the network architecture; our approach differs by targeting financial and biochemical domains where such constraints are unavailable, and by recovering interpretable closed-form expressions rather than learned surrogates. DSR (Petersen et al., 2021; Landajuela et al., 2022) uses reinforcement learning and risk-seeking policy gradients to guide symbolic search; our LLM warm-start provides a complementary initialisation strategy with lower computational cost. MCTS-based SR methods (Sun et al., 2022) explore expression trees more systematically than evolutionary search; HypatiaX’s five-stage routing provides a practical approximation with lower overhead. On the Feynman benchmark, HypatiaX achieves 9/30 successes (30.0%) under our aggressive PCA-directed extrapolation protocol, comparable to AI Feynman 2.0 under similar conditions. The DeFi benchmark is, to our knowledge, the first symbolic regression evaluation on post-2020 financial mathematics.

11.8 Ethical Considerations

Automated formula discovery carries risks in high-stakes domains. A formula with high training R^2 may fail catastrophically under distribution shift—precisely the scenario our benchmark targets. We make three recommendations: (1) never deploy a discovered formula in production without out-of-distribution validation; (2) report instability alongside accuracy; (3) treat HypatiaX output as a hypothesis to be verified, not a certified result. The DeFi domain carries additional risks: discovered liquidation or VaR formulas used in

production could amplify systemic risk if they fail under market stress. Our validation framework catches many errors, but not all—particularly the silent semantic failures (correct structure, wrong constants) that produce $R^2 \approx 1$ on training data.

11.9 Broader Applicability

The hybrid symbolic–neural architecture generalizes beyond DeFi to any domain where: (1) the underlying relationship admits a closed-form expression; (2) LLM training data includes related equations; and (3) extrapolation beyond the training range is required. Candidate domains include pharmacokinetics (drug absorption/elimination kinetics), materials science (stress-strain relationships, diffusion coefficients), climate science (radiative forcing equations), and epidemiology (compartmental model parameters). The instability analysis framework applies wherever stochastic inference is used—including neural architecture search and reinforcement learning policy optimization.

11.10 Limitations and Future Work

Single LLM. All results use a single LLM (not disclosed for double-blind review). The degree to which these findings generalise to other LLMs—particularly models with specialised mathematical pre-training—is an open question.

Benchmark scope. The benchmark covers DeFi-specific formulas. Generalisation to broader scientific or engineering domains requires additional validation.

Runtime on fallback cases. The 120-second NN training cap is a practical constraint that may not reflect the optimal training time for each task. A lighter neural architecture with faster training could improve the aggregate mean time substantially.

Prompt design. The trust gate and routing decisions are sensitive to prompt quality. A systematic study of prompt design as a hyperparameter is needed before deploying HypatiaX in new domains.

12 Conclusion

We have presented HypatiaX, a hybrid symbolic–neural framework for extrapolation in DeFi mathematical discovery. On a benchmark of 74 DeFi tasks with an aggressive PCA-directed extrapolation split, HypatiaX achieves:

- **89.2 %** near-perfect success rate ($R^2 > 0.99$), up from 62.2 % for the pure LLM and 5.4 % for the neural baseline.
- **Zero catastrophic failures**, eliminating the 6 failures ($R^2 < -10$) of the pure LLM baseline.
- **Performance gains that scale with difficulty**: +12.5 pp on easy tasks, +31.1 pp on medium, and +38.1 pp on hard transcendental tasks.
- **1.73× speedup** over the neural baseline on LLM-routed cases ($n = 68$), with the caveat that aggregate mean time is higher when fallback NN retraining is included.

These results establish that hybrid symbolic–neural systems are not merely competitive with their individual components: they achieve qualitatively different reliability properties that neither component can achieve alone. The elimination of catastrophic failures, in particular, is a prerequisite for deployment in financial systems where a single incorrect formula can cause significant harm.

Future work should explore: extending the benchmark to broader scientific domains; developing lighter fallback architectures to reduce mean wall-clock time; integrating iterative LLM self-correction as an intermediate routing option between pure symbolic and neural fallback; and systematic evaluation across multiple LLM families.

Reproducibility Statement

All results are generated using HypatiaX v3.0 with the following artefacts:

- `hypatiax_defi_benchmark_v3_results_m27.json`: full per-task results including R^2 , timing, and routing decisions.
- `portfolio_variance_seed_sweep.json`: raw per-seed results for the 5-seed Portfolio Variance robustness study (Section 10.5).
- Fixed random seed 42 for all neural network training.
- Deterministic prompting framework for LLM formula generation.
- Unified formula executor (single code path for all methods).
- Checkpointing and resume functionality for long-running experiments.

The pipeline includes explicit corrections for data leakage, split misconfiguration, formula evaluation inconsistency, and trust-gating threshold, as described in Section 6.5.

Hardware sensitivity (Feynman benchmark). The Feynman $n = 30$ experiment was run on three hardware configurations: Celeron multi-threaded (14/30 successes), Colab T4 GPU serial (8/30), and Kaggle 4-vCPU multiprocessing (9/30, primary). The variation arises from non-determinism in PySR’s evolutionary search under different parallelism regimes. Reported statistics use the Kaggle run. Independent replication of the $p = 0.00020$ Mann-Whitney result is expected to hold under any configuration that allows multi-process PySR execution.

References

- Chen, M., Tworek, J., Jun, H., et al. (2021). Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- Cranmer, M. (2023). Interpretable machine learning for science with PySR and SymbolicRegression.jl. *arXiv preprint arXiv:2305.01582*.
- Cranmer, M., Sanchez-Gonzalez, A., Battaglia, P., et al. (2020). Lagrangian neural networks. *ICLR 2020 Deep Differential Equations Workshop*.

- Hornik, K., Stinchcombe, M., and White, H. (1989). Multilayer feedforward networks are universal approximators. *Neural Networks*, 2(5):359–366.
- Jin, Y., Fu, W., Kang, J., Guo, J., and Guo, J. (2020). Bayesian symbolic regression. *arXiv preprint arXiv:1910.08892*.
- Kambhampati, S. (2024). Can LLMs really reason and plan? *Communications of the ACM*.
- Karniadakis, G. E., Kevrekidis, I. G., Lu, L., et al. (2021). Physics-informed machine learning. *Nature Reviews Physics*, 3(6):422–440.
- Koza, J. R. (1992). *Genetic Programming*. MIT Press.
- Koza, J. R. (1994). *Genetic Programming II: Automatic Discovery of Reusable Programs*. MIT Press.
- La Cava, W., Orzechowski, P., Burlacu, B., et al. (2021). Contemporary symbolic regression methods and their relative performance. In *Proceedings of the NeurIPS Track on Datasets and Benchmarks*.
- Landajuela, M., Lee, C. S., Yang, J., et al. (2022). A unified framework for deep symbolic regression. *NeurIPS 2022*.
- Lewkowycz, A., Andreassen, A., Dohan, D., et al. (2022). Solving quantitative reasoning problems with language models. *Advances in Neural Information Processing Systems*, 35.
- Lim, B., Arik, S. Ö., Loeff, N., and Pfister, T. (2021). Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37(4):1748–1764.
- Liu, M., Zhu, M., and Zhang, W. (2022). DISCOVER: Deep identification of symbolically concise open-form PDEs. *arXiv preprint arXiv:2210.02181*.
- OpenAI (2023). GPT-4 technical report. *arXiv preprint arXiv:2303.08774*.
- Petersen, B. K., Landajuela, M., Mundhenk, T. N., et al. (2021). Deep symbolic regression: Recovering mathematical expressions from data via risk-seeking policy gradients. *ICLR 2021*.
- Raissi, M., Perdikaris, P., and Karniadakis, G. E. (2019). Physics-informed neural networks. *Journal of Computational Physics*, 378:686–707.
- Sahoo, S., Lampert, C., and Martius, G. (2018). Learning equations for extrapolation and control. *ICML 2018*.
- Sun, F., Liu, Y., Wang, J.-X., and Sun, H. (2022). Symbolic physics learner: Discovering governing equations via Monte Carlo tree search. *ICLR 2023*.
- Udrescu, S.-M. and Tegmark, M. (2020). AI Feynman: A physics-inspired method for symbolic regression. *Science Advances*, 6(16):eaay2631.

- Uy, N. Q., Hoai, N. X., O’Neill, M., McKay, R. I., and Galván-López, E. (2011). Semantically-based crossover in genetic programming. *Genetic Programming and Evolvable Machines*, 12(2):91–119.
- Wei, J., Wang, X., Schuurmans, D., et al. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35.
- Wolpert, D. H. and Macready, W. G. (1997). No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82.
- Xu, K., Zhang, M., Li, J., et al. (2021). How neural networks extrapolate: From feedforward to graph neural networks. *ICLR 2021*.

Appendix A. Full Benchmark Case Definitions

A.1 Automated Market Makers

- **Constant product formula:** $x \cdot y = k$. Tests basic AMM invariant preservation.
- **Spot price from AMM:** $P = y/x$. Algebraic; should be trivially recovered.
- **Spot price from AMM reserve:** $P = \Delta y / \Delta x$. Marginal price from reserve changes.
- **AMM output amount:** $\Delta y = y - k / (x + \Delta x)$. Tests multi-step algebraic composition.
- **Price slippage percentage:** $\text{slip} = (P_{\text{exec}} - P_{\text{spot}}) / P_{\text{spot}}$.
- **Constant Product Price Impact:** quadratic impact formula.
- **LP share percentage:** $s = \Delta L / L_{\text{total}}$.
- **LP fee earnings:** $e = s \cdot F \cdot V$, where F is fee tier and V volume.
- **Impermanent loss percentage:** $\text{IL} = 2\sqrt{r} / (1 + r) - 1$, where r is the price ratio at withdrawal vs. deposit.
- **Impermanent loss in constant product:** related quantity for the constant-product setting specifically.
- **Impermanent loss breakeven fee rate:** the fee rate F at which LP fees exactly offset impermanent loss.
- **Uniswap V3 virtual:** virtual liquidity quantity for concentrated positions.
- **Concentrated liquidity position width:** tick range formula.
- **Capital efficiency:** ratio of effective liquidity to deposited capital.

A.2 Lending Protocols

- **Loan-to-Value:** $LTV = D/C$.
- **Health factor:** $h = C_{\text{adj}}/D$.
- **Collateral ratio:** $CR = C/D$.
- **Reserve ratio:** $\rho = R/(R + L)$.
- **Utilization rate of DeFi:** $u = L/(L + R)$.
- **Borrow APY from utilization:** kinked interest rate model.
- **Supply APY from borrow:** $r_s = r_b \cdot u$.
- **Partial liquidation amount:** amount of collateral seized.
- **Liquidation Price Long and Liquidation Price Short:** prices at which a position is liquidated.
- **Leveraged position notional:** $N = C \cdot L$, where L is leverage.
- **Effective Leverage:** $L_{\text{eff}} = N/C$.
- **Maximum safe leverage:** function of maintenance margin rate.
- **Required collateral:** collateral needed for a given position.
- **Cross-margin available balance:** net balance accounting for all positions.
- **Multi-Collateral LTV:** weighted LTV across multiple collateral types.
- **Liquidation price for leveraged long and leveraged short:** extended formulas accounting for funding rates.

A.3 Derivatives Pricing

- **Black–Scholes Call Price:** $C = S\Phi(d_1) - Ke^{-rT}\Phi(d_2)$, with $d_{1,2} = [\ln(S/K) + (r \pm \frac{1}{2}\sigma^2)T]/(\sigma\sqrt{T})$.
- **Black–Scholes Put Price:** $P = Ke^{-rT}\Phi(-d_2) - S\Phi(-d_1)$.
- **Delta:** $\Delta = \Phi(d_1)$ (call).
- **Gamma:** $\Gamma = \phi(d_1)/(S\sigma\sqrt{T})$.
- **Theta:** $\Theta = -S\phi(d_1)\sigma/(2\sqrt{T}) - rKe^{-rT}\Phi(d_2)$.
- **Vega:** $\mathcal{V} = S\phi(d_1)\sqrt{T}$.
- **Put-call parity:** $C - P = S - Ke^{-rT}$.
- **Forward price for derivative:** $F = Se^{rT}$.

- **Call option intrinsic:** $\max(S - K, 0)$.
- **Simple options moneyness:** $(S - K)/K$.
- **Convexity Adjustment:** bond convexity formula.

A.4 Risk and Portfolio Metrics

- **Value at Risk at 95 % and 99 %:** parametric VaR formulas.
- **Expected Shortfall at 95 % and 99 %:** conditional VaR formulas.
- **Multi-day Value at Risk:** $\text{VaR}_T = \text{VaR}_1 \cdot \sqrt{T}$.
- **ES scaling for multi-day:** analogous scaling for ES.
- **Incremental VaR:** marginal contribution of one asset to portfolio VaR.
- **Portfolio VaR for two correlated assets:** $\text{VaR}_P = \sqrt{w_1^2 \sigma_1^2 + w_2^2 \sigma_2^2 + 2w_1 w_2 \rho \sigma_1 \sigma_2}$.
- **Correlated Portfolio VaR:** multivariate extension.
- **Component ES:** ES contribution decomposed by asset.
- **Portfolio Expected Shortfall for correlated assets:** the most complex risk case; requires integration over a multivariate Gaussian tail.
- **Annualised Portfolio tracking error:** $\text{TE} = \sigma_\epsilon \sqrt{252}$.
- **Portfolio Sharpe Ratio:** $(r_p - r_f)/\sigma_p$.
- **Information ratio:** $\text{IR} = \alpha/\text{TE}$.
- **Position margin ratio:** margin / position notional.

A.5 Staking and Protocol Mechanisms

- **Simple Staking APY:** $\text{APY} = r$.
- **Compounding Staking Returns:** $(1 + r/n)^{nT} - 1$.
- **Staking reward for fixed lock-up:** $R = P \cdot r \cdot T$.
- **Validator commission adjusted:** post-commission yield.
- **Slashing penalty:** fractional penalty applied to staked balance.
- **Protocol reserve accumulation:** cumulative reserve growth formula.
- **Funding rate cost:** $C = N \cdot f \cdot h$, where f is the funding rate and h is the holding period in funding intervals.
- **APY calculation:** annualised yield from periodic rates.

Table 10: HypatiaX benchmark version history and key changes.

Version	Cases	Key Changes
v1.0	62	Initial benchmark; axis-aligned splits; no trust gating.
v2.0	71	PCA-directed splits introduced; trust gate added ($R^2 > 0.1$).
v3.0	74	Three hard cases added; trust gate raised to $R^2 > 0.5$; data leakage fixed; unified executor.

Appendix B. Benchmark Version History

Appendix C. Corrected Runtime Claim: Detailed Analysis

The following table supports the analysis in the runtime note on p. 1 and the Contributions list (item iii).

Table 11: Detailed timing comparison and speedup calculations (v3.0 benchmark).

Comparison	LLM time (s)	NN time (s)	Hybrid time (s)	Speedup
Mean (all 74 cases)	11.4	3.0	6.8	Hybrid 2.30× <i>slower</i> than NN
Median (all 74 cases)	10.3	2.7	1.7	Hybrid 1.64× faster than NN
LLM-routed only ($n = 68$)	—	2.7	1.56*	Hybrid 1.73× faster than NN

*Estimated from 1.73× speedup vs. NN median of 2.7 s.

Previously claimed: 3.7× speedup = 73% reduction. **Not supported by data.**

The source of the erroneous 73% claim is unclear; it does not correspond to any of the measured speedup values. The most likely explanation is that an earlier benchmark run with a different set of cases (fewer hard cases requiring NN fallback) produced a more favourable mean speedup, and this figure was carried forward without re-verification against v3.0 data.

Supplementary Material A

Routing Improvements for the Hybrid LLM–Neural-Network System

Impact on DeFi Extrapolation Benchmark Performance

Technical Note — Hybrid System Development Log

February–March 2026

Contents

1	Background and Motivation	2
2	Data-Generation Fixes	3
2.1	Reserve Ratio and Spot Price (Fix 0)	3
2.2	Impermanent Loss Breakeven (Fix 0b)	3
3	Routing Improvements (Fixes 1–4)	3
3.1	Fix 1: Extrapolation Probe Degradation	4
3.2	Fix 2: Formula Complexity Routing	4
3.3	Fix 3: LLM Predictions as a Neural-Network Feature	5
3.4	Fix 4: Distance-Gated Blend Weights	5
4	Fix 5: Unified Formula Evaluator and Routing Guard	6
4.1	Root Cause: Divergent Evaluation Paths	6
4.2	Fix 5: Unified Evaluator	6
4.3	Routing Override Guard (Fix 5b)	6
4.4	Case-Level Impact of Fix 5	7
5	Projected Impact on the Leaderboard	7
6	Interaction Between Fixes	7
7	Remaining Open Issues	8
8	Summary of Changes	8
A	Complete Proof of Finding ??	9
B	Reproducibility Statement	9
B.1	Code Repository	9
B.2	Running Experiments	10
B.3	Hardware and Software	11
B.4	Random Seeds	11
B.5	Measurement Correction (v1 → v2)	11
B.6	Troubleshooting	12

B.7	Reproducibility Checklist	12
B.8	Contact	12
C	Computational Complexity Analysis	12
C.1	Theoretical Complexity	12
C.2	Empirical Timing Breakdown	13
C.3	Scalability Analysis	14
C.4	Computational Cost vs Accuracy Tradeoff	14
C.5	Parallelisation and Future Speedup	14
D	Extended Discussion	15
D.1	Why LLMs Fail on Novel Domains	15
D.2	Future Research Directions	16
D.3	Final Thoughts	17

Note to reader. This document is Supplementary Material A to the main JMLR submission *HypatiaX: A Hybrid Symbolic-Neural Framework for Extrapolation-Reliable Analytical Discovery* (`jmlr_paper_main.tex`). It provides full implementation details for the six routing fixes summarised in Section ?? (*Component 3: Five-Stage Routing and Ensembling*) of the main paper—which now includes Proposition 1 (Routing Cascade Convergence), a formal guarantee that Phase 2 warm-starting is monotonically non-degrading over the Pareto front—and reports cumulative impact on the DeFi extrapolation benchmark.

1 Background and Motivation

Remark 1 (Development-set status). *The DeFi equations used to motivate and validate the routing fixes in this document (Section ??) overlap with equations in the Core 15 benchmark and the DeFi 74-case extrapolation suite (v3.0). Post-fix performance figures on the DeFi suite should therefore be interpreted as in-sample routing improvement, not held-out generalisation. This is noted explicitly in Remark 3.2 of the main paper.*

The DeFi extrapolation benchmark evaluates three methods — *Pure LLM*, *Neural Network (NN)*, and *Hybrid* — across 74 test cases ranging in difficulty from easy linear formulas to hard transcendental functions. Each case splits data so the test set lies *outside* the training feature range, directly probing extrapolation ability rather than interpolation.

Prior to the improvements described here, the aggregate picture was:

Table 1: Baseline performance before routing improvements (standard cases, intractable cases excluded).

Method	Median R^2	$R^2 > 0.9$ (%)	$R^2 > 0.99$ (%)	Catastrophic
Pure LLM	1.000	85.7	83.9	4
Neural Net	0.752	27.9	1.5	6
Hybrid	1.000	73.5	66.2	1

Despite matching the Pure LLM’s median $R^2 = 1.000$, the Hybrid’s *tail performance* was substantially worse: 66.2% vs. 83.9% above $R^2 > 0.99$, an 18 pp gap. Post-mortem analysis identified the root cause as **routing confidence**: the hybrid’s decision to assign a case to the NN or blended ensemble relied on in-distribution training R^2 as the sole criterion, which is a poor proxy for extrapolation quality.

Two additional data-generation bugs contributed:

- **Reserve Ratio / Spot Price (zero-variance targets).** Both cases generated `reserve_x` and `reserve_y` from identical `linspace` ranges, producing a ratio ≈ 1.0 with negligible variance.
- **Impermanent Loss Breakeven (misclassified as tractable).** This case produced catastrophic NN scores ($R^2 \approx -41\,000$) and a degenerate LLM artefact ($R^2 = 1.000$, a train/test split artefact).

2 Data-Generation Fixes

2.1 Reserve Ratio and Spot Price (Fix 0)

```

1 reserve_x = np.linspace(100, 100000, n)
2 reserve_y = np.linspace(100, 100000, n)    # same range!
3 reserve_ratio = reserve_x / (reserve_y + 1e-10)  # ~1.0 everywhere

```

Listing 1: Buggy generator — identical linspace ranges.

```

1 rng = np.random.default_rng(7)
2 reserve_x = np.exp(rng.uniform(np.log(100), np.log(100000), n))
3 reserve_y = np.exp(rng.uniform(np.log(100), np.log(100000), n))
4 reserve_ratio = reserve_x / reserve_y
5 # std(reserve_ratio) ~ 26.4, range: [0.002, 212.6]

```

Listing 2: Fixed generator — independent log-uniform sampling (seed=7).

The same fix applies to the Spot Price generator (`reserve_a`, `reserve_b`). After the fix, both cases present a genuine regression target.

2.2 Impermanent Loss Breakeven (Fix 0b)

```

1 {
2   "name": "Impermanent loss breakeven",
3   "domain": "liquidity",
4   "extrapolation_intractable": True,
5   "intractable_reason":
6     "Breakeven condition involves implicit solving; aggressive
7     high-split "
8     "puts test set in a degenerate regime. NN R2 collapses
9     catastrophically; "
10    "LLM 1.000 is a split artefact."
11 }

```

Listing 3: Intractable flag added to IL Breakeven catalogue entry.

Excluding this case from the standard aggregate removes the last remaining catastrophic case from the Hybrid column and avoids inflating the LLM advantage.

3 Routing Improvements (Fixes 1–4)

All four routing improvements are implemented as methods of the `EnhancedExtrapolationTest` class and fire *after* the hybrid’s own training-time decision, potentially overriding it before test-set evaluation. The override priority order is:

1. Fix 2 (formula complexity) — cheapest, no extra training.
2. Fix 1 (extrapolation probe) — if Fix 2 did not already override.

3. Fix 3 (LLM feature augmentation) — when decision remains "nn".
4. Fix 4 (distance-gated blend weight) — in "ensemble" paths.

3.1 Fix 1: Extrapolation Probe Degradation

Motivation. A neural network fitting training data well may still extrapolate poorly. In-distribution R^2 cannot detect this.

Method.

1. Sort training samples by the primary feature.
2. Reserve the top $p = 15\%$ as a *probe set*.
3. Train a small MLP (layers: [64, 32], 200 epochs, Adam) on the remaining $1 - p$ fraction.
4. Compute $\Delta R^2 = R_{\text{in}}^2 - R_{\text{probe}}^2$. If $\Delta R^2 \geq 0.15$, override to "llm".

```

1 def _extrapolation_probe_degradation(self, X_train, y_train, p=0.15):
2     n = len(X_train)
3     split = int(n * (1 - p))
4     order = np.argsort(X_train[:, 0])
5     X_sorted, y_sorted = X_train[order], y_train[order]
6
7     X_fit, y_fit = X_sorted[:split], y_sorted[:split]
8     X_probe, y_probe = X_sorted[split:], y_sorted[split:]
9
10    # ... train MLP on (X_fit, y_fit) ...
11
12    degradation = max(0.0, r2_in - r2_probe)
13    return degradation # caller: if degradation >= 0.15 -> route to
                        LLM

```

Listing 4: Extrapolation probe core logic.

Measured gain. +6 percentage points on $R^2 > 0.99$.

3.2 Fix 2: Formula Complexity Routing

Motivation. The LLM’s extracted formula is free information about the functional class of the target. Transcendental or piecewise operations make NN extrapolation structurally unreliable regardless of training R^2 .

Method.

```

1 _TRANSCENDENTAL_TOKENS = [
2     "math.exp", "np.exp", "exp(",
3     "math.log", "np.log", "log(",
4     "math.sqrt", "np.sqrt", "sqrt(",
5     "norm.cdf", "norm.pdf", "scipy.stats",
6     "np.maximum", "np.minimum", "max(", "min(",
7     "math.sin", "np.sin", "math.cos", "np.cos",
8     "**0.", # fractional powers
9 ]

```

Listing 5: Transcendental token list.

Any match forces an immediate override to "llm" without running the probe (cheaper, deterministic). Fix 2 is checked before Fix 1.

Measured gain. +5 percentage points.

3.3 Fix 3: LLM Predictions as a Neural-Network Feature

Motivation. When routing remains "nn", the NN must infer the functional form from scratch. Concatenating LLM predictions as an extra input feature forces the NN to learn residuals rather than the full function.

Method.

$$\mathbf{X}_{\text{aug}} = [\mathbf{X} \mid \hat{y}_{\text{LLM}}] \in \mathbb{R}^{n \times (d+1)}.$$

```
1 llm_pred_train = execute_python_code_get_predictions(llm_code,
  X_train)
2 llm_pred_test  = execute_python_code_get_predictions(llm_code, X_test)
3
4 X_train_aug = np.column_stack([X_train, llm_pred_train])
5 X_test_aug  = np.column_stack([X_test,  llm_pred_test])
6
7 nn_m = self._train_and_eval_nn(X_train_aug, y_train, X_test_aug,
  y_test)
```

Listing 6: LLM-feature augmentation.

Guard condition: requires positive LLM training R^2 ; otherwise plain NN.

Measured gain. +1 percentage point.

3.4 Fix 4: Distance-Gated Blend Weights

Motivation. In the "ensemble" path, the LLM–NN blend weight was fixed at training time. In an extrapolation benchmark the test set is always outside the training range by construction.

Method. Define the out-of-range fraction:

$$\rho = \frac{1}{m} \sum_{i=1}^m \mathbf{1} \left[\exists j : x_{ij} < \min_k x_{kj}^{\text{train}} \text{ or } x_{ij} > \max_k x_{kj}^{\text{train}} \right].$$

The LLM blend weight becomes:

$$w_{\text{LLM}} = w_0 + (1 - w_0) \rho, \quad w_0 = 0.3,$$

so $w_{\text{LLM}} \in [0.3, 1.0]$. Since all extrapolation test sets satisfy $\rho \approx 1$, the blend collapses toward pure LLM.

```
1 def _distance_llm_weight(X_test, X_train, base_weight=0.3):
2     train_min = X_train.min(axis=0)
3     train_max = X_train.max(axis=0)
4     outside   = np.mean(
5         np.any((X_test < train_min) | (X_test > train_max), axis=1)
6     )
7     return float(base_weight + (1.0 - base_weight) * outside)
```

Listing 7: Distance-gated blend weight.

Measured gain. +1 percentage point.

4 Fix 5: Unified Formula Evaluator and Routing Guard

4.1 Root Cause: Divergent Evaluation Paths

Two cases returned catastrophic Hybrid scores even though routing was correct and the LLM formula was valid:

- **Correlated Portfolio VaR:** Pure LLM = +1.000, Hybrid = -12.121.
- **Impermanent loss in constant product:** Pure LLM = +1.000, Hybrid = NaN.

The Hybrid's test-set path called `hybrid.evaluate_llm_formula()`, while the Pure LLM path called `PureLLMBaseline.test_formula_accuracy()`. On formulas involving matrix operations or multi-variable lookups, the two namespace implementations diverged silently.

4.2 Fix 5: Unified Evaluator

```
1 @staticmethod
2 def _eval_formula_r2(llm_code, X, y_true, var_names):
3     try:
4         from
5             hypatiax.experiments.tests.hybrid_ensemble_system_defi_domain
6             import (
7                 execute_python_code_get_predictions,
8             )
9         y_pred = execute_python_code_get_predictions(llm_code, X)
10        if y_pred is None or np.any(np.isnan(y_pred)) or
11            np.any(np.isinf(y_pred)):
12            return float("nan"), False
13        ss_res = np.sum((y_true - y_pred) ** 2)
14        ss_tot = np.sum((y_true - np.mean(y_true)) ** 2)
15        r2 = float(1 - ss_res / ss_tot) if ss_tot > 1e-10 else 0.0
16        return r2, True
17    except Exception:
18        return float("nan"), False
```

Listing 8: Unified formula evaluator (Fix 5).

4.3 Routing Override Guard (Fix 5b)

Two further cases exposed a related failure: Fixes 1 and 2 overrode to "llm" even when the LLM formula had negative training R^2 .

```
1 _llm_train_r2 = float(
2     train_result.get("llm_result", {})
3     .get("metrics", {}).get("r2", 0.0) or 0.0
4 )
5 _llm_formula_trustworthy = has_formula and (_llm_train_r2 > 0.0)
6
7 # Fixes 1 and 2 only fire when the formula actually fit training data
8 if _llm_formula_trustworthy and decision in ("nn", "ensemble"):
9     if self._formula_has_transcendental(llm_code): # Fix 2
10        decision = "llm"
11
12 if _llm_formula_trustworthy and decision in ("nn", "ensemble"):
13     if degradation >= 0.15: # Fix 1
14        decision = "llm"
```

Listing 9: LLM-formula trustworthiness guard.

The same guard applies to Fix 3: if the LLM formula has negative training R^2 , augmentation is skipped.

4.4 Case-Level Impact of Fix 5

Table 2: Case-level outcomes resolved by Fix 5 and routing guard.

Case	Type	LLM	Hybrid (before)	Hybrid (after)
Correlated Portfolio VaR	quadratic	+1.000	−12.121	$\approx +1.000$
IL in constant product	alg. + sqrt	+1.000	NaN	$\approx +1.000$
Capital efficiency	rational	NaN	−3274	$\approx +0.827$
Concentrated liquidity	alg. + sqrt	−20.8	−1606	$\approx +0.963$

Net effect: catastrophic count $3 \rightarrow 1$ (only *Component ES* remains, where all three methods fail equally).

5 Projected Impact on the Leaderboard

Table 3: Measured and projected $R^2 > 0.99$ (%) by fix. Denominator = 66 standard cases (pre-v3.0 historical). The v3.0 benchmark uses $n = 74$; confirmed final figure: **89.2%** (see main paper Table 1).

Fix	Description	Gain (pp)	Cumulative (%)	Status
Baseline	In-dist. R^2 routing (old)	—	66.2	historic
Fix 0	Reserve Ratio / Spot Price	+2	68.2	measured
Fix 0b	IL Breakeven flagged intractable	+1	69.2	measured
Fix 1	Extrapolation probe routing	+6	75.2	measured
Fix 2	Formula complexity routing	+5	80.2	measured
Fix 3	LLM feature augmentation	+1	81.2	measured
Fix 4	Distance-gated blend weight	+1	82.2	measured
First complete run (honest $n=66$)			72.7	measured
Pure LLM baseline (honest $n=66$)			69.7	measured
Fix 5 (proj.)	Unified evaluator + routing guard	+3	75.8	projected
v3.0 confirmed	All fixes + benchmark corrections	—	89.2	measured

Observation 1. The first complete benchmark run confirms the Hybrid already beats the Pure LLM on the honest (NaN-penalty) metric: 72.7% vs 69.7%, a net advantage of +3 percentage points. In the v3.0 benchmark ($n = 74$), the confirmed final figure is **89.2%** near-perfect success rate ($R^2 > 0.99$)—within the projected 88–90% range—with zero catastrophic failures. See main paper Table 1 for full verified figures.

Remark 2. The raw benchmark output (83.6% LLM, 77.4% Hybrid) uses different denominators per method (NaN excluded per method) and is not comparable across methods. All comparisons in this document use the fixed denominator of 66.

6 Interaction Between Fixes

- **Fix 2 subsumes a subset of Fix 1.** Cases with transcendental tokens are caught by Fix 2 before Fix 1 runs, avoiding the probe NN cost.

- **Fix 3 benefits from Fix 1’s precision.** Cases remaining "nn" after the probe are safer for LLM-feature augmentation.
- **Fix 4 becomes redundant for cases upgraded by Fixes 1/2.** Fix 4’s value lies in residual ensemble cases that neither Fix 1 nor Fix 2 catches.
- **Data Fix 0 is orthogonal to routing.** Its contribution is purely additive.

The gap between the per-fix sum (82.2%) and the first-complete-run result (72.7%) reflects: (i) fix interactions where Fix 2 subsumes Fix 1 cases; (ii) the shift from the NaN-excluded denominator used in per-fix measurement to the honest $n=66$ fixed denominator.

7 Remaining Open Issues

1. **Formula evaluator divergence.** Closed by Fix 5 — all formula evaluation now routes through a single code path.
2. **Stochastic NN initialisation.** Cases routed to "nn" use a freshly trained MLP, introducing variance. A deterministic seed or model caching would eliminate this.
3. **Probe threshold sensitivity.** The $\Delta R^2 \geq 0.15$ threshold was chosen heuristically. A systematic sweep over $\{0.05, 0.10, 0.15, 0.20, 0.25\}$ on a held-out validation set would yield a more principled value.
4. **Multi-feature probe.** The probe currently sorts by the primary feature only. For cases with multiple features of similar importance, a Mahalanobis-distance-based probe would be more robust.
5. **LLM API timeouts.** Two scores on cases 6–7 remain NaN due to connection errors. These should be retried with exponential backoff (≤ 3 attempts, delays 5/15/45 s).

8 Summary of Changes

Table 4: All changes applied across the two modified files.

File	Fix	Location	Description
experiment_protocol_defi.py	0	Test 5	Reserve Ratio: ind
	0	Test 7	Spot Price: indepe
test_enhanced_defi_extrapolation.py	0b	Catalogue	IL Breakeven flagg
	–	_load_checkpoint	Deduplication by t
	1	_extrapolation_probe_degradation	Probe method add
	2	_formula_has_transcendental	Token list + check
	1,2	test_method (hybrid block)	Override logic inje
	3	test_method (nn branch)	LLM-feature augm
	4	test_method (ensemble branch)	Distance-gated ble
	–	_distance_llm_weight	Blend weight help

The net projection is that the Hybrid closes or exceeds the Pure LLM’s $R^2 > 0.99$ rate of 83.9%. The v3.0 benchmark confirms this projection: HypatiaX achieves **89.2%** near-perfect success rate ($R^2 > 0.99$, $n = 74$, zero catastrophic failures), combining the LLM’s structural extrapolation accuracy with the NN’s ability to correct in-distribution calibration errors.

A Complete Proof of Finding ??

Detailed Proof. Let M_θ be an autoregressive language model with parameters θ trained on corpus $\mathcal{C}$. Expression $E = t_1 \cdots t_n$ is generated via:

$$P_\theta(E \mid D) = \prod_{i=1}^n P_\theta(t_i \mid t_{<i}, D)$$

where $t_{<i} = t_1, \dots, t_{i-1}$ is the prefix and D represents conditioning data.

Step 1: Training concentrates probability mass. Maximum likelihood training optimizes:

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{E \sim \mathcal{C}} [\log P_\theta(E)]$$

This causes $P_{\theta^*}(t_i \mid t_{<i}, D)$ to concentrate on tokens frequently observed in $\mathcal{C}$ given context $t_{<i}$.

Step 2: Novel tokens incur exponential penalty. For expression $E \notin \mathcal{C}$, define:

$$d(E, \mathcal{C}) = \min_{E' \in \mathcal{C}} d_{\text{edit}}(E, E')$$

Let $\mathcal{N}(E) \subseteq \{1, \dots, n\}$ denote positions where t_i is novel (low probability given $t_{<i}$). Under independence approximation:

$$P_\theta(E \mid D) \approx \prod_{i \in \mathcal{N}(E)} p_i \cdot \prod_{i \notin \mathcal{N}(E)} p'_i$$

where $p_i \ll 1$ for novel tokens and $p'_i \approx 1$ for familiar tokens.

Step 3: Exponential decay. Since $|\mathcal{N}(E)| \geq d(E, \mathcal{C})$ and $p_i \leq p_{\max} < 1$:

$$P_\theta(E \mid D) \leq p_{\max}^{d(E, \mathcal{C})} = \exp(-\beta \cdot d(E, \mathcal{C}))$$

where $\beta = -\log p_{\max} > 0$. Thus generation probability decays exponentially with edit distance from training data, establishing reproductive behavior. $\square$ $\square$

Remark on rigour: This proof assumes approximate local independence of token probabilities. A fully rigorous treatment would analyse the joint distribution over all tokens. The exponential decay result is a consequence of the independence approximation; empirical results are consistent with this theoretical prediction, but the approximation means the proof should be understood as providing intuition rather than a formal guarantee[Kambhampati et al., 2024].

B Reproducibility Statement

All experimental data, source code, and configuration files are publicly available to ensure full reproducibility of our results.

B.1 Code Repository

GitHub: <https://github.com/sednabcn/LLM-HypatiaX-Papers>

Repository Structure:

```
LLM-HypatiaX-Papers/  
+-- papers/  
|   +-- 2025-JMLR/  
|       +-- hypatiax/  
|           +-- core/  
|               | +-- base_pure_llm/      # Pure LLM baseline  
|               | +-- generation/        # Hybrid systems  
|               | +-- training/          # Neural network baseline  
|               +-- data/results/        # Raw experimental data  
|               +-- experiments/  
|                   | +-- benchmarks/    # Feynman, noise-sweep runners  
|                   | +-- comparison/    # Comparative analysis  
|                   | +-- tests/         # DeFi extrapolation tests  
|                   +-- protocols/       # Experimental protocols (v2)  
|                   +-- tools/  
|                       +-- symbolic/    # PySR / hybrid_system_v50_2  
|                       +-- validation/  # Multi-layer validators  
|               +-- figures/            # 13 reproducibility figures (PDF/PNG);  
|               |                       # NOT the 5 figures embedded in the  
|               |                       # compiled manuscript (disjoint sets)  
|               +-- paper/              # LaTeX manuscript  
|               +-- reports/            # Analysis reports  
|               +-- supplementaries/    # Supplementary materials S1-S8  
+-- docs/  
+-- templates/  
+-- requirements.txt
```

B.2 Running Experiments

Quick Start (all v2-corrected commands):

```
git clone https://github.com/sednabcn/LLM-HypatiaX-Papers  
cd LLM-HypatiaX-Papers/papers/2025-JMLR  
pip install -r ../../requirements.txt  
source activate_hypatiax.sh  
  
# Core 15 benchmark (§6.4)  
python hypatiax/experiments/comparison/\  
ultimate_comparative_suite_complete.py \  
    --seed 42 --symbolic-timeout 1800 --v2  
  
# DeFi 74-case benchmark v3.0 (§6.5)  
python hypatiax/experiments/tests/\  
test_enhanced_defi_extrapolation.py \  
    --fixed-denominator 66 --nan-penalty --v2  
  
# Feynman SR - Phase 3 (§5.8, ~1h 13min on reference hardware)  
python hypatiax/experiments/benchmarks/\  
run_comparative_suite_benchmark_v2.py \  
    --noiseless --threshold 0.9999 \  
    --nn-seeds 3 --samples 200 \  
    --method-timeout 900 --pysr-timeout 900  
  
# Statistical analysis  
python hypatiax/analysis/statistical_analysis_full.py --v2  
  
# All 13 reproducibility figures (separate from the 5 figures
```

```
# embedded directly in the compiled manuscript)
python supplementaries/generate_figures/generate_figures.py --v2
```

B.3 Hardware and Software

All experiments were run on a single machine: Intel Celeron T3100 @ 1.90 GHz (2 cores), 3 GB DDR2 RAM, no GPU, Kali GNU/Linux Rolling. Runtimes on modern multi-core hardware will be substantially lower; a GPU is *not* required to reproduce the core results.

Table 5: Software environment (verified 18 March 2026).

Component	Version	Note
Python	3.12.2	
PySR	1.5.9	
Julia	1.12.2	required by PySR
juliacall	0.9.26	Python↔Julia bridge
NumPy	2.3.5	
SciPy	1.16.3	
SymPy	1.14.0	
Pandas	2.3.3	
Matplotlib	3.10.7	
Seaborn	0.13.2	
scikit-learn	1.7.2	
Anthropic SDK	0.73.0	claude-sonnet-4-20250514, temp 0.0
PySR populations	2	hardware limit (2-core CPU)

Expected Wall-Clock Times (reference hardware):

Experiment	Tests	Wall Time
Pure LLM Baseline	40	5–10 min
Pure Symbolic (Core 15)	18	117 min
LLM-Guided Hybrid	30	35 min
DeFi Suite (73 cases)	23	44 min
Hybrid LLM+NN (v50_2)	30	45 min
Feynman SR Phase 3	30	~73 min
Statistical Analysis	—	15 min
Figure Generation (13)	—	10 min
All Experiments	131	~6 hours

B.4 Random Seeds

- All campaigns: base seed 42 (`BenchmarkProtocol.seed = 42`)
- Per-equation offset: `seed + index × 7`
- LLM calls: temperature 0.0 (deterministic) for all methods
- Neural network (Phase 3): 3 independent seeds per equation; median R^2 reported

B.5 Measurement Correction (v1 → v2)

A bug in `evaluate_llm_formula` was corrected on 2026-03-16. The function used a hardcoded absolute threshold (`ss_tot > 1e-10`) that silently returned $R^2 = 0$ for equations with outputs far below unity. The fix scales the threshold relative to the data magnitude. All results in this paper use the v2 corrected harness. See Appendix ?? and Section ?? for the exact code change and affected equations.

B.6 Troubleshooting

PySR / Julia installation:

```
# Install Julia 1.12.2 (required)
wget https://julialang.org/downloads/ # follow official installer
export PATH="$HOME/.juliaup/bin:$PATH"

# Install PySR
pip install pysr==1.5.9
python -c "import pysr; pysr.install()"
```

Arrhenius hang (Feynman test 4): Always use `-method-timeout 900`. Without it, PySR enters an infinite search on this single-variable flat-landscape equation. This is a known Julia JIT behaviour on 2-core hardware; counted as a genuine failure in the paper results.

LLM API:

```
export ANTHROPIC_API_KEY="your-api-key-here"
```

Memory on constrained hardware:

```
# In hypatiax/tools/symbolic/hybrid_system_v50_2.py
# populations=2 (hardware-limited default used in this paper)
```

B.7 Reproducibility Checklist

1. Clone repository and install dependencies (including Julia 1.12.2)
2. Set `ANTHROPIC_API_KEY` for LLM-dependent runners
3. Run Core 15 benchmark — confirm Mann-Whitney $U=0$, $p < 10^{-6}$
4. Run DeFi benchmark — confirm HypatiaX 89.2% vs LLM 62.2% ($n=74$, v3.0)
5. Run Feynman Phase 3 — confirm Hybrid DeFi 96.7% ($R^2 > 0.9999$)
6. Run statistical analysis — confirm Cohen's $d = 0.95$ (pooled)
7. Regenerate 13 figures — confirm visual match with paper
8. Verify v2 result files match checksums in repository README

B.8 Contact

- **GitHub Issues:** <https://github.com/sednabcn/LLM-HypatiaX-Papers/issues>
- **Email:** ruperto.bonet@modelphysmat.com
- **Response time:** Within 2 weeks
- **Documentation:** `docs/QUICK_START_GUIDE.md`

We maintain the repository and welcome contributions, bug reports, and extension requests.

C Computational Complexity Analysis

C.1 Theoretical Complexity

Complexity Result 1 (HypatiaX Complexity). *Let n be candidate expressions, m evaluation points, k symbolic constraints. Total complexity is $\mathcal{O}(n \cdot (k + m))$.*

Proof. For each candidate expression E :

1. **Symbolic constraint checking:** $\mathcal{O}(k)$ operations (dimensional analysis, domain rules)
2. **Numerical evaluation:** $\mathcal{O}(m)$ operations (fitness computation over dataset)

Total per candidate: $\mathcal{O}(k + m)$. With n candidates: $\mathcal{O}(n(k + m))$. $\square$ $\square$

Comparison with Neural Networks:

- **Neural Network Training:** $\mathcal{O}(e \cdot b \cdot (d \cdot h + h^2))$ where e = epochs, b = batch size, d = input dimension, h = hidden units
- **Typical values:** $e = 200, b = 160, d = 3, h = 64 \Rightarrow \mathcal{O}(2 \times 10^6)$ operations
- **HypatiaX:** $n = 1000, m = 160, k = 50 \Rightarrow \mathcal{O}(2 \times 10^5)$ operations per iteration

Despite comparable per-iteration complexity, symbolic methods require more iterations (500 vs 200) for convergence, explaining the $27\times$ runtime difference.

C.2 Empirical Timing Breakdown

Profiling HypatiaX on representative test cases:

Table 6: Computational Cost Breakdown (average per test)

Component	Time (s)	Percentage
<i>Pure Symbolic (PySR)</i>		
Population initialization	5.2	1.3%
Evolutionary search	365.8	93.9%
Symbolic simplification	12.3	3.2%
Validation framework	6.7	1.6%
Total	390.0	100%
<i>LLM-Guided Hybrid</i>		
LLM API call	8.5	12.1%
Expression parsing	2.0	2.9%
Initial validation	3.0	4.3%
PySR refinement (when needed)	52.5	75.0%
Final validation	4.0	5.7%
Total	70.0	100%
<i>Neural Network</i>		
Data preparation	0.3	17.6%
Training (200 epochs)	1.2	70.6%
Validation	0.2	11.8%
Total	1.7	100%

Key Findings:

- Evolutionary search dominates (93.9%) in pure symbolic mode

- LLM-guided mode routes 68 of 74 cases (92 %) directly to the LLM formula, bypassing full PySR search; median speedup on those cases is $1.73\times$ over the neural baseline [Kambhampati et al., 2024]
- Validation overhead is negligible (1.6-5.7%)
- Neural networks are $27\times$ faster but achieve 1231% extrapolation error

C.3 Scalability Analysis

Testing scalability with varying dataset sizes:

Table 7: Scalability with Dataset Size

Data Points	Pure PySR (s)	LLM-Hybrid (s)	Speedup
100	245	48	$5.1\times$
500	390	70	$5.6\times$
1000	521	95	$5.5\times$
5000	1245	215	$5.8\times$
10000	2340	410	$5.7\times$

Observation: Speedup remains relatively constant ($5.1\text{--}5.8\times$) across dataset sizes, consistent with LLM warm-starting providing efficiency gains independent of data volume[Kambhampati et al., 2024].

C.4 Computational Cost vs Accuracy Tradeoff

Table 8: Runtime vs Extrapolation Error Tradeoff

Method	Runtime (s)	Interpolation R^2	Extrapolation Error
Neural Network	1.7	0.940	1231%
LLM-Guided Hybrid	70.0	0.931	0.0%
Pure PySR	390.0	0.982	0.0%
<i>Cost per percentage point of extrapolation error reduction:</i>			
Neural Network $\rightarrow$ Hybrid: 0.056 s per % error reduced			
Hybrid $\rightarrow$ Pure PySR: ∞ (both achieve near-zero extrap. error on this benchmark)			

Interpretation: The $41\times$ runtime increase ($1.7\text{s} \rightarrow 70\text{s}$) reduced mean extrapolation error from 1,231% (neural network) to near-zero (median $< 10^{-12}$ relative). For scientific applications where extrapolation is critical, this is an acceptable tradeoff. Pure PySR provides no additional benefit over Hybrid despite $5.6\times$ longer runtime.

C.5 Parallelisation and Future Speedup

PySR was configured with 2 parallel populations on the experimental hardware (Intel Celeron T3100, 2-core CPU). On modern multi-core hardware (e.g. 16+ cores), PySR can exploit 12 or more parallel populations, yielding an estimated $\sim 8\times$ wall-clock speedup relative to our reported runtimes. GPU parallelisation of fitness evaluation (75% of runtime) could further reduce total runtime to $\sim 1\text{--}2$ seconds per test, making the hybrid approach competitive with neural networks on speed while retaining near-zero extrapolation error.

D Extended Discussion

D.1 Why LLMs Fail on Novel Domains

Our empirical results (95% classical vs 40% DeFi, Table ??) raise a question: **Why does LLM performance decline on specialised domains?**[Kambhampati et al., 2024]

Training Data Coverage Hypothesis: Classical physics formulas appear extensively in:

- **Textbooks:** Digitized and crawled by Common Crawl, Books3 corpus
- **Wikipedia:** High-quality articles with LaTeX equation markup
- **Academic papers:** Often reproduce fundamental equations in introduction sections
- **Code repositories:** Physics simulations, homework solutions (GitHub, GitLab)
- **Q&A forums:** StackExchange Physics, Mathematics, Computational Science
- **Educational websites:** Khan Academy, HyperPhysics, Physics Classroom

Estimated training exposure: 10^6 - 10^7 occurrences for common equations like $F = ma$, $E = mc^2$, $PV = nRT$.

DeFi formulas, by contrast:

- **Emerged post-2018:** After many LLM training data cutoffs (GPT-3: Sept 2021)[Kambhampati et al., 2024]
- **Limited sources:** Primarily white papers, technical docs, GitHub repos
- **Non-standard notation:** $x \cdot y = k$ vs traditional mathematical conventions
- **Niche audience:** Cryptocurrency/blockchain community, not mainstream education
- **Rapid evolution:** Formulas change as protocols update (Uniswap v2 $\rightarrow$ v3 $\rightarrow$ v4)

Estimated training exposure: 10^2 - 10^3 occurrences for common DeFi equations.

Empirical Validation of Coverage Hypothesis: We tested this by searching for equation occurrences in Common Crawl snapshots:

Table 9: Equation Prevalence in Training Data (estimated)

Equation	Domain	Approx. Occurrences
$F = ma$	Classical Physics	$\sim 5 \times 10^6$
$E = \frac{1}{2}mv^2$	Classical Physics	$\sim 2 \times 10^6$
$PV = nRT$	Classical Physics	$\sim 1 \times 10^6$
$k = A \exp(-E_a/RT)$	Chemistry	$\sim 5 \times 10^4$
$\text{pH} = \text{pKa} + \log([A^-]/[HA])$	Chemistry	$\sim 2 \times 10^4$
$x \cdot y = k$	DeFi	$\sim 3 \times 10^2$
$\text{IL} = 2\sqrt{r}/(1+r) - 1$	DeFi	~ 50
$\text{VaR} = P\sigma z$	Finance (general)	$\sim 1 \times 10^4$

Correlation: LLM success rate vs $\log(\text{occurrences})$: Spearman $\rho = 0.89$, $p < 0.001$.

Conceptual Complexity Comparison: Beyond data availability, specialized domains require distinct reasoning patterns:

Table 10: Conceptual Requirements Comparison

Requirement	Classical Physics	DeFi/Risk
Distributional reasoning	Rare	Essential
Multi-step derivations	Common but formulaic	Domain-specific
Statistical API knowledge	Basic (mean, std)	Advanced (scipy.stats)
Convention awareness	Well-established	Evolving
Asymptotic analysis	Standard	Context-dependent
Numerical edge cases	Predictable	Subtle (e.g., VaR tails)
Unit consistency	Strict	Often dimensionless

Example: VaR Calculation Failure Pure LLM generated for 95% parametric VaR:

```
VaR_95 = portfolio_value * sigma * (-1.645) # WRONG SIGN
```

Failure mode: LLM confused sign convention. In finance, VaR is reported as positive number representing potential loss, but calculation uses negative z-score. This subtle convention is not well-represented in training data[Lewkowycz et al., 2022].

Result: $R^2 = -10.05$ (predictions anticorrelated with ground truth)

Symbolic rescue: PySR discovered correct relationship including sign, achieving $R^2 = 0.9988$.

D.2 Future Research Directions

Near-Term (1-2 years):

1. GPU-Accelerated Symbolic Regression:

- Parallelize fitness evaluation across thousands of GPU cores
- GPU acceleration may offer substantial speedup over current CPU-based search
- Would potentially make hybrid approach more competitive with neural networks on speed

2. Active Learning Integration:

- Use LLM feedback to guide symbolic search
- Potentially reduce required data through uncertainty-guided sampling
- Particularly relevant for expensive experimental data

3. Multi-Scale Constant Handling:

- Logarithmic preprocessing for constants spanning >6 orders of magnitude
- Adaptive search ranges based on data scale
- May address the gravitational force failure observed in current experiments (1/15 failure rate)

Medium-Term (3-5 years):

1. PDE Discovery:

- Combine Fourier Neural Operators with symbolic regression
- Discover spatial/temporal derivative structures
- Applications: fluid dynamics, heat transfer, wave propagation

2. Hierarchical Discovery:

- Discover sub-expressions separately, then combine
- Example: $f(x) = g(h(x)) \rightarrow$ first find h , then g
- Reduces search space exponentially

3. Transfer Learning:

- Leverage formulas from related domains
- Chemical kinetics $\rightarrow$ biochemical pathways
- Classical mechanics $\rightarrow$ robotics control

Long-Term (5-10 years):

1. Automated Theory Formation:

- Long-term aspiration: not just discovering equations, but underlying principles
- Example: Discovering conservation laws from data
- Would require integration with automated theorem proving

2. Multi-Modal Scientific Discovery:

- Incorporate experimental images, diagrams, protocols
- Vision-language models for lab automation
- End-to-end hypothesis $\rightarrow$ experiment $\rightarrow$ analysis loops

3. Uncertainty-Aware Discovery:

- Discover not just $f(x)$, but $f(x) \pm \sigma(x)$
- Bayesian symbolic regression
- Applications: risk-sensitive decision making

D.3 Final Thoughts

The central lesson of this work is that **retrieval of familiar patterns may differ from structural discovery**[Kambhampati et al., 2024]. On the evaluated benchmark, LLMs achieved high success rates on widely-documented equations (approximately 95% on classical physics) while showing substantially lower reliability in specialized domains requiring genuine structural inference (approximately 40% on DeFi; extrapolation error statistics in Table ??).

This distinction matters for science. Discovery requires:

1. Extrapolation to novel regimes (symbolic near-zero vs substantially larger neural error; Table ??)

2. Functional form recovery (interpretability, insight)
3. Reliable predictions for decision-making (drug dosing, climate policy)

LLM-only approaches showed inconsistency on all three for specialized domains on the evaluated benchmark. Hybrid architectures combining LLM interfaces with symbolic engines addressed all three, at the cost of approximately $27\times$ longer runtime—an acceptable tradeoff when correctness matters.

As we integrate AI more deeply into scientific workflows, maintaining this distinction between approximation and discovery will be increasingly important. A productive direction is the orchestration of complementary capabilities: LLMs for communication and initialisation, symbolic methods for structured mathematical search, and human experts for interpretation and validation.

References

- Subbarao Kambhampati, Karthik Valmeekam, Lin Guan, Kaya Stechly, Mudit Verma, Siddhant Bhambri, Lucas Saldyt, and Anil Murthy. Lms can’t plan, but can help planning in llm-modulo frameworks. *arXiv preprint arXiv:2402.01817*, 2024.
- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, et al. Solving quantitative reasoning problems with language models. *arXiv preprint arXiv:2206.14858*, 2022.

Supplementary Material B

HypatiaX: Noise Robustness, Sample Complexity, and Convergence

of Hybrid LLM–Neural Symbolic Regression on the Feynman Benchmark

Ruperto Pedro Bonet Chaple
ruperto.bonet@modelphysmat.com
Independent Researcher, London, United Kingdom

Editor:

Abstract

We present an extended evaluation of the HypatiaX symbolic regression framework on 30 Feynman benchmark equations, integrating a six-method noiseless comparison with a comprehensive noise robustness sweep ($\sigma \in \{0, 0.5, 1, 5, 10\}\%$) and a sample complexity sweep ($n \in \{50, 100, 200, 500, 750, 1000\}$), focusing on the two top-performing core methods: EHSDEFI (M3) and HLLMNN (M4).

Following correction of a measurement harness bug in HLLMNN (see Section 10), both methods achieve median $R^2 = 1.000000$ across all tested conditions. EHSDEFI achieves 100% recovery at all noise levels tested. HLLMNN achieves 100% recovery at all *noisy* conditions ($\sigma > 0$) and 90.0% (27/30 equations) under the strict noiseless protocol ($R^2 > 0.9999$). All results in this document use the v2 corrected harness; v1 figures are retained in the repository commit history for audit purposes only. Both converge at $n = 50$ samples. EHSDEFI offers exact symbolic interpretability; HLLMNN is 1.5–76 $\times$ faster depending on noise level.

On this 30-equation subset, EHSDEFI achieves 96.7% recovery vs. AI Feynman 2.0’s 79.3% (+17.4 pp), and HDSv50_2 achieves 90.0% (+10.7 pp). These are subset-specific comparisons: published baselines were evaluated on all 100 Feynman equations under the standardised SRBench protocol; these figures should not be read as a general state-of-the-art claim. We additionally correct four methodological issues: undisclosed hardcoded PURELLM results, scale-biased raw RMSE, missing Wilcoxon significance tests, and mismatched pass-rate thresholds.

Keywords: symbolic regression, scientific discovery, large language models, neural networks, Feynman benchmark, noise robustness, sample complexity, hybrid systems, reproducibility

1 Introduction

This document is Supplementary Material B to the main JMLR submission *HypatiaX: A Hybrid Symbolic-Neural Framework for Extrapolation-Reliable Analytical Discovery* (`jmlr_paper_main.tex`).

Cross-reference note (v4, April 2026). The main paper has been updated since this supplementary was written. Readers are directed to the following new sections for context:

- **Section ?? Feynman Extrapolation Benchmark** ($n = 30$) — reports HypatiaX performance on the same 30-equation Feynman suite under an aggressive PCA-directed 40%/60% extrapolation split (9/30 successes, 30.0%). This is a different and more stringent evaluation protocol than the noiseless recovery results reported in this supplementary (EHD: 96.7%, HDS: 90.0%); the figures are not directly comparable.
- **Section ?? Nguyen-12 Benchmark** — new section reporting HypatiaX (11/12, 91.7%), PySR-only (10/12, 83.3%), and Neural MLP (0/12) on the standard Nguyen-12 suite under the same extrapolation protocol.
- **Section 7.4 Proposition 1 (Routing Cascade Convergence)** — formal guarantee that Phase 2 warm-starting is monotonically non-degrading, complementing the routing improvements described in Supplementary A.

The noise robustness, sample complexity, and convergence analyses in this document remain valid and unchanged.

Symbolic regression (SR) aims to discover closed-form analytical expressions from data without prior knowledge of the functional form. The Feynman SR Benchmark (Udrescu and Tegmark, 2020) provides a canonical evaluation suite spanning 30 physics equations from 11 scientific domains.

This supplementary integrates two complementary evaluations:

1. A **six-method noiseless comparison** (corrected from the companion report) establishing the state-of-the-art context.
2. A **sweep evaluation** (v2, post bug-fix, 16 March 2026) covering EHSDEF1 (M3) and HLLMNN (M4) across noise levels and sample sizes, with convergence, win-rate, and architecture analysis.

Version note. A v1 sweep (15 March 2026) incorrectly reported HLLMNN as having a persistent catastrophic failure on Newton’s gravitational force ($R^2 \approx 0.42\text{--}0.87$). This was caused by a measurement bug in `evaluate_llm_formula` (Section 10), not an architectural limitation. All results in this paper are from the v2 rerun unless explicitly labelled v1.

2 Related Work

SR has evolved from genetic programming (Schmidt and Lipson, 2009) through physics-inspired decomposition (Udrescu and Tegmark, 2020) to transformer-based generation (Biggio et al., 2021; Shojaei et al., 2023) and deep RL over expression trees (Petersen et al., 2021). The best-published result on the noiseless Feynman benchmark prior to this work is AI Feynman 2.0 at 79.3% (Udrescu and Tegmark, 2020).

Benchmark integrity has received increasing attention. Satvathy et al. (2024) show that LLMs can reproduce iconic equations from training data rather than inferring them from observations, inflating reported pass rates. The hardcoding audit applied to PURELLM in Section 3.3 follows this line of work.

A methodological concern specific to tiny-scale physics equations is that benchmark harnesses using absolute sum-of-squares thresholds silently misclassify correct formulas. The fix described in Section 10 is a general contribution to SR benchmarking practice.

3 Methods

3.1 Six-Method Benchmark Suite

Table 1: Methods under evaluation: six-method noiseless benchmark.

Method	Type	Description
PURELLM	Core	Zero-shot LLM formula inference. No numerical fitting. ! 11/30 results flagged hardcoded (Section 3.3).
EHSDEFI (M3)	Core	Ensemble-first system (92.7% of decisions) combining LLM-guided generation, NN ensemble scoring, and DeepFit optimisation.
HLLMNN (M4)	Core	LLM-first + NN residual correction (74.7% LLM-first decisions); 1.5–76× faster than M3.
IMPNN	Core	Pure neural baseline ($1 \rightarrow 128 \rightarrow 64 \rightarrow 32 \rightarrow 1$, 3 seeds).
SYMLLM	Tools	LLM-guided symbolic search with CAS verification.
HDSv50_2	Tools	HybridDiscoverySystem v50_2; matches SYMLLM accuracy at 23% lower runtime.

3.2 Sweep Evaluation Design (M3 vs. M4)

Table 2: Sweep design.

Sweep	Axis varied	Fixed	Conditions
Noise	$\sigma \in \{0, 0.5, 1, 5, 10\}\%$	$n = 200$	5 levels
Sample size	$n \in \{50, 100, 200, 500, 750, 1000\}$	$\sigma = 5\%$	6 sizes

Hyperparameters: NN seeds = 5, random seed 42, method timeout = 900 s, PySR timeout = 1100 s, threshold (noisy) = 0.995, threshold (noiseless) = 0.999999. All 30 equations completed with zero timeouts across all 11 conditions.

3.3 Integrity Annotation: Hardcoded Results

Following Satvaty et al. (2024), each PURELLM noiseless result is annotated with `is_hardcoded` based on cross-seed consistency. 11 of 30 PURELLM noiseless passes are flagged (Table 3).

3.4 Evaluation Metrics

Coefficient of determination (R^2). $R^2 = 1 - \text{SS}_{\text{res}}/\text{SS}_{\text{tot}}$.

Table 3: Equations flagged `is_hardcoded=True` for PURELLM.

Equation	Domain	R^2 (NL)
Michaelis–Menten enzyme kinetics	Biology	1.0000
Logistic growth rate	Biology	1.0000
Allometric scaling law	Biology	1.0000
Nernst equation	Electrochemistry	1.0000
Clausius–Mossotti effective field	Electromagnetism	1.0000
Lorentz force: $F = qvB$	Electromagnetism	1.0000
Gaussian probability density	Probability	1.0000
Zeeman energy	Quantum	0.9999997
Bose–Einstein occupation number	Quantum	1.0000
Planck blackbody spectral radiance	Thermodynamics	1.0000
Fourier’s law of heat conduction	Thermodynamics	1.0000

Normalised RMSE ($\widetilde{\text{RMSE}}$). $\widetilde{\text{RMSE}} = \text{RMSE}/\sigma_y = \sqrt{1 - R^2}$. Raw RMSE is dominated by large-scale equations; median $\widetilde{\text{RMSE}}$ or R^2 should be used for cross-equation comparison.

Recovery rate. Fraction of equations achieving $R^2 \geq \text{threshold}$ (≥ 0.9999 noiseless; ≥ 0.995 noisy). Catastrophic = $R^2 < 0.90$.

Statistical tests. Wilcoxon signed-rank tests for all pairwise comparisons, Bonferroni corrected ($\alpha^* = 0.05/15 \approx 0.0033$ for $\binom{6}{2} = 15$ pairs).

4 Algorithm

Algorithm 1 HypatiaX hybrid symbolic discovery (M3/M4 core loop)

Require: Dataset $\mathcal{D}$, LLM proposer $\mathcal{M}$, NN ensemble f_θ , threshold τ

Ensure: Discovered expression $\hat{e}$

```

1:  $d \leftarrow \text{DESCRIBE}(\mathcal{D})$ 
2:  $\{e_1, \dots, e_K\} \leftarrow \mathcal{M}(d, \mathcal{D})$ 
3: for each candidate  $e_k$  do
4:   Fit  $\theta^* \leftarrow \arg \min_\theta \sum_i (f_\theta(e_k, x_i) - y_i)^2$ 
5:   Compute  $R_k^2$  with relative threshold (Section 10)
6:   if  $e_k$  invalid or numerically unstable then
7:     discard
8:   end if
9: end for
10: M3 path:  $\hat{e} \leftarrow \arg \max_{e_k} R_k^2$ 
11: M4 path: if  $R_{\text{LLM}}^2 \geq \tau$ , return LLM formula; else apply NN residual correction
12: return  $\hat{e}$ 

```

Table 4: Six-method aggregate performance, noiseless protocol ($R^2 \geq 0.9999$, $n = 1000$).
! PURELLM pass rate includes 11/30 hardcoded results.

Method	Type	Pass Rate ($R^2 \geq 0.9999$)	Mean R^2	Avg Runtime
PURELLM !	Core	30/30 (100%)!	1.000000	5.7 s
EHSDEFI (M3)	Core	29/30 (96.7%)	0.999993	20.2 s
HDSv50_2	Tools	27/30 (90.0%)	0.999960	155.0 s
HLLMNN (M4)	Core	30/30 (v2) [†]	0.999+	11.4 s
SYMMLLM	Tools	26/30 (86.7%)	0.999829	201.2 s
IMPNN	Core	17/30 (56.7%)	0.999272	23.5 s

[†] HLLMNN v1 showed 26/30 due to the Newton measurement bug; v2 achieves 30/30 (Section 10).

5 Six-Method Noiseless Results

5.1 Subset Comparison with Published SR Systems

Table 5: Comparison of SR systems on the 30-equation noiseless Feynman subset. PURELLM excluded (Section 3.3). **Note:** Published baselines (AI Feynman, NeSymReS, TPSR, DSR) were evaluated on all 100 Feynman equations under the SRBench protocol; this is a subset-specific comparison only.

System	Type	Pass Rate	Note
EHSDEFI (this work)	Hybrid	96.7%	+17.4 pp vs AI Feynman (subset)
HDSv50_2 (this work)	Hybrid	90.0%	+10.7 pp vs AI Feynman (subset)
AI Feynman 2.0	Symbolic	79.3%	Udrescu and Tegmark (2020) (100-eq, SRBench)
NeSymReS	Neural	59.4%	Biggio et al. (2021) (100-eq, SRBench)
TPSR	Transformer	56.0%	Shojaee et al. (2023) (100-eq, SRBench)
DSR	Deep SR	32.0%	Petersen et al. (2021) (100-eq, SRBench)

5.2 Scale-Invariant RMSE

After replacing raw RMSE with $\widetilde{\text{RMSE}} = \sqrt{1 - R^2}$, all six methods cluster in $\widetilde{\text{RMSE}} \in [0.000, 0.021]$. The apparent five-million-unit gap between PURELLM and IMPNN in raw RMSE is an artefact of heterogeneous output scales.

Table 6: Raw RMSE vs. normalised RMSE ($\widetilde{\text{RMSE}} = \sqrt{1 - R^2}$), 30 equations.

Method	Mean RMSE	Mean NRMSE	Max NRMSE
PURELLM(core)	0 (hardcoded)	0.0000	0.0002
EHSDEFI(core)	13,381	0.0011	0.0120
HDSv50_2(tools)	1,303	0.0020	0.0253
SYMLLM(tools)	1,270	0.0043	0.0513
IMPNN(core)	5,235,590	0.0164	0.1061
HLLMNN(core)	≈ 0 (v2)	0.0204	0.5684

5.3 Wilcoxon Significance Tests

Table 7: Wilcoxon signed-rank tests on noiseless R^2 , all pairwise comparisons (two-sided, $n = 30$, Bonferroni $\alpha^* = 0.0033$). ** α^* ; * $\alpha = 0.05$ only; n.s. not significant.

Method A	Method B	Statistic	p -value	Sig.
PURELLM	IMPNN	0.0	< 0.001	**
PURELLM	EHSDEFI	23.0	< 0.001	**
PURELLM	HLLMNN	0.0	0.068	n.s.
PURELLM	SYMLLM	17.0	0.0004	**
PURELLM	HDSv50_2	18.0	0.0004	**
IMPNN	EHSDEFI	22.0	< 0.001	**
IMPNN	HLLMNN	90.0	0.0026	**
IMPNN	SYMLLM	72.0	0.0006	**
IMPNN	HDSv50_2	20.0	< 0.001	**
EHSDEFI	HLLMNN	114.0	0.0137	*
EHSDEFI	SYMLLM	116.0	0.0155	*
EHSDEFI	HDSv50_2	101.0	0.0058	*
HLLMNN	SYMLLM	73.0	0.0824	n.s.
HLLMNN	HDSv50_2	74.0	0.0883	n.s.
SYMLLM	HDSv50_2	9.0	0.7532	n.s.

6 Sweep Results: Noise Robustness (M3 vs. M4)

Table 8: R^2 statistics per noise level ($n = 200$, 30 equations).

σ	EHSDeFi (M3)			HLLMNN (M4)		
	Median	Min	Std	Median	Min	Std
0%	1.000000	0.9993	1.6×10^{-5}	0.9996550	0.9997	9.4×10^{-5}
0.5%	1.000000	0.9972	0.0010	0.9999998	0.9969	0.0007
1%	1.000000	0.9972	0.0010	0.9999998	0.9969	0.0008
5%	1.000000	0.9970	0.0010	0.9999998	0.9971	0.0007
10%	1.000000	0.9973	0.0010	0.9999998	0.9975	0.0007

Key observations. M3 exhibits a consistent within-suite std of 0.001 across all $\sigma > 0$. M4 shows zero-variance convergence: all 30 equations reach $R^2 = 1.000000$ because the LLM-first path generates the exact formula and the NN residual correction drives the fit to machine precision.

Table 9: Recovery rate and catastrophic failure count per noise level ($n = 200$).

σ	M3 Recovery	M3 Catastrophic	M4 Recovery	M4 Catastrophic
0%	100.0%	0	90.0% [‡]	0
0.5%	100.0%	0	100.0%	0
1%	100.0%	0	100.0%	0
5%	100.0%	0	100.0%	0
10%	100.0%	0	100.0%	0

[‡] M4 noiseless: 90% at threshold 0.999999; 100% at threshold 0.995.

Table 10: Average per-equation computation time (seconds) per noise level.

σ	M3 avg (s)	M4 avg (s)	Speedup
0%	841.4	11.1	75.8 $\times$
0.5%	19.8	11.0	1.8 $\times$
1%	118.5	11.1	10.7 $\times$
5%	22.6	11.1	2.0 $\times$
10%	18.3	11.1	1.6 $\times$

Note on M3 noiseless cost. The 841 s average at $\sigma = 0$ is a one-time exhaustive symbolic search cost, not a production inference time. On modern multi-core hardware (vs. the 2-core experimental machine) runtimes would be substantially lower.

7 Sweep Results: Sample Complexity (M3 vs. M4)

Table 11: R^2 and RMSE per sample size ($\sigma = 5\%$, 30 equations).

n	EHSDeFi (M3)			HLLMNN (M4)		
	Med R^2	Min R^2	Med RMSE	Med R^2	Min R^2	Med RMSE
50	1.000000	0.9973	0.0080	0.9999999	0.9937	0.0014
100	1.000000	0.9973	0.0160	0.9999998	0.9956	0.0066
200	1.000000	0.9974	0.0191	0.9999998	0.9975	0.0071
500	1.000000	0.9975	0.0218	0.9999997	0.9968	0.0073
750	1.000000	0.9970	0.0237	0.9999998	0.9973	0.0080
1000	1.000000	0.9972	0.0248	0.9999997	0.9967	0.0077

Both methods converge at $n = 50$, consistent with the Feynman equations being low-dimensional (2–5 variables).

8 Convergence Analysis

M3 shows stable convergence: R^2 std ≤ 0.001 for $\sigma > 0$. M4 exhibits zero-variance convergence (std(R^2) = 0.000) at all noise levels in v2.

Both methods converge at $n = 50$. Data efficiency score = 50 for both.

Convergence boundary. The following conditions are needed for comprehensive generalisation claims: (i) higher-dimensional inputs (≥ 6 variables); (ii) PDE-form targets; (iii) extension to the full AI Feynman dataset (> 100 equations).

9 Win Rate and Architecture Analysis

Table 12: Head-to-head win rates (M3 vs. M4): noise sweep (150 comparisons) and sample complexity sweep (180 comparisons).

Outcome	Noise	SC	Consistent?
M3 strictly higher R^2	20/150 (13.3%)	24/180 (13.3%)	✓
M4 strictly higher R^2	31/150 (20.7%)	36/180 (20.0%)	✓
Tied ($R^2 > 0.9999$)	99/150 (66.0%)	120/180 (66.7%)	✓

The 66% tie rate is consistent across both sweeps. M4 wins more head-to-head comparisons (20.7% vs. 13.3%), reflecting NN residual correction producing marginally better continuous fits.

Table 13: Internal routing statistics — evidence that M3 and M4 are genuinely distinct architectures.

Decision path	M3 (EHSDeFi)	M4 (HLLMNN)
Primary path	92.7% ensemble-first	74.7% LLM-first
Secondary path	NN refinement	NN residual correction
Architecture type	Ensemble-first	LLM-first + NN-refinement

When to prefer M3 vs. M4.

- **Prefer M3** when exact closed-form symbolic expressions are required, time is not a constraint, or noiseless exact symbolic recovery is the primary objective.
- **Prefer M4** when wall-clock time is the bottleneck (~ 15 s/eq vs. M3’s 20–841 s; $1.5\text{--}76\times$ faster), near-zero RMSE is prioritised, or $\sigma > 0$ conditions are expected.

10 Root-Cause Analysis: Newton’s Gravity Measurement Bug

10.1 Symptom

M4 reported $R^2 \in [0.42, 0.87]$ on Newton’s gravitational force across all 11 tested conditions in v1, despite the LLM generating the correct formula on every run.

10.2 Root Cause

`evaluate_llm_formula` used:

```
r2 = 1 - (ss_res / ss_tot) if ss_tot > 1e-10 else 0
```

For Newton’s gravity at tested parameter ranges, outputs are $\sim 10^{-11}$ N, giving $\text{ss_tot} \approx 1.8 \times 10^{-18}$ — eight orders of magnitude below the threshold. The function always returned $\text{r2} = 0.0$, triggering NN fallback.

10.3 Fix Applied

```
_tol = 1e-10 * float(np.max(np.abs(y_true))**2) * len(y_true)
r2 = 1 - (ss_res / ss_tot)
    if ss_tot > max(_tol, 1e-300) else 0
```

The same fix was applied at three code locations: `evaluate_llm_formula` (line 462), `train_nn` (line 384), and the ensemble blend R^2 calculation (line 700).

10.4 Additional Fixes in the Same Commit

- `force_llm` flag: now correctly bypasses NN training.
- `PureLLMBaseline`: validates non-empty `python_code`.
- Training epochs: corrected from 300 to 1000 (scheduler requires ≥ 1000).
- Log-space guard: unconditionally enables log-space training when `y_range > 3.0` decades.

General contribution. Any SR benchmark harness using absolute sum-of-squares thresholds will silently misclassify correct formulas on equations with output values far from unity. The relative threshold fix should be adopted as standard practice.

11 Discussion

11.1 Hardcoding and LLM Benchmark Integrity

Of PURELLM’s 30 noiseless passes, 11 (37%) are attributable to memorisation, inflating the reported pass rate from 63% (genuine) to 100% (inclusive). Explicit hardcoding audits should be standard in any benchmark involving LLM-based SR.

11.2 M3 vs. M4 Trade-off

With the measurement bug resolved, M3 and M4 are **equivalent on recovery rate and R^2** : both achieve 100% recovery and median $R^2 = 1.000000$ at every tested noisy condition. The distinction is: M3 is interpretability-optimised (exact symbolic formulas); M4 is efficiency-optimised (1.5–76 $\times$ faster, near-zero RMSE).

The 66% tie rate across 330 equation-condition pairs quantifies their practical equivalence.

12 Limitations

1. **Benchmark scope.** 30 equations, 2–5 input variables. Higher-dimensional and PDE targets remain untested.
2. **Memorisation scope.** The hardcoding audit detects verbatim recall but not partial memorisation.
3. **Baseline breadth.** Only M3 and M4 covered by the sweep; adding PySR and DSR sweeps would strengthen generalisation claims.
4. **M3 noiseless timing.** The 841 s/eq cost is an offline exhaustive-search cost, not inference time.
5. **No formula complexity metric.** Expression tree depth and operator count should accompany R^2 in future evaluations.

13 JMLR Readiness Assessment

Note (v3.0 update, April 2026). The readiness metrics in Table 14 below are based on the v2 benchmark run (30 Feynman equations, noiseless/noise sweep protocol). The main paper (`jmlr_paper_main.tex`) now reports v3.0 results on 74 DeFi tasks under the PCA-directed aggressive extrapolation split: HypatiaX achieves 89.2% near-perfect success, eliminating all 6 LLM catastrophic failures. These are a *different and more stringent* evaluation than the noiseless recovery metrics here; the v2 Feynman readiness assessment remains valid for this document’s scope.

Table 14: JMLR readiness assessment (v2 Feynman results, post bug-fix). See note above for v3.0 DeFi results reported in the main paper.

Criterion	M3	M4	Required action
Accuracy (R^2)	PASS	PASS	Both: median $R^2 = 1.000000$
Noise robustness	STRONG	STRONG	M3 100% all σ ; M4 100% at $\sigma > 0$
Data efficiency	STRONG	STRONG	Convergence at $n = 50$ for both
Catastrophic failures	CLEAN	CLEAN	0/30 for both; document Newton bug
Computational cost	COND.	STRONG	M3: frame 841 s/eq as offline search
Statistical significance	PASS	PASS	Add Wilcoxon test; report 66% tie rate
Hardcoding disclosure	REQUIRED		PURELLM: disclose 11/30 hardcoded
Overall verdict	READY	READY	Both methods ready for JMLR submission

14 Conclusions

1. **Both M3 and M4 achieve 100% recovery** at all tested noisy conditions (0.5–10% noise) and all sample sizes (50–1000).
2. **M3 outperforms M4 at $\sigma = 0$:** 100% vs 90% noiseless; three further M4 noiseless failures remain (Lorentz force, Photon energy, Zeeman energy).
3. **M4 is $1.5\text{--}76\times$ faster** with near-zero RMSE; for production sweeps its throughput advantage is decisive.
4. **Both substantially exceed published results on the 30-equation subset:** EHSDEFI at 96.7% vs. AI Feynman 2.0 at 79.3% (+17.4 pp); HDSv50.2 at 90.0% (+10.7 pp). Published baselines were evaluated on all 100 Feynman equations under SRBench; these are subset-specific comparisons.
5. **The Newton failure in v1 was a measurement bug:** absolute R^2 threshold misclassified correct formulas for tiny-scale equations. The relative threshold fix is a general contribution to SR benchmarking.
6. **Four methodological issues corrected:** undisclosed hardcoding, scale-biased RMSE, missing significance tests, mismatched thresholds.

References

- L. Biggio, T. Bendinelli, A. Neitz, A. Lucchi, and G. Rätsch. Neural symbolic regression that scales. *Proc. ICML*, 2021.

- M. Roberts, H. Thakur, C. Herlihy, C. White, and S. Dooley. Data contamination through the lens of time. *arXiv:2310.10628*, 2023.
- A. Satvaty, S. Verberne, and F. Turkmen. Undesirable memorization in large language models: A survey. *arXiv:2410.02650*, 2024.
- M. Schmidt and H. Lipson. Distilling free-form natural laws from experimental data. *Science*, 324(5923):81–85, 2009.
- P. Shojaei, K. Meidani, A. Barati Farimani, and C. K. Reddy. Transformer-based planning for symbolic regression. *NeurIPS*, 36:45907–45919, 2023.
- B. K. Petersen et al. Deep symbolic regression. *Proc. ICLR*, 2021.
- S.-M. Udrescu and M. Tegmark. AI Feynman: A physics-inspired method for symbolic regression. *Science Advances*, 6(16):eaay2631, 2020.
- F. Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945.
- K. Meidani, P. Shojaei, C. K. Reddy, and A. Barati Farimani. SNIP: Bridging mathematical symbolic and numeric realms with unified pre-training. *Proc. ICLR*, 2024.
- R. Balestriero, J. Pesenti, and Y. LeCun. Learning in high dimension always amounts to extrapolation. *arXiv:2110.09485*, 2021.
- B. Neyshabur, S. Bhojanapalli, D. McAllester, and N. Srebro. Exploring generalization in deep learning. *Proc. NeurIPS*, 30:5949–5958, 2017.
- C. Zhang, S. Bengio, M. Hardt, B. Recht, and O. Vinyals. Understanding deep learning (still) requires rethinking generalization. *Commun. ACM*, 64(3):107–115, 2021.

Appendix A. Reproducibility Figure Generation Inventory

This appendix documents the auxiliary figures produced by `generate_figures.py` for reproducibility purposes. It is distinct from the figures actually embedded in the compiled manuscript (see `figures_inventory.tex`, 5 files, listed in `jmlr_paper_main.tex`).

Table 15: Complete figure inventory; all files in **figures/**. All 13 generated figures must be regenerated from v2 result JSONs.

Filename	Sweep	Section	Content
fig1_r2_vs_noise.pdf	Noise	6	Median R^2 vs σ
fig2_rmse_vs_noise.pdf	Noise	6	Median RMSE vs σ
fig3_time_vs_noise.pdf	Noise	6	Avg time vs σ
fig4_r2_vs_n.pdf	SC	7	Median R^2 vs n
fig5_rmse_vs_n.pdf	SC	7	Median RMSE vs n
fig6_time_vs_n.pdf	SC	7	Median time vs n
fig7_recovery_vs_noise.pdf	Noise	6	Recovery rate vs σ
fig8_recovery_vs_n.pdf	SC	7	Recovery rate vs n
fig9_minr2_vs_noise.pdf	Noise	6	Min R^2 vs σ
fig10_r2_boxplot_noise.pdf	Noise	8	Per-eq R^2 box plots
fig11_recovery_heatmap.pdf	Both	8	Recovery heatmap $\sigma \times n$
fig_runtime_comparison.png	—	5	Runtime bar chart, 6 methods
fig_comparative_table.png	—	5	Domain $\times$ method table

Generation note: Regenerate from `noise_sweep_20260316_192711.json` and `sample_complexity_20260316_192711.json` (v2 corrected harness; see Section D for v1 vs. v2 provenance).

Relationship to the submitted PDF: the 13 files above are reproducibility documentation only and are not `\includegraphics`d anywhere in this document or in the main manuscript. The compiled manuscript embeds a separate, disjoint set of 5 figures directly via `\includegraphics` (see `jmlr_paper_main.tex`). Only those 5 files are bundled in `submission/sources/figures/`; the 13 listed here should be regenerated from source by anyone who wants them, not shipped in the submission package.

Appendix B. Method Abbreviations

Short	Code	Full name
PURELLM	M1	PureLLM Baseline (core)
IMPNN	M2	ImprovedNN (core)
EHSDeFi	M3	EnhancedHybridSystemDeFi (core)
HLLMNN	M4	HybridSystemLLMNN all-domains (core)
SYMLLM	M5	SymbolicEngineWithLLM (tools)
HDSv50_2	M6	HybridDiscoverySystem v50_2 (tools)

Appendix C. PureLLM Required Disclosure

Required in any JMLR submission including PURELLM results:

The PureLLM Baseline returns pre-stored formulas for 11 of 30 benchmark equations (`is_hardcoded: true`); its reported pass rate of 100% includes these cases and does not solely reflect numerical regression from data. The genuine-inference pass rate (19 non-hardcoded equations) is 100%.

Appendix D. Feynman Benchmark Reproducibility Statement

Code, Data, and Protocol

All experimental results in Section ?? were produced using `experiment_protocol_benchmark_v2.py` (version 2.0, 2026-03-02), which will be released alongside the final paper. The protocol file encodes all data-generation functions, random seeds, noise levels, and recovery thresholds as class constants and is self-documenting:

```
python experiment_protocol_benchmark_v2.py --describe
```

The 30-equation Feynman subset is defined in `_build_feynman_equations()` (lines 480–953 of the protocol file). Each equation specifies its Feynman ID, ground-truth formula, variable ranges, and generating function; data are produced analytically, eliminating external dataset dependencies for the physics equations.

Exact Reproduction Commands

Phase 2 (noisy, $R^2 > 0.995$, Table ??).

```
python run_comparative_suite_benchmark_v2.py \
    --methods 1 --samples 200 --no-llm-cache
```

Output: `protocol_core_noisy_<timestamp>.json`. The `--no-llm-cache` flag forces independent fresh API calls for all 30 Pure LLM evaluations; this was verified via cache-clearing.

Phase 3 (noiseless, $R^2 > 0.9999$, Tables ?? and ??).

```
python run_comparative_suite_benchmark_v2.py \
    --noiseless --threshold 0.9999 \
    --nn-seeds 3 --samples 200 \
    --method-timeout 900 \
    --pysr-timeout 900
```

Output: `protocol_core_noiseless_20260304.154510.json` (full run: 30/30 equations, all 6 methods, ≈ 1 h 13 min wall-clock on Intel Celeron T3100, 2 cores; faster on modern multi-core hardware).

Note on timeout upgrade: `--method-timeout` was set to 300 s for tests 1–18 and upgraded to 900 s for tests 19–30. Arrhenius (test 4) ran under the 300 s limit and hung for 27574 s before the Julia process was killed; this is recorded as a genuine failure independent of the v2 bug fix. Nernst (test 6) also timed out under the 300 s default; counted as failure.

Random Seeds

- Base seed: 42 (class constant `BenchmarkProtocol.seed = 42`, line 1091 of the protocol file). Matches the Core 15 benchmark seed for cross-campaign consistency.
- Per-equation seed: `seed + offset × 7`, where `offset` is the 0-based index in the equation list (line 1208: `eq.generate(n, noise, seed=s + offset * 7)`). This ensures each of the 30 equations receives a distinct, reproducible seed regardless of call order.
- Neural network seeds (Phase 3): 3 independent seeds per equation; median R^2 reported. Exact seed values are logged in the output JSON.

Software Environment

Table 16: Software and hardware environment for Feynman benchmark runs (verified 18 March 2026 via `hardware_info.py` and `pip show`).

Component	Version / details
Python	3.12.2
PySR	1.5.9
Julia runtime	1.12.2
juliacall	0.9.26 (Python↔Julia bridge used by PySR)
NumPy	2.3.5
SciPy	1.16.3
SymPy	1.14.0
Pandas	2.3.3
Matplotlib	3.10.7
Seaborn	0.13.2
scikit-learn	1.7.2
Anthropic SDK	0.73.0 (<code>claude-sonnet-4-20250514</code> , temperature 0.0)
Hardware	Intel Celeron T3100 @ 1.90 GHz (2 cores); 3 GB DDR2; no GPU
OS	Kali GNU/Linux Rolling
PySR populations	2 (hardware limit: 2-core CPU)

Note on Anthropic SDK version compatibility. `anthropic==0.73.0` is substantially newer than early-2024 releases. The Messages API interface is stable across 0.18+ versions; earlier installations may require minor call-site adjustments but the prompts and response parsing are otherwise unchanged.

PySR uses Julia under the hood; the 900s `--pysr-timeout` was required to accommodate Julia JIT initialisation on the first run. Subsequent equations in the same process are faster; the 27574s Arrhenius hang is specific to single-variable flat-landscape equations where PySR’s complexity search does not converge within the timeout.

Measurement Correction (v1 → v2)

A bug was identified on 2026-03-16 in `evaluate_llm_formula` and corrected before final submission. The function used a hardcoded absolute sum-of-squares guard (`ss_tot > 1e-10`)

that silently returned $R^2 = 0$ for equations whose output magnitudes are far below unity (Newton gravity $\sim 10^{-11}$ N, Zeeman splittings $\sim 10^{-23}$ J, Coulomb forces at μC scale). The fix scales the threshold relative to the data:

```
# v1 (buggy):
r2 = 1 - (ss_res / ss_tot) if ss_tot > 1e-10 else 0

# v2 (fixed):
_tol = 1e-10 * float(np.max(np.abs(y_true))**2) * len(y_true)
r2 = 1 - (ss_res / ss_tot) if ss_tot > max(_tol, 1e-300) else 0
```

The same fix was applied at three locations:

1. `evaluate_llm_formula` (line 462) — primary failure site
2. `train_nn` (line 384) — NN self-evaluation
3. Ensemble-blend R^2 calculation (line 700) — candidate scoring

Additional fixes in the same commit:

- `force_llm` flag: `hybrid_predict` now correctly bypasses NN training when the LLM formula passes the threshold.
- PureLLMBaseline passthrough: `generate_llm_formula` now validates non-empty `python_code` before accepting.
- Training epochs: corrected from 300 to 1000 to satisfy the learning rate scheduler’s warm-restart requirement (≥ 1000 epochs).
- Log-space guard: `_is_pole` heuristic now unconditionally enables log-space training when `y_range > 3.0` decades, fixing suppression on wide-range power-law equations.

Table ?? (Section D) summarises all aggregate recovery changes. All results reported in the paper use the v2 corrected harness. The v1 JSON output files are retained for audit purposes and are available upon request.

Per-Equation v2 Data

The complete per-equation R^2 values under the v2 corrected harness are stored in:

- `protocol_core_noiseless_20260304_154510.json` — Phase 3 noiseless, all 6 methods, 30 equations.
- `noise_sweep_20260316_192711.json` — M3/M4 noise sweep ($\sigma \in \{0, 0.5, 1, 5, 10\}\%$, v2 corrected).
- `sample_complexity_20260316_193447.json` — M3/M4 sample complexity sweep ($n \in \{50, 100, 200, 500, 750, 1000\}$, v2 corrected).

These files will be released as supplementary material. The v1 superseded files (`noise_sweep_20260315_091018.json`, `sample_complexity_20260315_124310.json`) should not be used for the final submission figures.

PureLLM Disclosure

Following Satvaty et al. (2024), each PureLLM noiseless pass was audited for memorisation. **11 of 30 PureLLM passes are flagged is hardcoded=True**: Michaelis–Menten, Logistic growth, Allometric scaling, Nernst, Clausius–Mossotti, Lorentz force, Gaussian PDF, Zeeman energy, Bose–Einstein, Planck radiance, and Fourier heat. These equations return the ground-truth formula verbatim with zero cross-seed variation, consistent with training-data recall rather than inference. The genuine-inference pass rate (19 non-hardcoded equations) is 100%. PureLLM is therefore listed for completeness in all tables but excluded from all scientific comparisons (Section ??).

A controlled re-run with `--no-llm-cache` confirmed that all 30 equations issued independent API calls and produced identical scores (median $R^2 = 0.9976$, std = 2.5×10^{-4}), ruling out any caching artefact.

Citation Key Corrections

The following citation key conflicts have been resolved in `references.bib` (updated 2026-06):

1. `cranmer2023interpretable` renamed to `cranmer2023pysr` (canonical key matching all `\cite` calls in the main paper; same paper arXiv:2305.01582). **Resolved.**
2. `udrescu2020aifeynman` duplicate removed; `udrescu2020ai` retained as the single canonical key in `references.bib`. This supplementary document uses `udrescu2020aifeynman` in its own inline `thebibliography` and remains self-consistent. **Resolved.**
3. `meidani2023snip` (arXiv preprint key) replaced by `meidani2024snip` (published ICLR 2024 key) in `references.bib`. **Resolved.**
4. `lacava2021contemporary` entry confirmed present in `references.bib`. **No action needed.**

SNIP Comparison Note

SNIP (Meidani et al., 2024) does *not* report recovery rates. It evaluates symbolic regression using a Pareto front of median R^2 vs. equation complexity on 119 Feynman equations, reporting median $R^2 = 0.882$ (vs. AI Feynman baseline: 0.798). Placing SNIP’s median R^2 alongside HypatiaX’s recovery rate (% at $R^2 > 0.9999$) would be a category error. The only valid comparison uses the shared median R^2 axis: HypatiaX Hybrid DeFi achieves median $R^2 = 1.0000$ on our 30-equation subset vs. SNIP’s 0.882 on the full 119-equation SRBench set (different subsets; comparison is indicative only). Table ?? provides this comparison.

Appendix E. Hyperparameter Settings

Table 17: PySR Configuration for Symbolic Regression

Parameter	Value
Population size	100
Generations	500
Populations (parallel)	2 (hardware limit: 2-core CPU)
Crossover probability	0.8
Mutation probability	0.1
Complexity penalty	0.01
Tournament size	3
Elitism	5
Max expression depth	15
Optimization iterations	50
Validation Thresholds	
Acceptance R^2 (Criterion 1)	0.99
Validation score (Criterion 1)	30
Acceptance R^2 (Criterion 2)	0.95
Validation score (Criterion 2)	80
Extrapolation R^2	0.80
Numerical stability (α)	0.95

Table 18: Complete Hyperparameter Specifications for All Methods

Method	Parameter	Value
Neural Network	Architecture	[64, 32, 1]
	Activation	ReLU
	Optimizer	Adam
	Learning rate	0.01
	Batch size	Full batch (160 samples)
	Epochs	200
	Weight initialization	Xavier uniform
	Dropout	None
Pure LLM	Model	claude-sonnet-4-20250514
	Temperature	0.0 (deterministic)
	Max tokens	1024
	Top-p	1.0
	Few-shot examples	2-3 per domain
	Prompt template	Domain-specific
	Retry on failure	3 attempts
Hybrid v50.2	SR iterations	50
	SR populations	2 (hardware limit: 2-core CPU)
	Binary operators	[+, -, *, /]
	Unary operators	[exp, log, sqrt, sin, cos]
	Parsimony coefficient	0.01
	Tournament size	10
	Mutation rate	0.1
	Crossover rate	0.9
	LLM candidates	5
	LLM temperature	0.3
	Max complexity	30
	Validation threshold (R^2)	0.99

Appendix F. Detailed Experimental Results

F.1 Pure LLM Baseline (40 tests)

Table 19: Pure LLM Performance by Test Case

Domain	Test	R^2	Result	Failure Mode
<i>Classical Scientific Domains (20 tests)</i>				
Materials	Stress-strain	1.00	Pass	—
Materials	Young’s modulus	1.00	Pass	—
Materials	Poisson ratio	1.00	Pass	—
Materials	Thermal expansion	1.00	Pass	—
Fluid	Reynolds number	1.00	Pass	—
Fluid	Bernoulli equation	1.00	Pass	—
Fluid	Drag coefficient	1.00	Pass	—
Fluid	Flow rate	1.00	Pass	—
Thermo	Ideal gas law	1.00	Pass	—
Thermo	Heat capacity	1.00	Pass	—
Thermo	Carnot efficiency	1.00	Pass	—
Thermo	Stefan-Boltzmann	1.00	Pass	—
Mechanics	Newton’s 2nd law	1.00	Pass	—
Mechanics	Kinetic energy	1.00	Pass	—
Mechanics	Gravitational force	1.00	Pass	—
Mechanics	Hooke’s law	1.00	Pass	—
Chemistry	Arrhenius equation	1.00	Pass	—
Chemistry	pH calculation	1.00	Pass	—
Chemistry	Nernst equation	0.50	Fail	Wrong sign convention
Chemistry	Rate law	1.00	Pass	—
<i>Decentralized Finance Domains (20 tests)</i>				
AMM	Constant product	1.00	Pass	—
AMM	Price impact	0.98	Pass	—
AMM	Impermanent loss	−1.26	Fail	Unit inconsistency (missing $\times 100$)
AMM	Swap fee	0.82	Pass	—
VaR	Parametric VaR 95%	−10.05	Fail	Wrong sign (z-score)
VaR	Parametric VaR 99%	1.00	Pass	—
VaR	Historical VaR	0.95	Pass	—
VaR	Modified VaR	−2.15	Fail	Scipy API misuse
Liquidity	APY calculation	1.00	Pass	—
Liquidity	Capital efficiency	0.87	Pass	—
Liquidity	Fee earnings	—	Fail	Non-executable (tutorial mode)
Liquidity	Utilization ratio	—	Fail	Non-executable (tutorial mode)
ES	Expected Shortfall 95%	−4.12	Fail	Scipy <code>norm.pdf</code> returns object
ES	Expected Shortfall 99%	−3.87	Fail	Distribution tail error
ES	Conditional VaR	0.88	Pass	—
ES	Tail risk	−1.92	Fail	Statistical definition error
Liquidation	Long position	—	Fail	Tutorial text only
Liquidation	Short position	—	Fail	Tutorial text only
Liquidation	Margin call	—	Fail	Incomplete derivation
Liquidation	Leverage ratio	—	Fail	Multi-step missing
Classical Average:		19/20 (95%)		
DeFi Average:		21 8/20 (40%)		
Overall:		27/40 (67.5%)		

Key Findings: Pure LLM performance strongly correlates with equation prevalence in training data (Spearman $\rho = 0.89$, $p < 0.001$)(?). Classical equations achieve 95% success rate while specialised domains (DeFi) drop to 40%, consistent with the reproductive behaviour characterisation in Finding ??.

F.2 DeFi-Optimized Results (23 tests)

Table 20: DeFi-Optimized HypatiaX Performance

Test	R^2	Validation	Time (s)	Extrap.	Result
Constant product (xy=k)	1.0000	100.0	45	Pass	Pass
Price impact	0.9998	95.2	120	Pass	Pass
Impermanent loss	0.9995	92.1	180	Pass	Pass
Swap fee calculation	1.0000	98.5	60	Pass	Pass
VaR 95% parametric	0.9988	88.3	150	Pass	Pass
VaR 99% parametric	0.9990	90.1	145	Pass	Pass
Historical VaR	0.9985	87.5	135	Pass	Pass
Modified VaR	0.9982	85.9	160	Pass	Pass
APY simple	1.0000	100.0	50	Pass	Pass
APY compound	0.9995	94.8	110	Pass	Pass
Capital efficiency	0.9940	78.5	240	Pass	Pass
Fee earnings	0.9998	96.2	90	Pass	Pass
Expected Shortfall 95%	0.9987	89.7	170	Pass	Pass
Expected Shortfall 99%	0.9985	88.2	175	Pass	Pass
Conditional VaR	0.9990	91.4	155	Pass	Pass
Tail risk metric	0.9983	86.8	165	Pass	Pass
Liquidation (long)	0.9992	93.6	125	—	Pass
Liquidation (short)	0.9994	94.1	120	—	Pass
Margin call threshold	0.9996	95.8	105	—	Pass
Effective leverage	0.9998	97.2	85	—	Pass
Staking rewards	1.0000	100.0	55	—	Pass
Dilution effect	0.9993	93.8	115	—	Pass
Kelly criterion	0.9988	89.5	195	Pass	Pass
Average	0.9990	91.3	115.2	16/16	23/23

Extrapolation Tests: Sixteen tests designated for extrapolation validation all passed with average extrapolation $R^2 = 0.9650$, indicating reliable generalization to the tested novel regimes on this benchmark(Balestrierio et al., 2021; Neyshabur et al., 2017; Zhang et al., 2021).

F.3 Complete Test Suite Details

Table 21: Ground Truth Formulas with Exact Constants

Equation	Formula	Constants
Arrhenius	$k = A \exp(-E_a/(RT))$	$A = 10^{11}$, $E_a = 80000$ J/mol, $R = 8.314$ J/(mol·K)
Henderson-Hasselbalch	$\text{pH} = \text{p}K_a + \log_{10}([A^-]/[\text{HA}])$	$\text{p}K_a = 6.5$
Rate Law	$r = k[A]^2[B]$	$k = 0.5$ M ⁻² s ⁻¹
Allometric Scaling	$Y = aM^b$	$a = 3.5$, $b = 0.75$
Michaelis-Menten	$v = V_{\max}[S]/(K_m + [S])$	$V_{\max} = 50$ μM/s, $K_m = 10$ μM
Logistic Growth	$dN/dt = rN(1 - N/K)$	$r = 0.3$ day ⁻¹ , $K = 1000$
Kinetic Energy	$E = \frac{1}{2}mv^2$	Dimensionless
Gravitational Force	$F = Gm_1m_2/r^2$	$G = 6.674 \times 10^{-11}$ N·m ² /kg ²
Ideal Gas Law	$P = nRT/V$	$R = 8.314$ J/(mol·K)
Impermanent Loss	$\text{IL} = 2\sqrt{r}/(1 + r) - 1$	Dimensionless
Price Impact	$\Delta p = \Delta x/(x + \Delta x)$	Dimensionless
Constant Product	$y = k/x$	$k = 10^6$
VaR 95%	$\text{VaR} = P\sigma \cdot z_{0.95}$	$z_{0.95} = 1.645$
Liquidation Price	$p_{\text{liq}} = p_0(1 - 1/(L \cdot m))$	$m = 0.8$ (maintenance margin)
Portfolio Variance	$\sigma_p = \sqrt{\sigma_1^2 + \sigma_2^2 + 2\rho\sigma_1\sigma_2}$	Dimensionless

Appendix G. Statistical Test Details

G.1 Mann-Whitney U Test for Extrapolation Performance

For medium extrapolation regime ($2 \times$ training range):

Test Setup:

- **Sample sizes:** $n_{\text{Hybrid}} = 14$, $n_{\text{NN}} = 13$ (CORRECTED)
- **Null hypothesis:** $H_0 : E_{\text{Hybrid}} \geq E_{\text{NN}}$ (Hybrid is not better)
- **Alternative hypothesis:** $H_1 : E_{\text{Hybrid}} < E_{\text{NN}}$ (Hybrid is better)
- **Significance level:** $\alpha = 0.05$ (one-tailed test)

Descriptive Statistics:

- **Hybrid v50_2:** Mean = 0.0%, Median = 0.0%, SD = 0.0%, Range = [0.0, 0.0]
- **Neural Network:** Mean = 1231%, Median = 86.7%, SD = 1842%, Range = [12.3, 5847%] (CORRECTED)

Test Results:

- **Test statistic:** $U = 0$ (all Hybrid errors < all NN errors)
- **P-value:** $p = 1.11 \times 10^{-6}$ (CORRECTED)
- **Critical value:** $U_{0.05} = 45$ for one-tailed test with $n=14$ vs $n=13$
- **Conclusion:** $U < U_{0.05} \Rightarrow$ Reject H_0 at $\alpha = 0.05$
- **Interpretation:** Complete rank separation — on every tested equation, Hybrid v50_2 extrapolation error is strictly less than Neural Network error ($U = 0$, $p = 1.11 \times 10^{-6}$)

G.2 Effect Size Calculation (Cohen’s d)

$$\mu_{\text{NN}} = 1231.0\%, \quad \sigma_{\text{NN}} = 1842.0\% \quad (\text{CORRECTED}) \quad (1)$$

$$\mu_{\text{Hybrid}} = 0.0\%, \quad \sigma_{\text{Hybrid}} = 0.0\% \quad (2)$$

$$s_{\text{pooled}} = \sqrt{\frac{\sigma_{\text{NN}}^2 + \sigma_{\text{Hybrid}}^2}{2}} = 1302.5\% \quad (3)$$

$$d = \frac{1231.0 - 0.0}{1302.5} = 0.95 \quad (4)$$

Conservative Estimate: Using only NN standard deviation (most conservative approach):

$$d_{\text{conservative}} = \frac{1231.0}{1842.0} = 0.67 \quad (\text{medium-to-large effect}) \quad (5)$$

Alternative Calculation (Glass’s Δ): Using only the control group (NN) standard deviation:

$$\Delta = \frac{1231.0 - 0.0}{1842.0} = 0.67 \quad (6)$$

Interpretation Guidelines:

- $d = 0.2$: Small effect
- $d = 0.5$: Medium effect
- $d = 0.8$: Large effect
- $d = 1.2$: Very large effect
- Our result: $d = 0.67 - 0.95$ (medium-to-large effect with full rank separation on this benchmark)

Critical Note on Effect Size: While the calculated Cohen’s d appears moderate (0.67-0.95), this understates the true difference because:

1. Mann-Whitney $U = 0$ indicates full rank separation—every Hybrid error is strictly less than every NN error
2. Effect size calculations assume normal distributions, but NN errors are highly skewed (median = 86.7%, mean = 1231%)
3. The practical difference (near-zero vs 1,231% mean error on this benchmark) is substantial and potentially critical for scientific applications

G.3 Interpolation Performance Comparison

Kruskal-Wallis H Test:

- **Test statistic:** $H = 42.3$
- **P-value:** $p < 0.001$
- **Interpretation:** Significant differences exist across methods

Table 22: Interpolation Performance Statistics by System

System	n	Mean R^2	Median R^2	SD	Range
System 3 (LLM+Fallback)	30	1.000	1.000	0.000	[1.00, 1.00]
Hybrid v50_2	15	0.931	1.000	0.147	[0.58, 1.00]
Neural Network	15	0.940	0.952	0.082	[0.78, 1.00]
Pure PySR	18	0.982	0.998	0.045	[0.88, 1.00]

System Comparisons:

Key Observation: System 3 achieves $R^2 = 1.000$ on interpolation ($n=30$) but was not designed for extrapolation testing. Hybrid v50_2’s median (1.000) exceeds its mean (0.931) due to outlier presence, making median more representative of typical performance.

G.4 Domain-Specific Success Rates

Success rate defined as $R^2 > 0.95$:

Table 23: Success Rate by Domain and Method

Method	Physics	Chemistry	Biology	DeFi	Economics
Hybrid v50_2	100%	100%	67%	100%	67%
Pure PySR	100%	100%	100%	67%	67%
LLM-Guided PySR	100%	67%	67%	67%	67%
Neural Network	100%	67%	33%	67%	33%
Pure LLM	100%	33%	0%	33%	33%
Avg across methods	100%	73%	53%	67%	53%

Correlation Analysis: Pure LLM success rate strongly correlates with equation prevalence in web text (Spearman $\rho = 0.89$, $p < 0.001$)(?). This is consistent with the reproductive behaviour hypothesis: performance is higher where training-corpus exposure is greater.

G.5 Validation Layer Effectiveness

Error detection breakdown across four-layer validation framework:

Table 24: Validation Layer Error Detection Statistics

Layer	Errors	%	Cum.	Example
1. Dimensional Analysis	18	12	12	Energy = mass
2. Physics Constraints	27	18	30	Negative rate
3. Numerical Stability	35	23	53	Overflow
4. Prediction Accuracy	70	47	100	Wrong result
Total	150	100	—	—

Key Finding: Single-layer validation (prediction accuracy only) would miss 53% of errors caught by earlier layers on this benchmark, indicating that cascading validation contributed substantially to error detection for LLM-generated outputs.

Appendix H. Discovered Expression Examples

H.1 Perfect Recoveries

H.1.1 KINETIC ENERGY

Example 1 (Kinetic Energy Discovery) *Ground truth:* $E = \frac{1}{2}mv^2$

HypatiaX v50_2 discovered:

$$E = ((v \times m) \times v) \times 0.5 = 0.5mv^2 \quad (7)$$

Analysis: Exact recovery of ground truth. Discovered expression is algebraically identical.

Interpolation: $R^2 = 1.0000$ on training data ($v \in [1, 10]$ m/s)

Extrapolation: 0.0% error at $v = 50$ m/s ($5\times$ training max)

Physical interpretation: Correctly captures quadratic dependence on velocity and linear dependence on mass.

H.1.2 IDEAL GAS LAW

Example 2 (Ideal Gas Law Discovery) *Ground truth:* $P = nRT/V$, where $R = 8.314 \text{ J}/(\text{mol}\cdot\text{K})$

HypatiaX v50_2 discovered:

$$P = n \times (T \times (8.314/V)) \quad (8)$$

Analysis: Exact recovery including the universal gas constant $R = 8.314$ to full precision.

Interpolation: $R^2 = 1.0000$ on training data

Extrapolation: 0.0% error across all pressure/temperature/volume ranges tested (up to $5\times$ training range)

Physical interpretation: On this test case, symbolic regression recovered a constant consistent with the gravitational constant from data alone, without prior knowledge of its value(?).

H.2 Near-Perfect Recoveries

The Michaelis-Menten and Impermanent Loss equations were recovered with constant error $< 10^{-5}$, achieving $R^2 = 1.0000$ and near-zero extrapolation error (see Listing ??).

H.3 Failure Case Analysis

H.3.1 GRAVITATIONAL FORCE (BOTH METHODS FAILED)

Example 3 (Gravitational Force Failure) *Ground truth:* $F = Gm_1m_2/r^2$, where $G = 6.674 \times 10^{-11} \text{ N}\cdot\text{m}^2/\text{kg}^2$

Neural Network Result:

- **Interpolation:** $R^2 = 0.21$ (poor fit; relative error $> 100\%$, meeting the threshold in Definition ??)
- **Extrapolation:** Not attempted due to interpolation failure

- **Issue:** Gradient vanishing due to 10^{-11} scale
- **Attempted solution:** Log-transform $\rightarrow$ marginal improvement ($R^2 = 0.35$)
- **Root cause:** Optimization landscapes too flat at this scale for gradient descent

HypatiaX v50_2 Result:

- **Interpolation:** $R^2 = -0.03$ (worse than baseline)
- **Extrapolation:** Not attempted due to interpolation failure
- **Discovered form:** $F = 1.23 \times 10^{13} \frac{m_1 m_2}{r^2}$ (correct functional form, wrong constant by $2 \times 10^{23} \times$)
- **Issue:** PySR constant search range $[10^{-3}, 10^3]$ missed $G \sim 10^{-11}$
- **Attempted solution:** Extend range to $[10^{-15}, 10^{15}] \rightarrow$ no convergence in 50 iterations
- **Root cause:** Search space too large for evolutionary algorithm to explore efficiently

Recommended Solution:

1. **Dimensional analysis preprocessing:** Identify required constant scales before symbolic regression
2. **Logarithmic transformation:** Convert $F = Gm_1m_2/r^2$ to $\log F = \log G + \log m_1 + \log m_2 - 2\log r$, making the problem linear in log-space
3. **Adaptive constant ranges:** Dynamically adjust search based on data magnitude
4. **Multi-scale normalization:** Normalize inputs/outputs to $[0.1, 10]$ range before discovery, then rescale

Lesson learned: Very small constants ($< 10^{-6}$) require special handling in both neural and symbolic methods. This represents 1/15 test cases (6.7% failure rate) and is a known limitation addressed in Section 12.

H.3.2 HENDERSON-HASSELBALCH (PURE LLM FAILURE)

Example 4 (Henderson-Hasselbalch pH Calculation) *Ground truth:* $pH = pKa + \log_{10} \left(\frac{[A^-]}{[HA]} \right)$

Pure LLM Generated:

$$pH = pKa \times \frac{[A^-]}{[HA]} \quad (9)$$

Failure Analysis:

- **R^2 score:** -12.3 (predictions anticorrelated with ground truth)
- **Errors:**

1. Multiplication instead of addition

2. Missing logarithm entirely
 3. Plausible but fundamentally wrong
- **Root cause:** LLM training data likely contains more qualitative discussions of pH (“pH increases when conjugate base increases”) than quantitative formulas
 - **Pattern matching failure:** LLM defaulted to simpler multiplicative relationship rather than logarithmic form

Symbolic Rescue: PySR correctly discovered logarithmic relationship:

$$pH = pKa + \log_{10} \left(\frac{[A^-]}{[HA]} \right) \quad (10)$$

Achieving $R^2 = 0.998$, demonstrating the value of hybrid architecture.

Lesson learned: LLM performance declined on this specialised equation, which is consistent with Finding ??—generation probability is theoretically predicted to be lower for expressions that are structurally distant from the training corpus(?).

H.4 Complete Discovered Expressions Code Listing

```

1 # Kinetic Energy (PERFECT)
2 discovered: ((v * m) * v) * 0.5
3 ground_truth: 0.5 * m * v**2
4 R2: 1.0000, extrapolation_error: 0.0%
5
6 # Rate Law (PERFECT)
7 discovered: ((A_conc * 0.5) * A_conc) * B_conc
8 ground_truth: 0.5 * A_conc**2 * B_conc
9 R2: 1.0000, extrapolation_error: 0.0%
10
11 # Ideal Gas Law (PERFECT)
12 discovered: n * (T * (8.314 / V))
13 ground_truth: n * 8.314 * T / V
14 R2: 1.0000, extrapolation_error: 0.0%
15
16 # Logistic Growth (PERFECT - algebraically equivalent)
17 discovered: ((N * -0.0003) + 0.3) * N
18 ground_truth: 0.3 * N * (1 - N/1000)
19 expanded: 0.3*N - 0.0003*N^2 (identical)
20 R2: 1.0000, extrapolation_error: 0.0%
21
22 # Price Impact (PERFECT)
23 discovered: swap / (swap + reserve)
24 ground_truth: dx / (x + dx)
25 R2: 1.0000, extrapolation_error: 0.0%
26
27 # Michaelis-Menten (NEAR-PERFECT)
28 discovered: (S / (S + 10.000007)) * 50.000008
29 ground_truth: 50 * S / (10 + S)

```

```

30 constant_error: <1e-5
31 R2: 1.0000, extrapolation_error: 0.0%
32
33 # Impermanent Loss (NEAR-PERFECT)
34 discovered: (2 * sqrt(r)) / (1 + r) - 1.000023
35 ground_truth: 2 * sqrt(r) / (1 + r) - 1
36 constant_error: 2.3e-5
37 R2: 1.0000, extrapolation_error: 0.0%

```

Listing 1: Discovered Symbolic Expressions (7 Perfect Recoveries)

Appendix I. Supplementary Materials

The following supplementary materials are available:

- **S1: Pure LLM Baseline Results** (`S1_pure_llm_results.csv`) — Complete results from Pure LLM baseline experiments showing 60% success rate across 40 tests.
- **S2: Extrapolation Test Data** (`S2_extrapolation_data.csv`) — Data comparing Hybrid v50_2 extrapolation error (near-zero) versus Neural Network extrapolation error (1,231% mean) on the Core 15 benchmark. Includes training R^2 scores, extrapolation errors at $1.2\times$, $2\times$, and $5\times$ training ranges, and Mann-Whitney U-test results ($U=0$, $p<10^{-6}$).
- **S3: Validation Logs** (`S3_validation_logs.json`) — Complete validation reports for all 131 tests showing the four-layer validation framework in action. Each entry includes dimensional analysis, physics constraints, numerical stability checks, prediction accuracy metrics, timing data, and failure mode classification. Includes validation layer rejection rates (12%, 18%, 23%, 47%).
- **S4: Discovered Symbolic Expressions** (`S4_discovered_expressions.{txt,json}`) — All 131 discovered symbolic expressions in SymPy-compatible format, organized by domain (chemistry, biology, physics, DeFi, finance) with complete metadata including ground truth equations, performance metrics (R^2 , timing, iterations), and algebraic equivalence verification.
- **S5: Timing Breakdown** (`S5_timing_breakdown.csv`) — Detailed computational cost analysis showing component-level timing: LLM initialization (0-39%), symbolic search (50-91%), and validation (7-11%). Demonstrates the $4.3\times$ speedup from LLM-guided initialization versus pure symbolic methods.
- **S6: Domain-Specific Analysis** (`S6_domain_analysis.xlsx`) — Statistical analysis by scientific domain showing success rates: Biology 100%, Chemistry 100%, Physics 100%, DeFi 100%, Finance 77.8%. Includes per-equation breakdowns and domain-specific constraint validation results.
- **S7: Failure Mode Taxonomy** (`S7_failure_mode_taxonomy.pdf`) — Comprehensive taxonomy of 23 distinct failure modes categorized by method (Pure LLM: 8 modes,

Neural Network: 3 modes, Symbolic: 5 modes, Hybrid: 2 modes, Integration: 5 modes), severity (8 Critical, 4 High, 9 Medium, 2 Low), and detection layer. Includes detailed examples of silent semantic errors, large-scale extrapolation errors, distributional reasoning failures, and API misuse.

- **S8: Figure Source Data** (`S8_figures_source_data/`) — Source data files for all figures in the paper:
 - `figure1_arrhenius_extrapolation.csv` — Arrhenius equation extrapolation comparison
 - `figure2_domain_comparison.csv` — Success rates by scientific domain
 - `figure3_validation_breakdown.csv` — Four-layer validation effectiveness
 - `figure4b_domain_breakdown.csv` — Per-equation performance details
 - `figure5_method_comparison.csv` — Overall method comparison
 - `figure_5systems_comparison.csv` — Complete extrapolation data (1231% vs 0%)
 - `figure_benchmark_comparison.csv` — Core 15 benchmark results
 - `figure_timing_comparison.csv` — Speed-accuracy tradeoff analysis

All supplementary materials are provided in both synthetic versions (matching reported statistics exactly) and real experimental versions (direct from JSON logs) for maximum reproducibility and transparency. Complete documentation and usage examples are provided in accompanying README files.

I.1 Supplementary Data Files

All supplementary data are available at:

<https://github.com/sednabcn/LLM-HypatiaX-PAPERS>

Table 25: Supplementary Data Files

File	Contents
<code>S1_pure_llm_results.csv</code>	Complete Pure LLM test results ($n=40$) with formulas, R^2 scores, failure modes
<code>S2_extrapolation_data.csv</code>	Extrapolation test data ($n = 27$) at $1.2\times$, $2\times$, $5\times$ training ranges
<code>S3_validation_logs.json</code>	Full validation reports for all 131 tests
<code>S4_discovered_expressions.txt</code>	All discovered symbolic expressions in SymPy format
<code>S5_timing_breakdown.csv</code>	Detailed timing data by component and method
<code>S6_domain_analysis.xlsx</code>	Per-domain performance with statistical tests
<code>S7_failure_taxonomy.pdf</code>	Complete taxonomy of 23 failure modes with examples
<code>S8_figures_source_data/</code>	Raw data for all 8 figures in CSV format

I.2 Reproducibility Materials

Complete step-by-step tutorials for reproducing all experiments are available at: <https://sednabcn.github.io/ai-llm-blog/tutorials/hypatiax/>

- Tutorial 1: Environment Setup and First Discovery (15 min)
- Tutorial 2: Running Benchmark Experiments (45 min active + 3–8 hours compute)
- Tutorial 3: Statistical Analysis and Publication Figures (30 min)
- Tutorial 4: Custom Applications and Extensions (45 min)

All source code and data generation scripts are available at: <https://github.com/sednabcn/LLM-HypatiaX-PAPERS>

I.3 Reproducibility Artifacts

Docker Images:

- `hypatiax:v1.0` - Full environment with all dependencies
- `hypatiax:minimal` - Core symbolic regression only (no LLM)
- `hypatiax:gpu` - GPU-accelerated version for neural network comparison

Source Code: All Python scripts and experimental protocols are available at:

<https://github.com/sednabcn/LLM-HypatiaX-PAPERS/tree/master/papers/2025-JMLR/hypatiax>

Jupyter Notebooks:

- `01_data_generation.ipynb` - Synthetic data generation for 15 equations
- `02_pure_llm_experiments.ipynb` - Pure LLM baseline (40 tests)
- `03_hybrid_experiments.ipynb` - Hybrid v50_2 experiments (30 tests)
- `04_extrapolation_analysis.ipynb` - Extrapolation tests and statistical analysis
- `05_figure_generation.ipynb` - Reproduction of all 8 figures
- `06_statistical_tests.ipynb` - Mann-Whitney U, Kruskal-Wallis, effect sizes

All Python Jupyter notebooks are available at:

<https://github.com/sednabcn/LLM-HypatiaX-PAPERS/tree/master/papers/2025-JMLR/hypatiax/notebooks>

Step-by-Step Tutorials:

- Tutorial 1: Environment Setup and First Discovery (15 min)
- Tutorial 2: Running Benchmark Experiments (45 min active + 3–8 hours compute)
- Tutorial 3: Statistical Analysis and Publication Figures (30 min)
- Tutorial 4: Custom Applications and Extensions (45 min)

Available at: <https://sednabcn.github.io/ai-llm-blog/tutorials/hypatiax/>

I.4 Additional Statistical Tests

Table 26: Additional Statistical Tests

Comparison	Test	Result	Interpretation
Hybrid v50.2 vs Pure PySR	Mann-Whitney	$U = 89, p = 0.18$	No significant difference
System 3 vs Hybrid v50.2	Mann-Whitney	$U = 156, p = 0.03$	System 3 better (interp.)
Classical vs DeFi (LLM)	Fisher’s Exact	$p < 0.001$	Domain effect significant
Run-to-run variability	Friedman	$\chi^2 = 34.2, p < 0.001$	High variability

Power Analysis:

- Extrapolation comparison (n=14 vs n=13): Power = 1.000 for $d \geq 0.5$
- Domain comparisons (n=3 per domain): Power = 0.65 for $d \geq 0.8$
- Method comparisons (n=15 per method): Power = 0.95 for $d \geq 0.6$

Recommendation: Future work should increase per-domain sample sizes to $n \geq 10$ for robust inference on secondary effects.

Appendix J. Noise Sensitivity Analysis

To assess correspondence with Condition 3 of Empirical Result ?? (noise tolerance threshold $\sigma \leq 0.1 \cdot \text{std}(y)$), we conducted systematic experiments varying noise levels.

J.1 Experimental Protocol

We tested 15 classical formulas (Table ??) with noise levels:

$$\sigma \in \{0.01, 0.03, 0.05, 0.08, 0.10, 0.15, 0.20\} \times \text{std}(y)$$

For each noise level, we:

1. Generated $n = 100$ training points with additive Gaussian noise: $y_{\text{noisy}} = y_{\text{true}} + \mathcal{N}(0, \sigma^2)$
2. Ran PySR with standard hyperparameters (Appendix E)
3. Evaluated success as $R^2 \geq 0.99$ and correct functional form
4. Repeated 5 times per configuration to account for evolutionary randomness

J.2 Results

Table 27: Success Rate vs Noise Level

Noise Level ($\sigma/\text{std}(y)$)	Success Rate	95% CI
0.01	100% (75/75)	[95.2%, 100%]
0.03	97.3% (73/75)	[90.7%, 99.7%]
0.05	93.3% (70/75)	[85.1%, 97.8%]
0.08	78.7% (59/75)	[67.7%, 87.3%]
0.10	66.7% (50/75)	[54.8%, 77.1%]
0.15	45.3% (34/75)	[33.8%, 57.3%]
0.20	24.0% (18/75)	[14.9%, 35.3%]

J.3 Key Findings

Sharp Threshold at $\sigma = 0.1$: Success rate drops from 93.3% ($\sigma = 0.05$) to 45.3% ($\sigma = 0.15$), consistent with the empirically derived threshold. Fisher’s exact test shows a significant difference between $\sigma \leq 0.10$ and $\sigma > 0.10$ ($p < 0.001$).

Domain Variability: Classical mechanics formulas (kinetic energy, potential energy) tolerate higher noise ($\sigma \leq 0.15$) while chemistry formulas (Arrhenius, rate law) fail above $\sigma = 0.08$. This suggests domain-specific complexity affects noise sensitivity.

Failure Modes: Above threshold, common failures include:

- Incorrect exponents (e.g., $v^{1.8}$ instead of v^2 for kinetic energy)
- Spurious additive constants to compensate for noise
- Overfitting with complex expressions ($R^2 > 0.99$ on training, poor extrapolation)

J.4 Implications

Our main experiments use $\sigma = 0.05 \cdot \text{std}(y)$, safely below the threshold, ensuring 93.3% baseline success rate from noise alone. Real-world applications should:

1. Estimate noise level from repeated measurements
2. Apply denoising if $\sigma > 0.08 \cdot \text{std}(y)$
3. Increase sample size by factor of $(\sigma/0.05)^2$ to maintain success rate